

Sri Sivasubramaniya Nadar College of Engineering, Chennai
(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester	V
Subject Code & Name	ICS1512 & Machine Learning Algorithms Laboratory		
Academic year	2025-2026 (Odd)	Batch:2023-2028	Due date: 27.01.2025

Experiment 1: Working with Python packages-Numpy, Scipy, Scikit-Learn, Matplotlib

Aim: To study and explore the fundamental Python libraries used in data science and machine learning, and to understand how different datasets can be analyzed and mapped to appropriate machine learning models using exploratory data analysis techniques.

Libraries Used:

- **NumPy**: Used for numerical computations and efficient handling of multi-dimensional arrays.
- **Pandas**: Used for data manipulation, cleaning, and analysis using DataFrames.
- **Matplotlib**: Used for creating visualizations such as line graphs, bar charts, and histograms.
- **Seaborn**: Used for advanced statistical data visualization with attractive and informative plots.

Objectives performed:

1 Iris Dataset

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder

from sklearn.feature_selection import SelectKBest, chi2
from sklearn.feature_selection import f_classif
from sklearn.model_selection import train_test_split
```

```
[5]: df=pd.read_csv('Datasets/Iris.csv')
df.drop(columns=['Id'], inplace=True)

df.head()
```

```
[5]:    SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm      Species
0           5.1          3.5          1.4          0.2  Iris-setosa
1           4.9          3.0          1.4          0.2  Iris-setosa
2           4.7          3.2          1.3          0.2  Iris-setosa
3           4.6          3.1          1.5          0.2  Iris-setosa
4           5.0          3.6          1.4          0.2  Iris-setosa
```

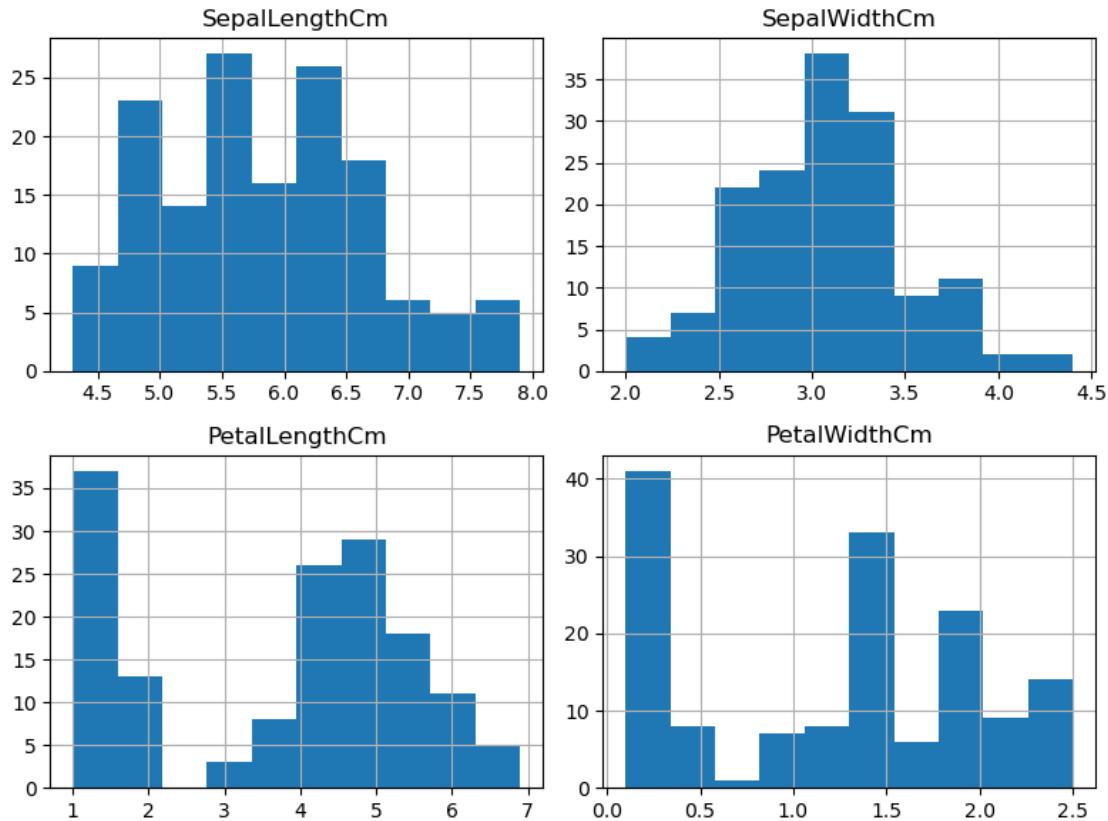
```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   SepalLengthCm    150 non-null    float64
 1   SepalWidthCm     150 non-null    float64
 2   PetalLengthCm   150 non-null    float64
 3   PetalWidthCm    150 non-null    float64
 4   Species          150 non-null    object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
[7]: df.describe()
```

```
[7]:    SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm
count      150.000000    150.000000    150.000000    150.000000
mean       5.843333     3.054000     3.758667     1.198667
std        0.828066     0.433594     1.764420     0.763161
min        4.300000     2.000000     1.000000     0.100000
25%        5.100000     2.800000     1.600000     0.300000
50%        5.800000     3.000000     4.350000     1.300000
75%        6.400000     3.300000     5.100000     1.800000
max        7.900000     4.400000     6.900000     2.500000
```

```
[8]: df[df.columns].hist(figsize=(8,6))
plt.tight_layout()
plt.savefig("histogram.png", dpi=300, bbox_inches="tight")
plt.show()
```

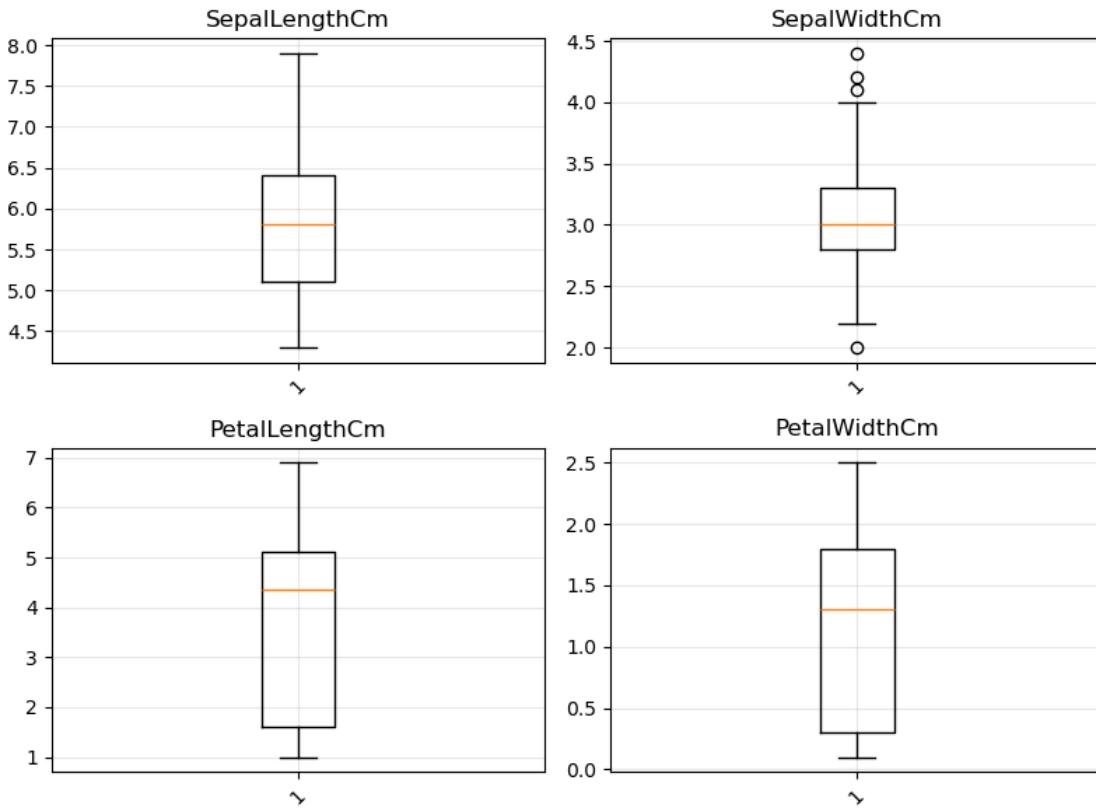


```
[9]: num_cols = df.select_dtypes(include=['int64', 'float64']).columns

fig, axes = plt.subplots(2,2, figsize = (8,6))
axes = axes.flatten()

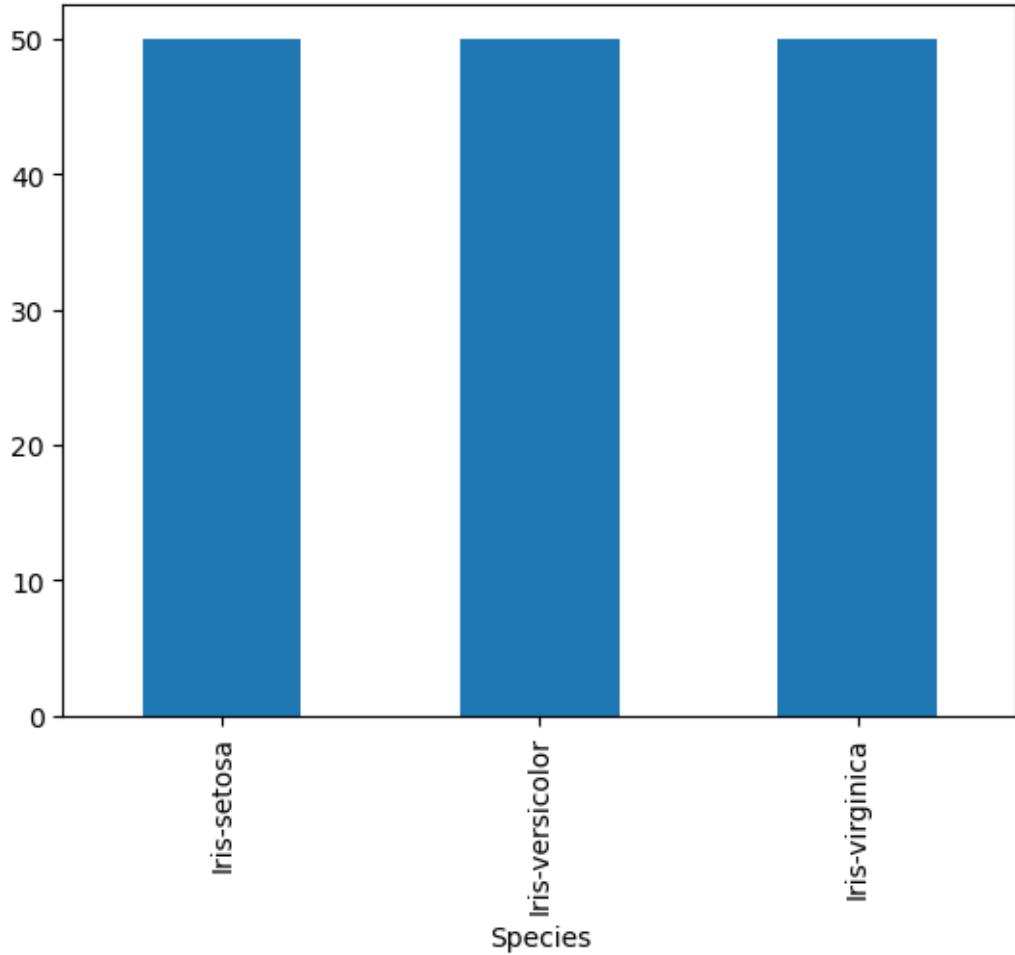
for i, col in enumerate(num_cols):
    axes[i].boxplot(df[col])
    axes[i].set_title(col)
    axes[i].tick_params(axis='x', rotation=45)
    axes[i].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("box_plot.png", dpi=300, bbox_inches="tight")
plt.show()
```

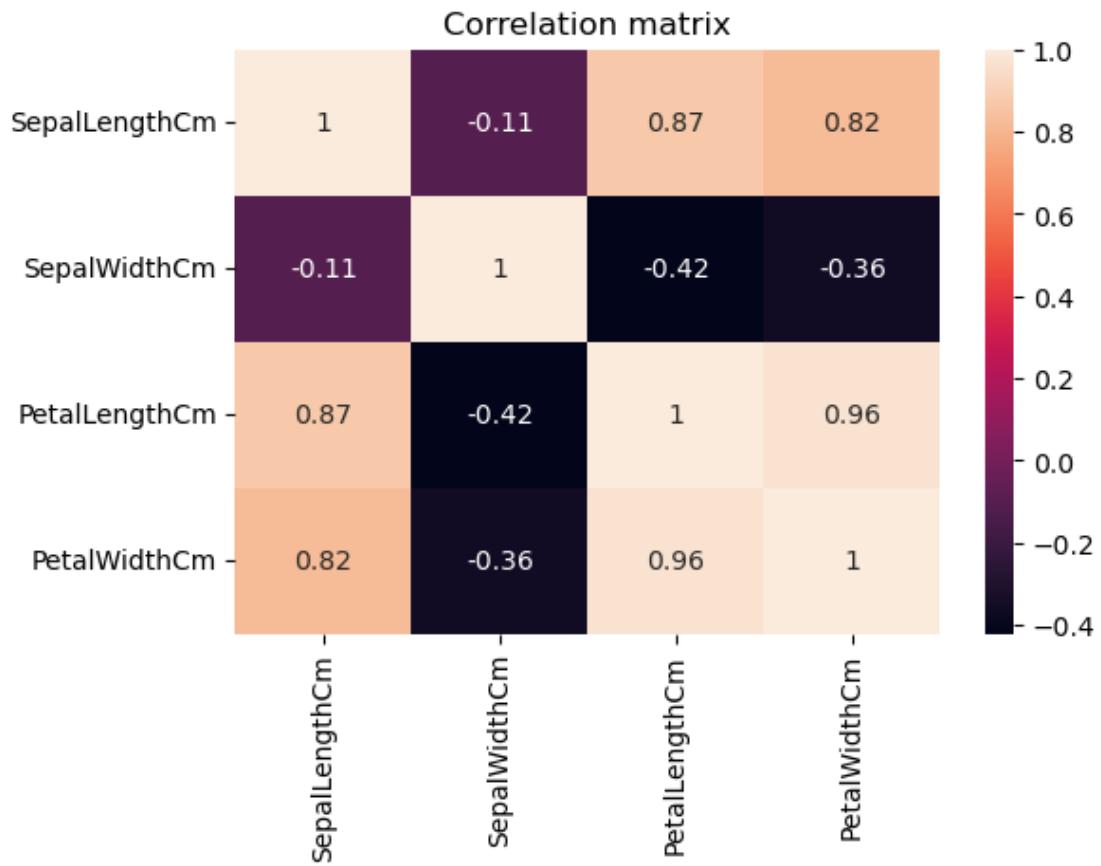


Not much outliers to be concerned about

```
[10]: df['Species'].value_counts().plot(kind='bar')
axes[i].set_title("Species")
axes[i].set_xlabel('')
axes[i].set_ylabel('Count')
plt.savefig("bar_plot.png", dpi=300, bbox_inches="tight")
```



```
[11]: plt.figure(figsize=(6,4))
sns.heatmap(df[num_cols].corr(), annot=True)
plt.title("Correlation matrix")
plt.savefig("correlational_matrix.png", dpi=300, bbox_inches="tight")
plt.show()
```



[]:

2 Loan Amount Prediction

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.impute import SimpleImputer
```

```

from sklearn.feature_selection import SelectKBest, chi2
from sklearn.feature_selection import f_classif
from sklearn.model_selection import train_test_split

```

[]: #Loan dataset

```

df = pd.read_csv("/content/loan_data.csv")
print(df)

```

	person_age	person_gender	person_education	person_income	\
0	22.0	female	Master	71948.0	
1	21.0	female	High School	12282.0	
2	25.0	female	High School	12438.0	
3	23.0	female	Bachelor	79753.0	
4	24.0	male	Master	66135.0	
...
44995	27.0	male	Associate	47971.0	
44996	37.0	female	Associate	65800.0	
44997	33.0	male	Associate	56942.0	
44998	29.0	male	Bachelor	33164.0	
44999	24.0	male	High School	51609.0	

	person_emp_exp	person_home_ownership	loan_amnt	loan_intent	\
0	0	RENT	35000.0	PERSONAL	
1	0	OWN	1000.0	EDUCATION	
2	3	MORTGAGE	5500.0	MEDICAL	
3	0	RENT	35000.0	MEDICAL	
4	1	RENT	35000.0	MEDICAL	
...
44995	6	RENT	15000.0	MEDICAL	
44996	17	RENT	9000.0	HOMEIMPROVEMENT	
44997	7	RENT	2771.0	DEBTCONSOLIDATION	
44998	4	RENT	12000.0	EDUCATION	
44999	1	RENT	6665.0	DEBTCONSOLIDATION	

	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	\
0	16.02	0.49	3.0	
1	11.14	0.08	2.0	
2	12.87	0.44	3.0	
3	15.23	0.44	2.0	
4	14.27	0.53	4.0	
...
44995	15.66	0.31	3.0	
44996	14.07	0.14	11.0	
44997	10.02	0.05	10.0	
44998	13.23	0.36	6.0	
44999	17.05	0.13	3.0	

	credit_score	previous_loan_defaults_on_file	loan_status
--	--------------	--------------------------------	-------------

```

0          561           No      1
1          504           Yes     0
2          635           No      1
3          675           No      1
4          586           No      1
...
...          ...
44995       645           No      1
44996       621           No      1
44997       668           No      1
44998       604           No      1
44999       628           No      1

```

[45000 rows x 14 columns]

```
[ ]: print("Shape")
df.shape
```

Shape

```
[ ]: (45000, 14)
```

```
[ ]: print("Describe")
df.describe()
```

Describe

```

[ ]:      person_age  person_income  person_emp_exp  loan_amnt \
count  45000.000000  4.500000e+04  45000.000000  45000.000000
mean    27.764178  8.031905e+04  5.410333  9583.157556
std     6.045108  8.042250e+04  6.063532  6314.886691
min    20.000000  8.000000e+03  0.000000  500.000000
25%   24.000000  4.720400e+04  1.000000  5000.000000
50%   26.000000  6.704800e+04  4.000000  8000.000000
75%   30.000000  9.578925e+04  8.000000  12237.250000
max   144.000000  7.200766e+06  125.000000  35000.000000

      loan_int_rate  loan_percent_income  cb_person_cred_hist_length \
count  45000.000000            45000.000000            45000.000000
mean    11.006606              0.139725            5.867489
std     2.978808              0.087212            3.879702
min    5.420000              0.000000            2.000000
25%   8.590000              0.070000            3.000000
50%  11.010000              0.120000            4.000000
75%  12.990000              0.190000            8.000000
max   20.000000              0.660000            30.000000

      credit_score  loan_status
count  45000.000000  45000.000000
mean    632.608756  0.222222

```

```
std      50.435865    0.415744
min     390.000000    0.000000
25%    601.000000    0.000000
50%    640.000000    0.000000
75%    670.000000    0.000000
max     850.000000    1.000000
```

Outliers found! age = 144, person_income = 72 lakh

```
[ ]: print("columns")
df.columns
```

columns

```
[ ]: Index(['person_age', 'person_gender', 'person_education', 'person_income',
       'person_emp_exp', 'person_home_ownership', 'loan_amnt', 'loan_intent',
       'loan_int_rate', 'loan_percent_income', 'cb_person_cred_hist_length',
       'credit_score', 'previous_loan_defaults_on_file', 'loan_status'],
      dtype='object')
```

```
[ ]: df.isnull().sum()
(df.isnull().mean() * 100).sort_values(ascending=False)
```

```
[ ]: person_age          0.0
person_gender        0.0
person_education      0.0
person_income         0.0
person_emp_exp        0.0
person_home_ownership 0.0
loan_amnt            0.0
loan_intent           0.0
loan_int_rate         0.0
loan_percent_income   0.0
cb_person_cred_hist_length 0.0
credit_score          0.0
previous_loan_defaults_on_file 0.0
loan_status           0.0
dtype: float64
```

No Null values.

```
[ ]: #including previous_loan_defaults_on_file
df['previous_loan_defaults_on_file'] = df['previous_loan_defaults_on_file'].
    ↪map({
        'Yes': 1,
        'No': 0
    })

#column selection
```

```

num_cols = df.select_dtypes(include=['int64','float64']).columns
cat_cols = df.select_dtypes(include=['object']).columns

num_cols, cat_cols

```

```

[ ]: (Index(['person_age', 'person_income', 'person_emp_exp', 'loan_amnt',
       'loan_int_rate', 'loan_percent_income', 'cb_person_cred_hist_length',
       'credit_score', 'previous_loan_defaults_on_file', 'loan_status'],
       dtype='object'),
      Index(['person_gender', 'person_education', 'person_home_ownership',
             'loan_intent'],
       dtype='object'))

```

```

[ ]: df[num_cols].describe()
#df[num_cols].median()

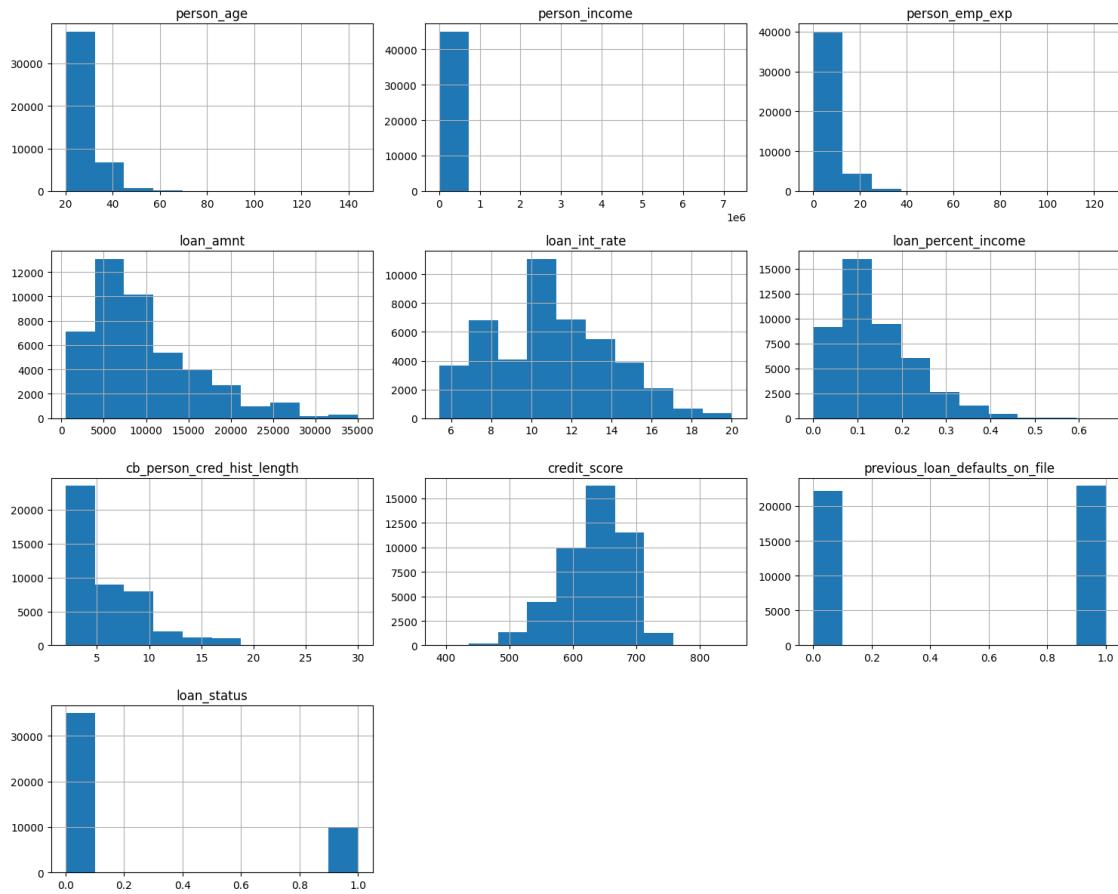
```

	person_age	person_income	person_emp_exp	loan_amnt	\
count	45000.000000	4.500000e+04	45000.000000	45000.000000	
mean	27.764178	8.031905e+04	5.410333	9583.157556	
std	6.045108	8.042250e+04	6.063532	6314.886691	
min	20.000000	8.000000e+03	0.000000	500.000000	
25%	24.000000	4.720400e+04	1.000000	5000.000000	
50%	26.000000	6.704800e+04	4.000000	8000.000000	
75%	30.000000	9.578925e+04	8.000000	12237.250000	
max	144.000000	7.200766e+06	125.000000	35000.000000	

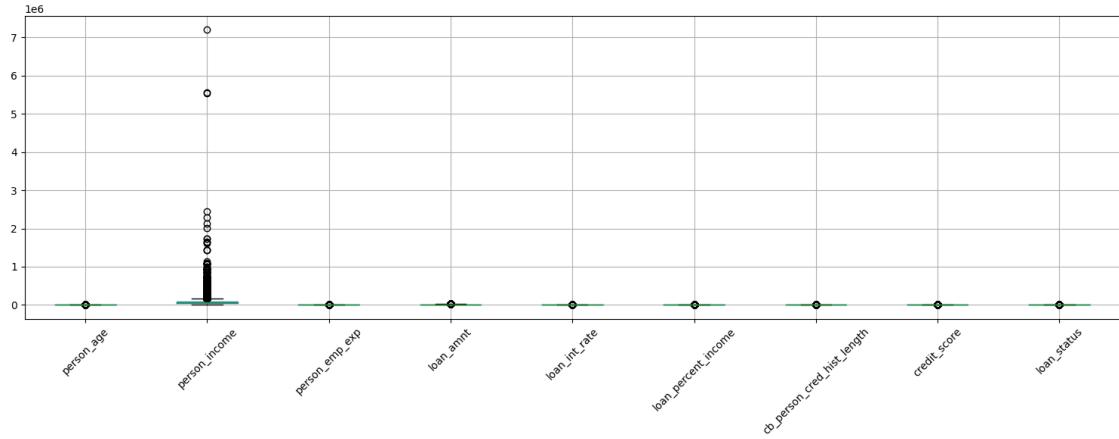
	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	\
count	45000.000000	45000.000000	45000.000000	
mean	11.006606	0.139725	5.867489	
std	2.978808	0.087212	3.879702	
min	5.420000	0.000000	2.000000	
25%	8.590000	0.070000	3.000000	
50%	11.010000	0.120000	4.000000	
75%	12.990000	0.190000	8.000000	
max	20.000000	0.660000	30.000000	

	credit_score	loan_status
count	45000.000000	45000.000000
mean	632.608756	0.222222
std	50.435865	0.415744
min	390.000000	0.000000
25%	601.000000	0.000000
50%	640.000000	0.000000
75%	670.000000	0.000000
max	850.000000	1.000000

```
[ ]: df[num_cols].hist(figsize=(15,12))
plt.tight_layout()
plt.show()
```



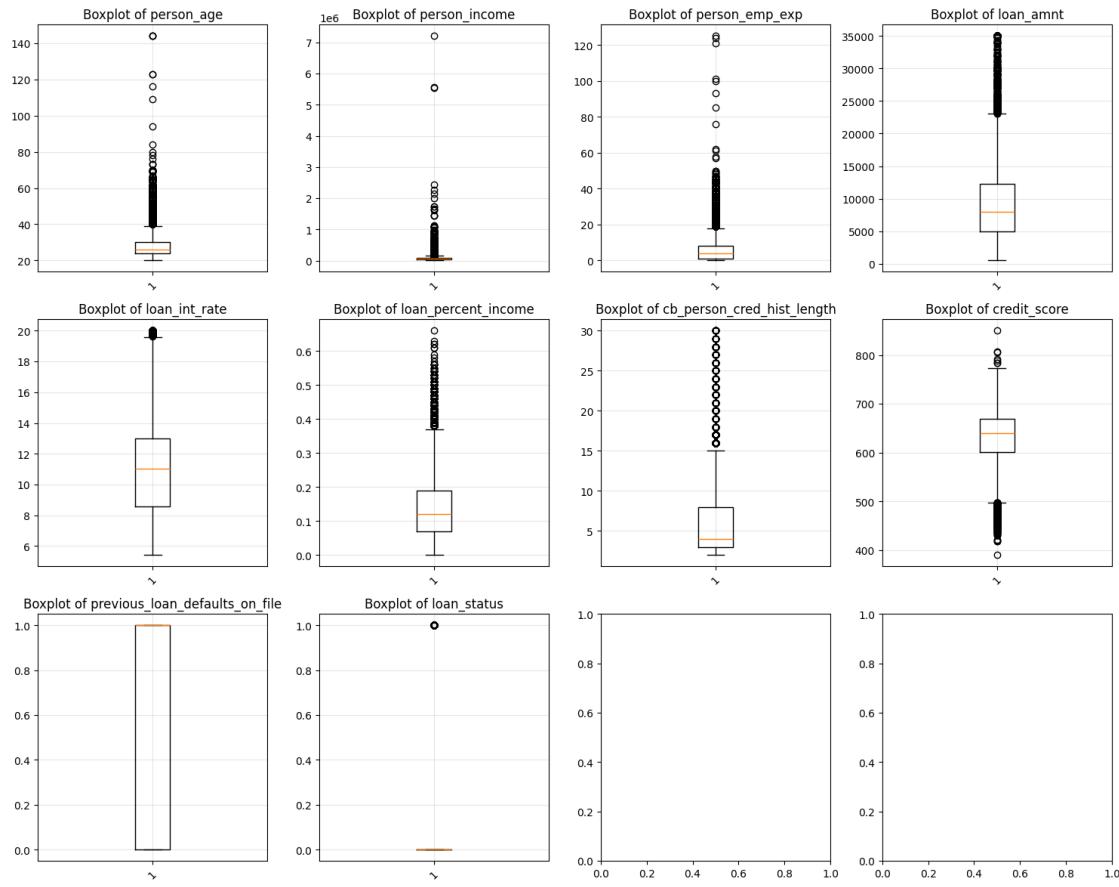
```
[ ]: #box plot
plt.figure(figsize=(15, 6))
df[num_cols].boxplot(rot=45)
plt.tight_layout()
plt.show()
```



```
[ ]: fig, axes = plt.subplots(3, 4, figsize=(15, 12))
axes = axes.flatten()

for i, col in enumerate(num_cols):
    axes[i].boxplot(df[col])
    axes[i].set_title(f"Boxplot of {col}")
    axes[i].tick_params(axis='x', rotation=45)
    axes[i].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



```
[ ]: outlier_summary = {}

for col in num_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    lower_outliers = (df[col] < lower_bound).sum()
    upper_outliers = (df[col] > upper_bound).sum()

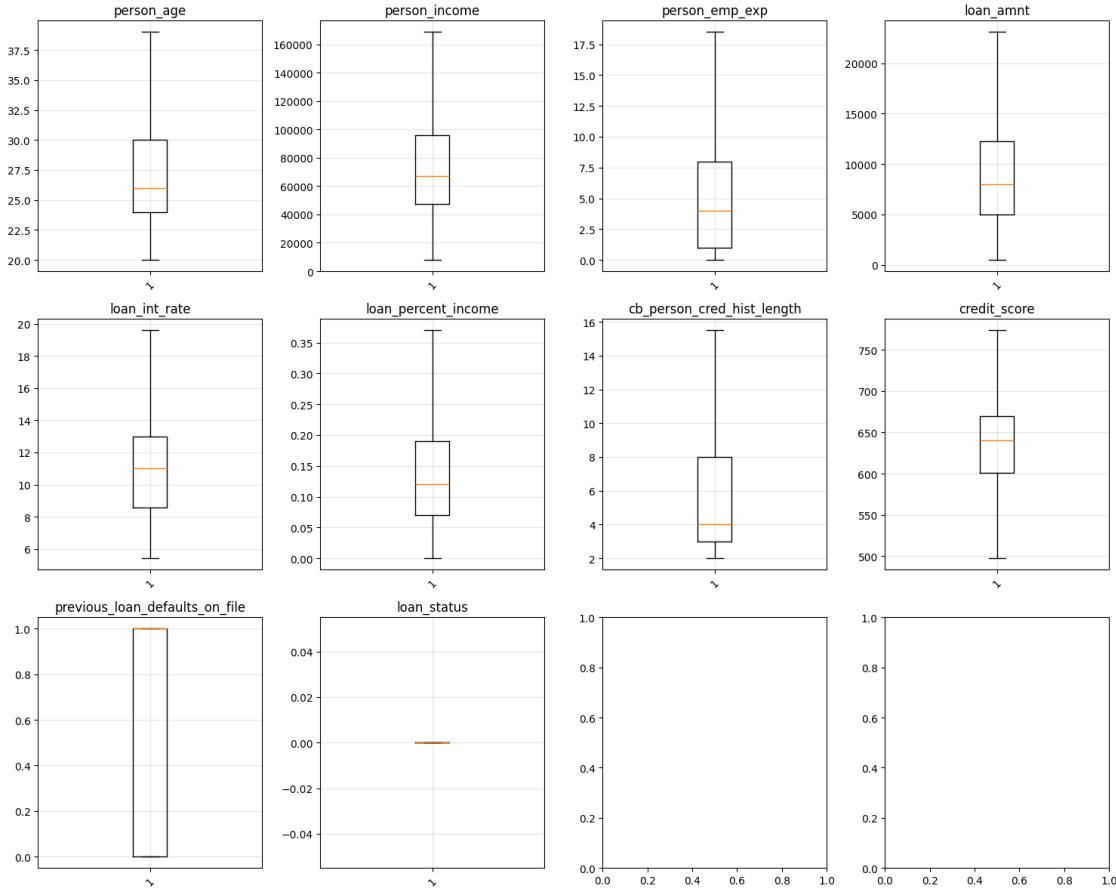
    outlier_summary[col] = {
        "Lower Outliers": lower_outliers,
        "Upper Outliers": upper_outliers,
        "Total Outliers": lower_outliers + upper_outliers
    }
```

```
outlier_df = pd.DataFrame(outlier_summary).T  
outlier_df
```

```
[ ]:      Lower Outliers  Upper Outliers  Total Outliers  
person_age                      0          2188          2188  
person_income                     0          2218          2218  
person_emp_exp                   0          1724          1724  
loan_amnt                        0          2348          2348  
loan_int_rate                     0           124           124  
loan_percent_income                0           744           744  
cb_person_cred_hist_length        0          1366          1366  
credit_score                      460            7           467  
previous_loan_defaults_on_file    0           0             0  
loan_status                       0          10000          10000
```

```
[ ]: #winorization  
df_capped = df.copy()  
  
for col in num_cols:  
    Q1 = df[col].quantile(0.25)  
    Q3 = df[col].quantile(0.75)  
    IQR = Q3 - Q1  
  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
  
    df_capped[col] = df_capped[col].clip(lower_bound, upper_bound)
```

```
[ ]: fig, axes = plt.subplots(3, 4, figsize=(15, 12))  
axes = axes.flatten()  
  
for i, col in enumerate(num_cols):  
    axes[i].boxplot(df_capped[col])  
    axes[i].set_title(col)  
    axes[i].tick_params(axis='x', rotation=45)  
    axes[i].grid(True, alpha=0.3)  
  
plt.tight_layout()  
plt.show()
```



```
[ ]: ncols = 5
nrows = (len(cat_cols) + ncols - 1) // ncols

fig, axes = plt.subplots(
    nrows=nrows,
    ncols=ncols,
    figsize=(20, 4 * nrows)
)

# Make axes always a flat array
if isinstance(axes, np.ndarray):
    axes = axes.flatten()
else:
    axes = [axes]

# Plot each categorical column
for i, col in enumerate(cat_cols):
    df[col].value_counts().plot(kind='bar', ax=axes[i])
    axes[i].set_title(col)
```

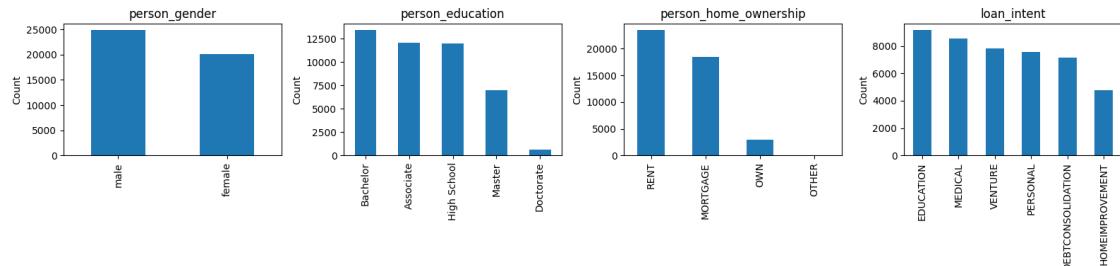
```

        axes[i].set_xlabel('')
        axes[i].set_ylabel('Count')

# Remove any unused subplots
for j in range(len(cat_cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

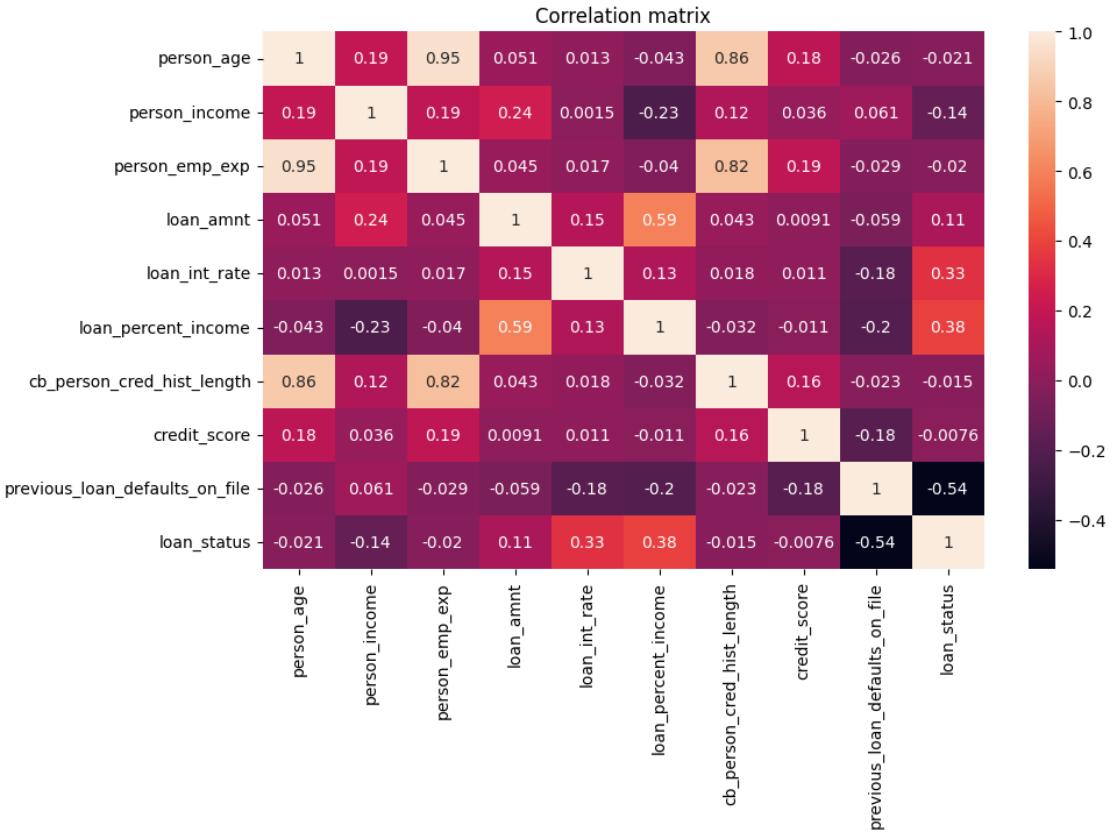
```



```

[ ]: plt.figure(figsize=(10,6))
sns.heatmap(df[num_cols].corr(), annot=True)
plt.title("Correlation matrix")
plt.show()

```



```
[ ]: y = df['loan_status']
X = df.drop(columns=['loan_status'])

num_cols = [col for col in num_cols if col != 'loan_status']
```

```
[ ]: # ANOVA

selector = SelectKBest(score_func=f_classif, k=5)
X_anova_selected = selector.fit_transform(X[num_cols], y)

anova_scores = pd.DataFrame({
    'Feature': num_cols,
    'ANOVA F-Score': selector.scores_
}).sort_values(by='ANOVA F-Score', ascending=False)

anova_scores
```

```
[ ]:          Feature  ANOVA F-Score
8  previous_loan_defaults_on_file  18824.727466
5      loan_percent_income     7824.794030
```

```

4          loan_int_rate    5574.454260
1          person_income    845.525887
3          loan_amnt        528.213632
0          person_age       20.763596
2          person_emp_exp   18.883771
6  cb_person_cred_hist_length    9.926174
7          credit_score     2.631606

```

```

[ ]: selected_features = [
    'previous_loan_defaults_on_file',
    'loan_percent_income',
    'loan_int_rate',
    'person_income',
    'loan_amnt'
]

```

```

[ ]: from sklearn.model_selection import train_test_split

X_train, X_temp, y_train, y_temp = train_test_split(
    X,
    y,
    test_size=0.30,           # 30% + temp
    random_state=42,
    stratify=y                # IMPORTANT for imbalanced classes
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp,
    y_temp,
    test_size=0.50,           # split 30% into 15% + 15%
    random_state=42,
    stratify=y_temp
)

print("Training set:", X_train.shape)
print("Validation set:", X_val.shape)
print("Test set:", X_test.shape)

```

Training set: (31500, 13)
Validation set: (6750, 13)
Test set: (6750, 13)

```

[ ]: #to verify class imbalance
def class_distribution(y, name):
    print(f"{name} class distribution:")
    print(y.value_counts(normalize=True))
    print()

```

```
class_distribution(y_train, "Train")
class_distribution(y_val, "Validation")
class_distribution(y_test, "Test")
```

```
Train class distribution:
loan_status
0    0.777778
1    0.222222
Name: proportion, dtype: float64
```

```
Validation class distribution:
loan_status
0    0.777778
1    0.222222
Name: proportion, dtype: float64
```

```
Test class distribution:
loan_status
0    0.777778
1    0.222222
Name: proportion, dtype: float64
```

Similar score, so there is no class imbalance

3 Predicting Diabetes

```
[ ]: # DiabetesData
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder

from sklearn.feature_selection import SelectKBest, chi2
from sklearn.feature_selection import f_classif
from sklearn.model_selection import train_test_split
```

```
[ ]: df=pd.read_csv('diabetes.csv')
df
```

```
[ ]:      gender   age  hypertension  heart_disease smoking_history    bmi \
0     Female  80.0           0            1        never  25.19
```

```

1      Female  54.0          0          0      No Info  27.32
2       Male   28.0          0          0      never    27.32
3      Female  36.0          0          0     current  23.45
4       Male  76.0          1          1     current  20.14
...
...     ...    ...
99995  Female  80.0          0          0      No Info  27.32
99996  Female   2.0          0          0      No Info  17.37
99997  Male   66.0          0          0     former   27.83
99998  Female  24.0          0          0      never   35.42
99999  Female  57.0          0          0     current  22.43

      HbA1c_level  blood_glucose_level  diabetes
0            6.6              140          0
1            6.6              80          0
2            5.7              158          0
3            5.0              155          0
4            4.8              155          0
...
...     ...
99995        6.2              90          0
99996        6.5             100          0
99997        5.7              155          0
99998        4.0              100          0
99999        6.6              90          0

```

[100000 rows x 9 columns]

```
[ ]: print(df.columns)
```

```
Index(['gender', 'age', 'hypertension', 'heart_disease', 'smoking_history',
       'bmi', 'HbA1c_level', 'blood_glucose_level', 'diabetes'],
      dtype='object')
```

```
[ ]: print(df.describe())
```

	age	hypertension	heart_disease	bmi	\
count	100000.000000	100000.000000	100000.000000	100000.000000	
mean	41.885856	0.07485	0.039420	27.320767	
std	22.516840	0.26315	0.194593	6.636783	
min	0.080000	0.00000	0.000000	10.010000	
25%	24.000000	0.00000	0.000000	23.630000	
50%	43.000000	0.00000	0.000000	27.320000	
75%	60.000000	0.00000	0.000000	29.580000	
max	80.000000	1.00000	1.000000	95.690000	

	HbA1c_level	blood_glucose_level	diabetes
count	100000.000000	100000.000000	100000.000000
mean	5.527507	138.058060	0.085000
std	1.070672	40.708136	0.278883

```

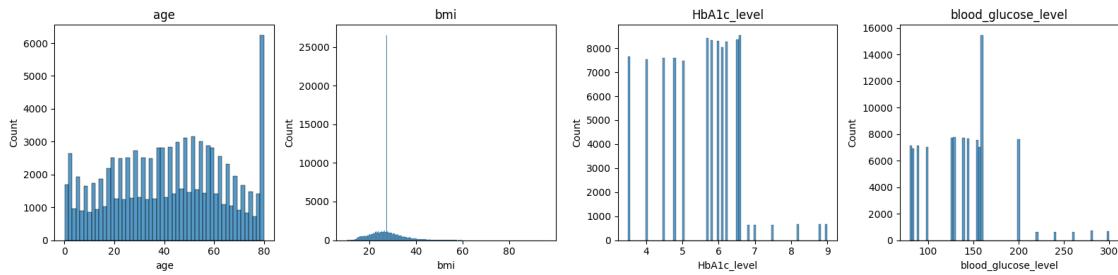
min           3.500000          80.000000        0.000000
25%          4.800000         100.000000        0.000000
50%          5.800000         140.000000        0.000000
75%          6.200000         159.000000        0.000000
max           9.000000         300.000000        1.000000

```

```

[ ]: import math
num_cols = ['age', 'bmi', 'HbA1c_level', 'blood_glucose_level']
ncols=5
nrows=math.ceil(len(num_cols)/ncols)
fig,axes=plt.subplots(nrows=nrows,ncols=ncols,figsize=(20,4*nrows))
axes=axes.flatten()
for i,col in enumerate(num_cols):
    sns.histplot(df[col],ax=axes[i])
    axes[i].set_title(col)
for j in range(i+1,len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()

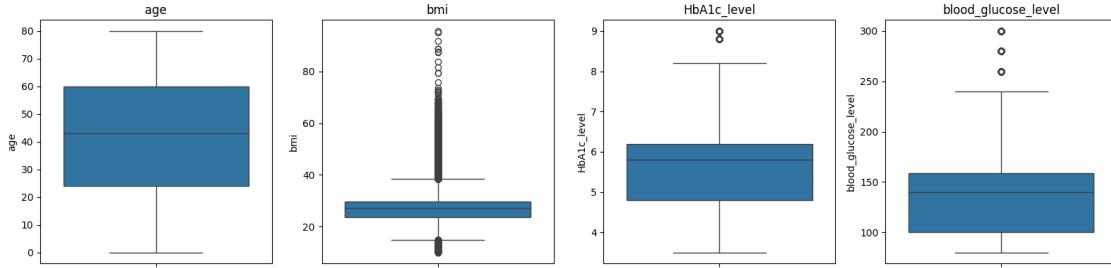
```



```

[ ]: ncols=5
nrows=math.ceil(len(num_cols)/ncols)
fig,axes=plt.subplots(nrows=nrows,ncols=ncols,figsize=(20,4*nrows))
axes=axes.flatten()
for i,col in enumerate(num_cols):
    sns.boxplot(df[col],ax=axes[i])
    axes[i].set_title(col)
for j in range(i+1,len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()

```

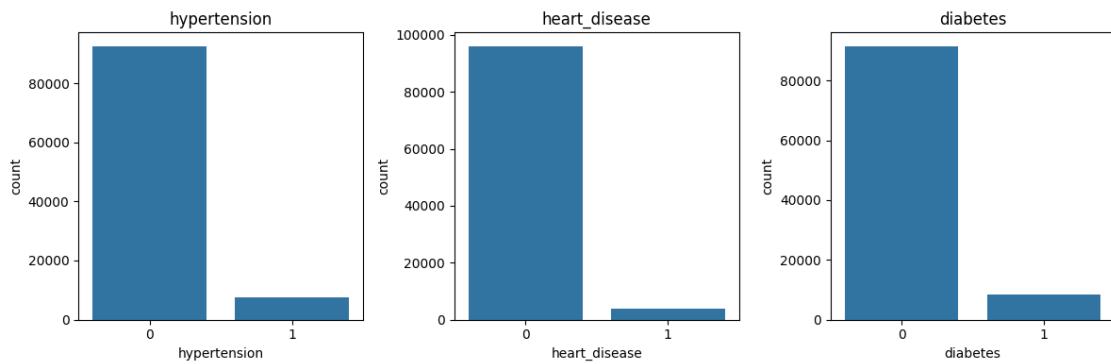


```
[ ]: cat_cols = ['hypertension', 'heart_disease', 'diabetes']

fig, axes = plt.subplots(nrows=1, ncols=5, figsize=(20,4))
axes = axes.flatten()
for i, col in enumerate(cat_cols):
    sns.countplot(x=df[col], ax=axes[i])
    axes[i].set_title(col)

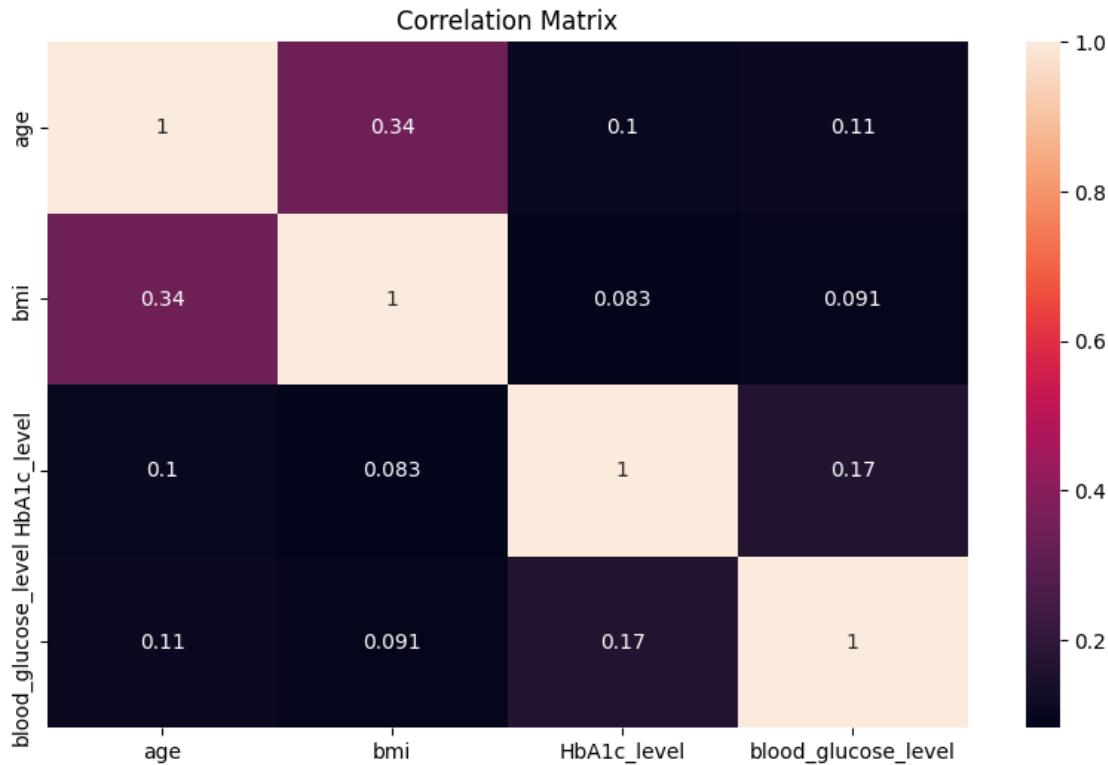
for j in range(i+1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



Heavy imbalanced

```
[ ]: plt.figure(figsize=(10, 6))
sns.heatmap(df[num_cols].corr(), annot=True)
plt.title("Correlation Matrix")
plt.show()
```



4 Classification of Email Spam

```
[19]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import math
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
[20]: df = pd.read_csv('spam_ham_dataset.csv')
df
```

```
[20]:      Unnamed: 0  label          text \
0           605   ham  Subject: enron methanol ; meter # : 988291\r\n...
1          2349   ham  Subject: hpl nom for january 9 , 2001\r\n( see...
2          3624   ham  Subject: neon retreat\r\nho ho ho , we ' re ar...
3          4685  spam  Subject: photoshop , windows , office . cheap ...
4          2030   ham  Subject: re : indian springs\r\nthis deal is t...
...         ...    ...
5166        1518   ham  Subject: put the 10 on the ft\r\nthe transport...
5167        404   ham  Subject: 3 / 4 / 2000 and following noms\r\nnhp...
```

```
5168      2933  ham  Subject: calpine daily gas nomination\r\n>\r\n...
5169      1409  ham  Subject: industrial worksheets for august 2000...
5170      4807  spam  Subject: important online banking alert\r\nndea...
```

```
label_num  
0          0  
1          0  
2          0  
3          1  
4          0  
...        ...  
5166      0  
5167      0  
5168      0  
5169      0  
5170      1
```

[5171 rows x 4 columns]

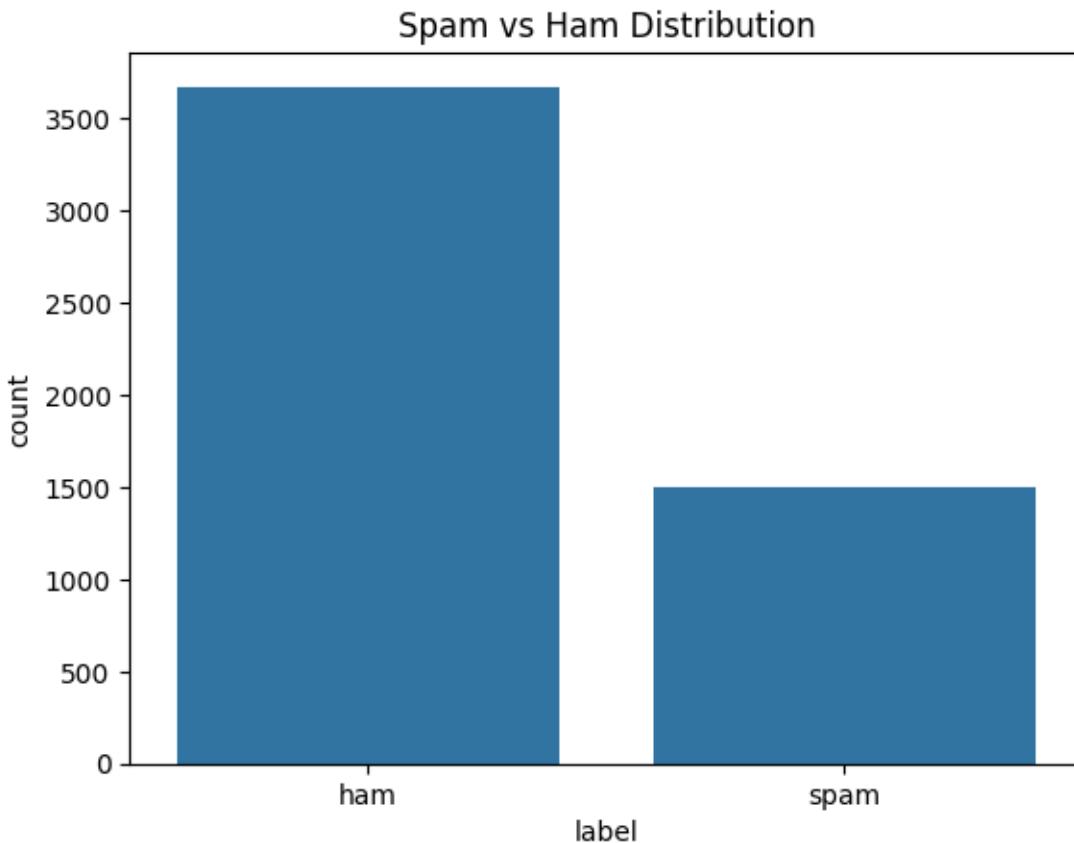
[21]: df.info()

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5171 entries, 0 to 5170  
Data columns (total 4 columns):  
 #   Column      Non-Null Count  Dtype     
---  --    
 0   Unnamed: 0    5171 non-null   int64    
 1   label        5171 non-null   object    
 2   text         5171 non-null   object    
 3   label_num    5171 non-null   int64    
dtypes: int64(2), object(2)  
memory usage: 161.7+ KB
```

[22]: df.describe()

```
Unnamed: 0  label_num  
count  5171.000000  5171.000000  
mean   2585.000000  0.289886  
std    1492.883452  0.453753  
min    0.000000  0.000000  
25%   1292.500000  0.000000  
50%   2585.000000  0.000000  
75%   3877.500000  1.000000  
max   5170.000000  1.000000
```

[23]: sns.countplot(x='label', data=df)
plt.title('Spam vs Ham Distribution')
plt.show()



```
[24]: import re
def clean_text(text):
    text = text.lower()
    text = re.sub(r"http\S+", "", text)
    text = re.sub(r"[^a-z\s]", "", text)
    return text

df['clean_text'] = df['text'].apply(clean_text)
df['word_count'] = df['clean_text'].apply(lambda x: len(x.split()))
df['char_count'] = df['clean_text'].apply(len)

df.head()
```

```
[24]:      Unnamed: 0  label
0          605    ham  Subject: enron methanol ; meter # : 988291\r\n...
1         2349    ham  Subject: hpl nom for january 9 , 2001\r\n( see...
2         3624    ham  Subject: neon retreat\r\nho ho ho , we ' re ar...
3         4685   spam  Subject: photoshop , windows , office . cheap ...
4         2030    ham  Subject: re : indian springs\r\nthis deal is t...
```

```

label_num          clean_text word_count \
0      0 subject enron methanol meter  \r\nthis is a ...      49
1      0 subject hpl nom for january  \r\n see attach...      12
2      0 subject neon retreat\r\nho ho ho we re aroun...     460
3      1 subject photoshop windows office cheap mai...      44
4      0 subject re indian springs\r\nthis deal is to ...      64

char_count
0      302
1      80
2     2428
3     409
4     329

```

```
[25]: from collections import Counter

all_words = " ".join(df['clean_text']).split()
common_words = Counter(all_words).most_common(20)
common_words
```

```
[25]: [('the', 25613),
('to', 20332),
('ect', 13900),
('and', 12815),
('for', 10505),
('of', 10167),
('a', 9813),
('you', 8159),
('subject', 8060),
('in', 7699),
('on', 7308),
('hou', 7289),
('is', 7162),
('this', 7161),
('enron', 6555),
('i', 6379),
('be', 5060),
('that', 4767),
('we', 4339),
('from', 4191)]
```

```
[26]: spam_words = " ".join(df[df['label']=="spam"]['clean_text']).split()
ham_words = " ".join(df[df['label']=="ham"]['clean_text']).split()
print("Spam words most common \n\n", Counter(spam_words).most_common(15))
print("\nNot Spam common words \n", Counter(ham_words).most_common(15))
```

Spam words most common

```
[('the', 7254), ('to', 5160), ('and', 4903), ('of', 4490), ('a', 3787), ('in', 3129), ('you', 2794), ('for', 2523), ('this', 2283), ('is', 2256), ('your', 1946), ('subject', 1657), ('with', 1470), ('that', 1348), ('s', 1316)]
```

Not Spam common words

```
[('the', 18359), ('to', 15172), ('ect', 13897), ('for', 7982), ('and', 7912), ('hou', 7281), ('enron', 6555), ('subject', 6403), ('on', 6049), ('a', 6026), ('of', 5677), ('you', 5365), ('i', 5241), ('is', 4906), ('this', 4878)]
```

```
[27]: df.columns
```

```
[27]: Index(['Unnamed: 0', 'label', 'text', 'label_num', 'clean_text', 'word_count', 'char_count'],  
          dtype='object')
```

```
[28]: df.drop(columns=['Unnamed: 0'], inplace=True)
```

```
[29]: vectorizer = TfidfVectorizer(  
    stop_words='english',  
    max_features=5000  
)
```

```
X_text = vectorizer.fit_transform(df['clean_text'])  
y = df['label_num']
```

```
[32]: from sklearn.feature_selection import SelectKBest, chi2
```

```
chi2_scores, p_values = chi2(X_text, y)
```

```
[33]: chi2_df = pd.DataFrame({  
    'word': vectorizer.get_feature_names_out(),  
    'chi2_score': chi2_scores,  
    'p_value': p_values  
)  
  
chi2_df = chi2_df.sort_values(by='chi2_score', ascending=False)  
  
chi2_df.head(15)
```

```
[33]:      word  chi2_score      p_value  
1447      ect  120.299194  5.440422e-28  
2203      http  103.009438  3.335553e-24  
1539     enron   90.682972  1.686383e-21  
2187      hpl   77.925914  1.069787e-18  
4967      xls   70.555826  4.474242e-17  
2179      hou   65.285626  6.479222e-16  
4956      www   60.633799  6.874349e-15  
1163     deal   54.796253  1.336954e-13
```

```

2888    meter    52.872480  3.559216e-13
3174    online   51.076549  8.883323e-13
811     click    47.859892  4.577915e-12
2857    meds    44.986378  1.984098e-11
1911    gas     44.786746  2.197048e-11
320     attached 42.445020  7.269727e-11
4790    viagra   41.864447  9.782529e-11

```

5 Handwritten Character Recognition / MNIST

```
[19]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import math
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
[20]: df = pd.read_csv('spam_ham_dataset.csv')
df
```

```
[20]:      Unnamed: 0 label
0          605  ham  Subject: enron methanol ; meter # : 988291\r\n...
1          2349  ham  Subject: hpl nom for january 9 , 2001\r\n( see...
2          3624  ham  Subject: neon retreat\r\nnho ho ho , we ' re ar...
3          4685  spam  Subject: photoshop , windows , office . cheap ...
4          2030  ham  Subject: re : indian springs\r\nthis deal is t...
...
5166        1518  ham  Subject: put the 10 on the ft\r\nthe transport...
5167        404   ham  Subject: 3 / 4 / 2000 and following noms\r\nnhp...
5168        2933  ham  Subject: calpine daily gas nomination\r\n>\r\n...
5169        1409  ham  Subject: industrial worksheets for august 2000...
5170        4807  spam  Subject: important online banking alert\r\ndea...

      label_num
0            0
1            0
2            0
3            1
4            0
...
5166        0
5167        0
5168        0
5169        0
5170        1
```

```
[5171 rows x 4 columns]
```

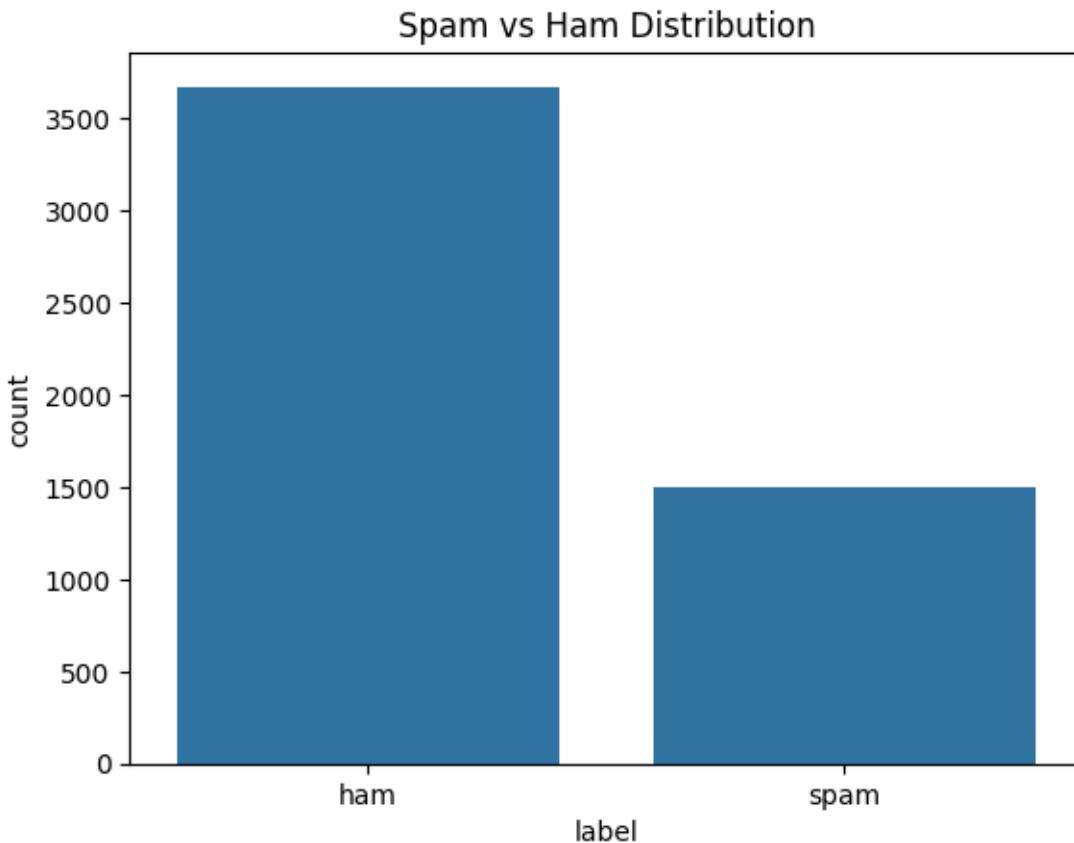
```
[21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5171 entries, 0 to 5170
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Unnamed: 0    5171 non-null   int64  
 1   label        5171 non-null   object  
 2   text         5171 non-null   object  
 3   label_num    5171 non-null   int64  
dtypes: int64(2), object(2)
memory usage: 161.7+ KB
```

```
[22]: df.describe()
```

```
Unnamed: 0    label_num
count    5171.000000  5171.000000
mean     2585.000000    0.289886
std      1492.883452    0.453753
min      0.000000    0.000000
25%    1292.500000    0.000000
50%    2585.000000    0.000000
75%    3877.500000    1.000000
max    5170.000000    1.000000
```

```
[23]: sns.countplot(x='label', data=df)
plt.title('Spam vs Ham Distribution')
plt.show()
```



```
[24]: import re
def clean_text(text):
    text = text.lower()
    text = re.sub(r"http\S+", "", text)
    text = re.sub(r"[^a-z\s]", "", text)
    return text

df['clean_text'] = df['text'].apply(clean_text)
df['word_count'] = df['clean_text'].apply(lambda x: len(x.split()))
df['char_count'] = df['clean_text'].apply(len)

df.head()
```

```
[24]:      Unnamed: 0  label
0          605    ham  Subject: enron methanol ; meter # : 988291\r\n...
1         2349    ham  Subject: hpl nom for january 9 , 2001\r\n( see...
2         3624    ham  Subject: neon retreat\r\nho ho ho , we ' re ar...
3         4685   spam  Subject: photoshop , windows , office . cheap ...
4         2030    ham  Subject: re : indian springs\r\nthis deal is t...
```

```

label_num          clean_text word_count \
0      0 subject enron methanol meter  \r\nthis is a ...      49
1      0 subject hpl nom for january  \r\n see attach...      12
2      0 subject neon retreat\r\nho ho ho we re aroun...     460
3      1 subject photoshop windows office cheap mai...      44
4      0 subject re indian springs\r\nthis deal is to ...      64

char_count
0      302
1      80
2     2428
3     409
4     329

```

```
[25]: from collections import Counter

all_words = " ".join(df['clean_text']).split()
common_words = Counter(all_words).most_common(20)
common_words
```

```
[25]: [('the', 25613),
('to', 20332),
('ect', 13900),
('and', 12815),
('for', 10505),
('of', 10167),
('a', 9813),
('you', 8159),
('subject', 8060),
('in', 7699),
('on', 7308),
('hou', 7289),
('is', 7162),
('this', 7161),
('enron', 6555),
('i', 6379),
('be', 5060),
('that', 4767),
('we', 4339),
('from', 4191)]
```

```
[26]: spam_words = " ".join(df[df['label']=="spam"]['clean_text']).split()
ham_words = " ".join(df[df['label']=="ham"]['clean_text']).split()
print("Spam words most common \n\n", Counter(spam_words).most_common(15))
print("\nNot Spam common words \n", Counter(ham_words).most_common(15))
```

Spam words most common

```
[('the', 7254), ('to', 5160), ('and', 4903), ('of', 4490), ('a', 3787), ('in', 3129), ('you', 2794), ('for', 2523), ('this', 2283), ('is', 2256), ('your', 1946), ('subject', 1657), ('with', 1470), ('that', 1348), ('s', 1316)]
```

Not Spam common words

```
[('the', 18359), ('to', 15172), ('ect', 13897), ('for', 7982), ('and', 7912), ('hou', 7281), ('enron', 6555), ('subject', 6403), ('on', 6049), ('a', 6026), ('of', 5677), ('you', 5365), ('i', 5241), ('is', 4906), ('this', 4878)]
```

```
[27]: df.columns
```

```
[27]: Index(['Unnamed: 0', 'label', 'text', 'label_num', 'clean_text', 'word_count', 'char_count'],  
          dtype='object')
```

```
[28]: df.drop(columns=['Unnamed: 0'], inplace=True)
```

```
[29]: vectorizer = TfidfVectorizer(  
    stop_words='english',  
    max_features=5000  
)
```

```
X_text = vectorizer.fit_transform(df['clean_text'])  
y = df['label_num']
```

```
[32]: from sklearn.feature_selection import SelectKBest, chi2  
  
chi2_scores, p_values = chi2(X_text, y)
```

```
[33]: chi2_df = pd.DataFrame({  
    'word': vectorizer.get_feature_names_out(),  
    'chi2_score': chi2_scores,  
    'p_value': p_values  
)  
  
chi2_df = chi2_df.sort_values(by='chi2_score', ascending=False)  
  
chi2_df.head(15)
```

	word	chi2_score	p_value
1447	ect	120.299194	5.440422e-28
2203	http	103.009438	3.335553e-24
1539	enron	90.682972	1.686383e-21
2187	hpl	77.925914	1.069787e-18
4967	xls	70.555826	4.474242e-17
2179	hou	65.285626	6.479222e-16
4956	www	60.633799	6.874349e-15
1163	deal	54.796253	1.336954e-13

2888	meter	52.872480	3.559216e-13
3174	online	51.076549	8.883323e-13
811	click	47.859892	4.577915e-12
2857	meds	44.986378	1.984098e-11
1911	gas	44.786746	2.197048e-11
320	attached	42.445020	7.269727e-11
4790	viagra	41.864447	9.782529e-11

Results and Discussions

Dataset	Type of ML Task	Feature Selection Technique	Suitable ML Algorithm
Iris Dataset	Classification	SelectKBest	Logistic Regression
Loan Amount Prediction	Regression	Correlation Coefficient	Linear Regression
Predicting Diabetes	Classification	SelectKBest	Ridge Regression
Classification of Email Spam	Classification	Correlation Coefficient	Naïve Bayes
Handwritten Character Recognition (MNIST)	Classification	PCA	CNN

Learning Practices

- Understand the fundamentals of data analysis using Python libraries such as NumPy, Pandas, Matplotlib, and Seaborn.
- Analyze datasets using descriptive statistics and exploratory data analysis techniques.
- Visualize data distributions and relationships using appropriate plots and charts..
- Identify suitable machine learning tasks such as classification and regression for different datasets.