

Sri Sivasubramaniya Nadar College of Engineering, Chennai
(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester VI
Subject Code & Name	UCS2612 – Machine Learning Algorithms Laboratory	
Academic Year	2025–2026 (Even)	Batch 2023–2027
Due Date	27.01.2026	

Experiment 2: Binary Classification using Naïve Bayes and K-Nearest Neighbors

Objective

To implement Naïve Bayes and K-Nearest Neighbors (KNN) classifiers for a binary classification problem, evaluate them using multiple performance metrics, visualize model behavior, and analyze overfitting, underfitting, and bias-variance characteristics.

Dataset

A benchmark binary classification dataset containing numerical features and two class labels is used.

Dataset reference:

- Kaggle: [Spambase Dataset](#)

```
[23]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
import math
from matplotlib import rcParams
from sklearn.model_selection import train_test_split, StratifiedKFold,
    GridSearchCV, RandomizedSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix, roc_curve, auc
rcParams['font.family']='Arial'
rcParams['font.weight']='bold'
rcParams['font.size']=15
rcParams['axes.labelweight']='bold'
rcParams['axes.titleweight']='bold'
rcParams['xtick.labelsize']=15
rcParams['ytick.labelsize']=15
```

```
from sklearn.metrics import (
accuracy_score, precision_score, recall_score, f1_score,
confusion_matrix, classification_report, roc_auc_score,
roc_curve, average_precision_score)
```

```
[18]: #Load the dataset
df=pd.read_csv('spambase_csv_Kaggle.csv')
df
```

```
[18]:    word_freq_make  word_freq_address  word_freq_all  word_freq_3d  \
0           0.00          0.64          0.64          0.0
1           0.21          0.28          0.50          0.0
2           0.06          0.00          0.71          0.0
3           0.00          0.00          0.00          0.0
4           0.00          0.00          0.00          0.0
...
4596         0.31          0.00          0.62          0.0
4597         0.00          0.00          0.00          0.0
4598         0.30          0.00          0.30          0.0
4599         0.96          0.00          0.00          0.0
4600         0.00          0.00          0.65          0.0

    word_freq_our  word_freq_over  word_freq_remove  word_freq_internet  \
0           0.32          0.00          0.00          0.00
1           0.14          0.28          0.21          0.07
2           1.23          0.19          0.19          0.12
3           0.63          0.00          0.31          0.63
4           0.63          0.00          0.31          0.63
...
4596         0.00          0.31          0.00          0.00
4597         0.00          0.00          0.00          0.00
4598         0.00          0.00          0.00          0.00
4599         0.32          0.00          0.00          0.00
4600         0.00          0.00          0.00          0.00

    word_freq_order  word_freq_mail  ...  char_freq_%3B  char_freq_%28  \
0           0.00          0.00  ...          0.000          0.000
1           0.00          0.94  ...          0.000          0.132
2           0.64          0.25  ...          0.010          0.143
3           0.31          0.63  ...          0.000          0.137
4           0.31          0.63  ...          0.000          0.135
...
4596         0.00          0.00  ...          0.000          0.232
4597         0.00          0.00  ...          0.000          0.000
4598         0.00          0.00  ...          0.102          0.718
4599         0.00          0.00  ...          0.000          0.057
```

```

4600      0.00      0.00  ...
               0.000      0.000
               char_freq_%5B  char_freq_%21  char_freq_%24  char_freq_%23 \
0          0.0          0.778      0.000      0.000
1          0.0          0.372      0.180      0.048
2          0.0          0.276      0.184      0.010
3          0.0          0.137      0.000      0.000
4          0.0          0.135      0.000      0.000
...
4596      0.0          0.000      0.000      0.000
4597      0.0          0.353      0.000      0.000
4598      0.0          0.000      0.000      0.000
4599      0.0          0.000      0.000      0.000
4600      0.0          0.125      0.000      0.000

               capital_run_length_average  capital_run_length_longest \
0                  3.756                  61
1                  5.114                 101
2                  9.821                 485
3                  3.537                  40
4                  3.537                  40
...
4596      1.142                  3
4597      1.555                  4
4598      1.404                  6
4599      1.147                  5
4600      1.250                  5

               capital_run_length_total  class
0                  278          1
1                 1028          1
2                 2259          1
3                  191          1
4                  191          1
...
4596                  88          0
4597                  14          0
4598                 118          0
4599                  78          0
4600                  40          0

```

[4601 rows x 58 columns]

[7]: *#Perform Exploratory Data Analysis (EDA)*
df.describe()

```
[7]:      word_freq_make word_freq_address word_freq_all word_freq_3d \
count    4601.000000   4601.000000   4601.000000   4601.000000
mean     0.104553     0.213015     0.280656     0.065425
std      0.305358     1.290575     0.504143     1.395151
min     0.000000     0.000000     0.000000     0.000000
25%    0.000000     0.000000     0.000000     0.000000
50%    0.000000     0.000000     0.000000     0.000000
75%    0.000000     0.000000     0.420000     0.000000
max     4.540000    14.280000    5.100000    42.810000

      word_freq_our word_freq_over word_freq_remove word_freq_internet \
count    4601.000000   4601.000000   4601.000000   4601.000000
mean     0.312223     0.095901     0.114208     0.105295
std      0.672513     0.273824     0.391441     0.401071
min     0.000000     0.000000     0.000000     0.000000
25%    0.000000     0.000000     0.000000     0.000000
50%    0.000000     0.000000     0.000000     0.000000
75%    0.380000     0.000000     0.000000     0.000000
max     10.000000    5.880000    7.270000    11.110000

      word_freq_order word_freq_mail ... char_freq_%3B char_freq_%28 \
count    4601.000000   4601.000000   ... 4601.000000   4601.000000
mean     0.090067     0.239413     ... 0.038575     0.139030
std      0.278616     0.644755     ... 0.243471     0.270355
min     0.000000     0.000000     ... 0.000000     0.000000
25%    0.000000     0.000000     ... 0.000000     0.000000
50%    0.000000     0.000000     ... 0.000000     0.065000
75%    0.000000     0.160000     ... 0.000000     0.188000
max     5.260000    18.180000     ... 4.385000    9.752000

      char_freq_%5B char_freq_%21 char_freq_%24 char_freq_%23 \
count    4601.000000   4601.000000   4601.000000   4601.000000
mean     0.016976     0.269071     0.075811     0.044238
std      0.109394     0.815672     0.245882     0.429342
min     0.000000     0.000000     0.000000     0.000000
25%    0.000000     0.000000     0.000000     0.000000
50%    0.000000     0.000000     0.000000     0.000000
75%    0.000000     0.315000     0.052000     0.000000
max     4.081000    32.478000    6.003000    19.829000

      capital_run_length_average capital_run_length_longest \
count          4601.000000           4601.000000
mean        5.191515             52.172789
std       31.729449            194.891310
min       1.000000             1.000000
25%      1.588000             6.000000
50%      2.276000            15.000000
```

```
75%           3.706000          43.000000
max         1102.500000        9989.000000
```

```
capital_run_length_total      class
count             4601.000000  4601.000000
mean              283.289285   0.394045
std               606.347851   0.488698
min               1.000000    0.000000
25%              35.000000    0.000000
50%              95.000000    0.000000
75%             266.000000   1.000000
max            15841.000000   1.000000
```

[8 rows x 58 columns]

[4]: df.info()

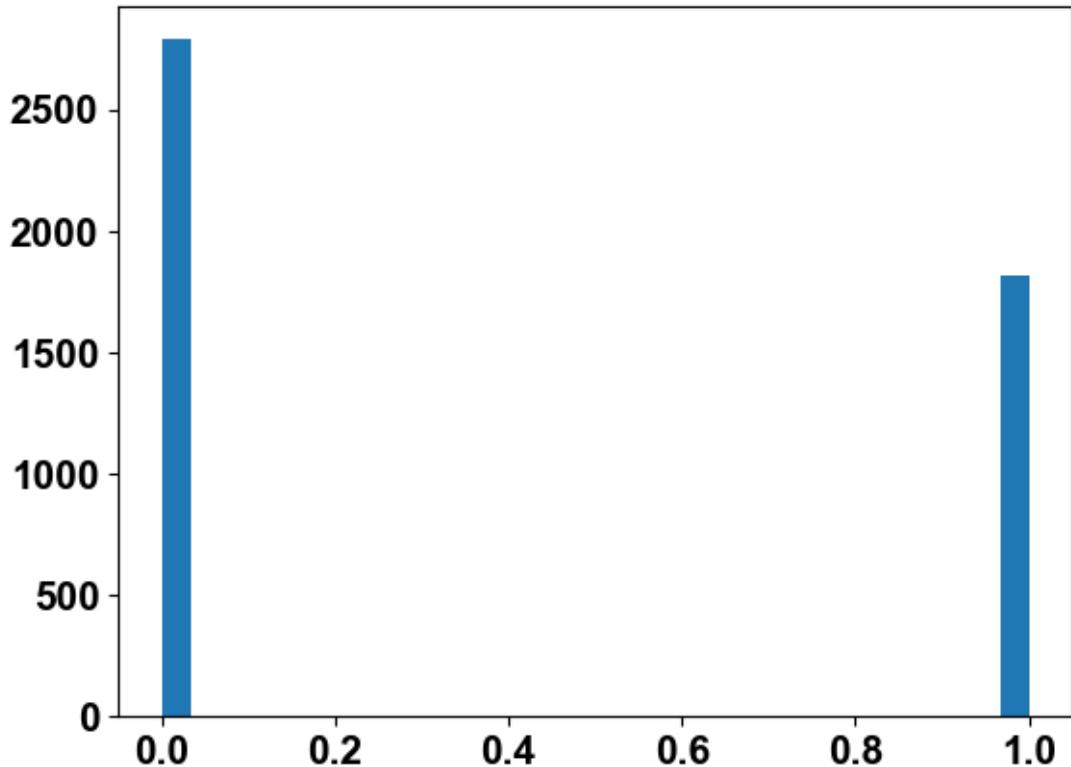
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4601 entries, 0 to 4600
Data columns (total 58 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   word_freq_make      4601 non-null   float64
 1   word_freq_address   4601 non-null   float64
 2   word_freq_all       4601 non-null   float64
 3   word_freq_3d        4601 non-null   float64
 4   word_freq_our       4601 non-null   float64
 5   word_freq_over      4601 non-null   float64
 6   word_freq_remove    4601 non-null   float64
 7   word_freq_internet  4601 non-null   float64
 8   word_freq_order     4601 non-null   float64
 9   word_freq_mail      4601 non-null   float64
 10  word_freq_receive   4601 non-null   float64
 11  word_freq_will      4601 non-null   float64
 12  word_freq_people    4601 non-null   float64
 13  word_freq_report    4601 non-null   float64
 14  word_freq_addresses 4601 non-null   float64
 15  word_freq_free      4601 non-null   float64
 16  word_freq_business  4601 non-null   float64
 17  word_freq_email     4601 non-null   float64
 18  word_freq_you       4601 non-null   float64
 19  word_freq_credit    4601 non-null   float64
 20  word_freq_your      4601 non-null   float64
 21  word_freq_font      4601 non-null   float64
 22  word_freq_000        4601 non-null   float64
 23  word_freq_money     4601 non-null   float64
 24  word_freq_hp         4601 non-null   float64
 25  word_freq_hpl        4601 non-null   float64
```

```

26 word_freq_george           4601 non-null   float64
27 word_freq_650              4601 non-null   float64
28 word_freq_lab              4601 non-null   float64
29 word_freq_labs             4601 non-null   float64
30 word_freq_telnet            4601 non-null   float64
31 word_freq_857              4601 non-null   float64
32 word_freq_data              4601 non-null   float64
33 word_freq_415              4601 non-null   float64
34 word_freq_85                4601 non-null   float64
35 word_freq_technology        4601 non-null   float64
36 word_freq_1999              4601 non-null   float64
37 word_freq_parts             4601 non-null   float64
38 word_freq_pm                4601 non-null   float64
39 word_freq_direct             4601 non-null   float64
40 word_freq_cs                4601 non-null   float64
41 word_freq_meeting            4601 non-null   float64
42 word_freq_original            4601 non-null   float64
43 word_freq_project             4601 non-null   float64
44 word_freq_re                 4601 non-null   float64
45 word_freq_edu                4601 non-null   float64
46 word_freq_table               4601 non-null   float64
47 word_freq_conference          4601 non-null   float64
48 char_freq_%3B                4601 non-null   float64
49 char_freq_%28                4601 non-null   float64
50 char_freq_%5B                4601 non-null   float64
51 char_freq_%21                4601 non-null   float64
52 char_freq_%24                4601 non-null   float64
53 char_freq_%23                4601 non-null   float64
54 capital_run_length_average    4601 non-null   float64
55 capital_run_length_longest     4601 non-null   int64
56 capital_run_length_total       4601 non-null   int64
57 class                         4601 non-null   int64
dtypes: float64(55), int64(3)
memory usage: 2.0 MB

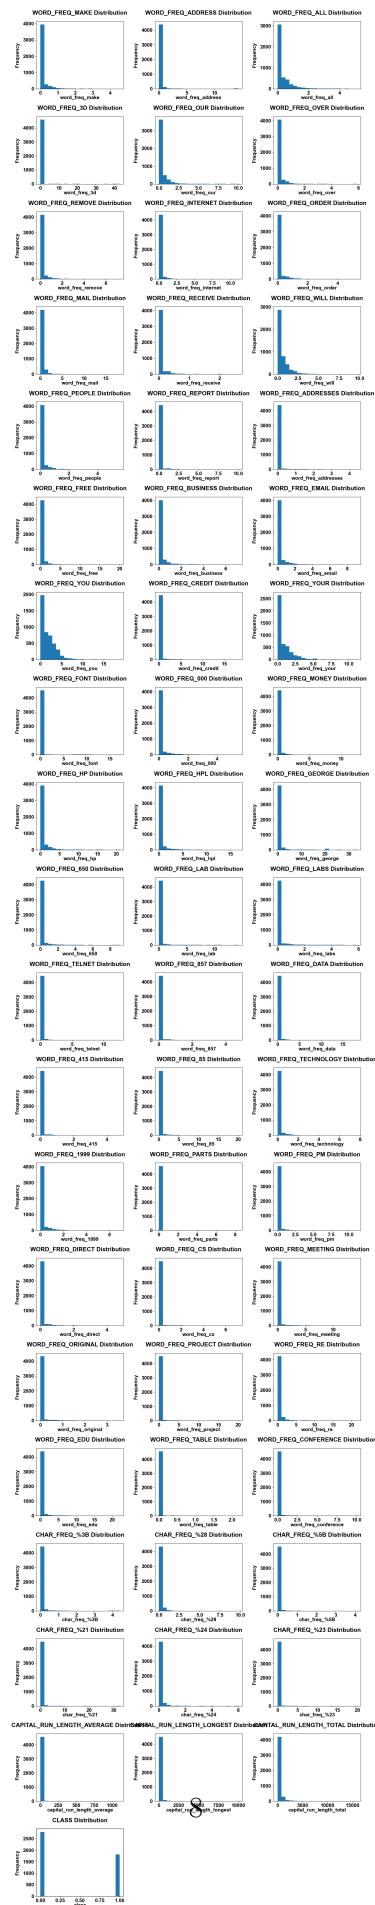
```

```
[24]: # visualize class distribution and feature behavior
plt.hist(df["class"], bins=30)
plt.savefig("class_distribution.png", dpi=300, bbox_inches="tight")
```



Histogram

```
[13]: bxwidth=1
rows=math.ceil(len(df.columns)/3)
fig,axes=plt.subplots(rows,3,figsize=(15,4*rows))
axes=axes.flatten()
columns=df.columns
for i,(ax,col) in enumerate(zip(axes,columns)):
    ax.hist(df[col],bins=20)
    ax.set_title(f'{col.upper()} Distribution',pad=20)
    ax.set_xlabel(col,labelpad=0)
    ax.set_ylabel("Frequency",labelpad=10)
    for spine in ax.spines.values():
        spine.set_linewidth(bxwidth)
    for j in range(i+1,len(axes)):
        fig.delaxes(axes[j])
plt.subplots_adjust(hspace=1.5,wspace=1.3)
plt.tight_layout()
plt.savefig("histogram.png", dpi=300, bbox_inches="tight")
plt.show()
```



Train-Test Split

```
[19]: X=df.iloc[:, :-1]
y=df.iloc[:, -1]
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
                                              ↵2,stratify=y,random_state=42)
print(X_train,X_test,y_train,y_test)
```

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	\
2940	0.05	0.00	0.45	0.0	
1303	0.17	0.26	1.21	0.0	
3468	0.00	0.00	0.00	0.0	
3181	0.00	0.00	0.00	0.0	
794	0.00	0.56	0.00	0.0	
...	
1861	0.00	0.00	4.00	0.0	
2366	0.00	0.00	0.00	0.0	
330	0.00	0.00	1.53	0.0	
536	0.00	0.00	0.00	0.0	
3114	0.00	0.00	0.00	0.0	

	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	\
2940	0.15	0.1	0.00	0.00	
1303	0.43	0.6	0.43	0.26	
3468	0.00	0.0	0.00	0.00	
3181	0.00	0.0	0.00	0.00	
794	0.56	0.0	0.00	0.00	
...	
1861	0.00	0.0	0.00	0.00	
2366	4.16	0.0	0.00	0.00	
330	0.00	0.0	0.00	0.00	
536	0.00	0.0	0.00	0.00	
3114	0.00	0.0	0.00	0.00	

	word_freq_order	word_freq_mail	...	word_freq_conference	\
2940	0.55	0.00	...	0.0	
1303	0.69	0.52	...	0.0	
3468	0.00	0.00	...	0.0	
3181	0.00	0.00	...	0.0	
794	1.01	0.56	...	0.0	
...	
1861	0.00	0.00	...	0.0	
2366	0.00	0.00	...	0.0	
330	0.00	0.00	...	0.0	
536	0.00	0.00	...	0.0	

3114	0.00	0.00	...	0.0
	char_freq_%3B	char_freq_%28	char_freq_%5B	char_freq_%21 \
2940	0.203	0.195	0.05	0.000
1303	0.000	0.108	0.00	0.271
3468	0.000	0.000	0.00	0.153
3181	0.000	0.000	0.00	0.000
794	0.000	0.186	0.00	0.056
...
1861	0.000	0.000	0.00	0.613
2366	0.000	0.689	0.00	0.689
330	0.000	0.000	0.00	1.434
536	0.000	0.407	0.00	0.203
3114	0.000	0.484	0.00	0.484
	char_freq_%24	char_freq_%23	capital_run_length_average \	
2940	0.014	0.000		2.880
1303	0.243	0.013		6.395
3468	0.000	0.000		1.933
3181	0.000	0.000		4.333
794	0.056	0.000		2.153
...
1861	0.000	0.000		1.000
2366	0.000	0.000		1.300
330	0.000	0.000		7.055
536	0.610	0.000		4.133
3114	0.000	0.000		2.500
	capital_run_length_longest	capital_run_length_total		
2940	45		1080	
1303	583		1375	
3468	7		58	
3181	20		26	
794	53		532	
...	
1861	1		14	
2366	4		13	
330	75		127	
536	17		62	
3114	15		65	
[3680 rows x 57 columns]	word_freq_make	word_freq_address	word_freq_all	
word_freq_3d \				
1472	0.00	0.00	0.00	0.0
258	0.00	0.00	0.33	0.0
3564	0.00	0.00	0.00	0.0
65	0.66	0.00	0.66	0.0
4303	0.00	0.00	0.00	0.0

...
1405	0.00	0.00	0.35	0.0
2312	0.00	2.59	1.29	0.0
2804	0.00	0.00	0.00	0.0
2047	0.00	0.00	0.00	0.0
2597	0.00	0.00	0.00	0.0
	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet
1472	0.00	0.00	0.00	0.00
258	0.99	0.99	0.33	0.33
3564	0.00	0.00	0.00	0.00
65	0.00	0.00	0.00	0.00
4303	0.00	0.00	0.00	0.00
...
1405	0.00	0.70	0.35	0.35
2312	1.29	0.00	0.00	0.00
2804	0.00	0.00	0.00	0.00
2047	0.34	0.00	0.00	0.00
2597	0.00	0.00	0.00	0.00
	word_freq_order	word_freq_mail	...	word_freq_conference
1472	0.0	0.00	...	0.00
258	0.0	0.00	...	0.00
3564	0.0	0.00	...	0.00
65	0.0	0.66	...	0.00
4303	0.0	2.77	...	0.00
...
1405	0.0	0.00	...	0.00
2312	0.0	0.00	...	0.00
2804	0.0	0.00	...	0.00
2047	0.0	0.00	...	0.34
2597	0.0	0.00	...	0.00
	char_freq_%3B	char_freq_%28	char_freq_%5B	char_freq_%21
1472	0.144	0.000	0.000	3.907
258	0.000	0.108	0.000	0.000
3564	0.000	0.000	0.000	0.000
65	0.000	0.000	0.000	2.205
4303	0.000	0.000	0.000	0.438
...
1405	0.000	0.061	0.000	0.061
2312	0.000	0.000	0.000	0.000
2804	0.000	0.000	0.000	0.000
2047	0.088	0.132	0.000	0.000
2597	0.142	0.000	0.142	0.000
	char_freq_%24	char_freq_%23	capital_run_length_average	
1472	0.000	0.000	13.928	

258	0.162	0.054	2.195
3564	0.000	0.000	1.214
65	0.000	0.000	3.184
4303	0.000	0.000	1.214
...
1405	0.000	0.122	2.302
2312	0.000	0.000	1.000
2804	0.000	0.000	2.000
2047	0.000	0.000	1.250
2597	0.000	0.000	1.717

	capital_run_length_longest	capital_run_length_total
1472	70	195
258	50	202
3564	4	17
65	34	121
4303	3	17
...
1405	21	99
2312	1	13
2804	4	6
2047	7	85
2597	12	67

```
[921 rows x 57 columns] 2940      0
1303    1
3468    0
3181    0
794     1
...
1861    0
2366    0
330     1
536     1
3114    0
Name: class, Length: 3680, dtype: int64 1472      1
258     1
3564    0
65      1
4303    0
...
1405    1
2312    0
2804    0
2047    0
2597    0
Name: class, Length: 921, dtype: int64
```

Z-score or Standard Scaling

```
[21]: scaler=StandardScaler()
X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.transform(X_test)
print(X_train,X_test,y_train,y_test)
```

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	\
2940	0.05	0.00	0.45	0.0	
1303	0.17	0.26	1.21	0.0	
3468	0.00	0.00	0.00	0.0	
3181	0.00	0.00	0.00	0.0	
794	0.00	0.56	0.00	0.0	
...	
1861	0.00	0.00	4.00	0.0	
2366	0.00	0.00	0.00	0.0	
330	0.00	0.00	1.53	0.0	
536	0.00	0.00	0.00	0.0	
3114	0.00	0.00	0.00	0.0	
	word_freq_our	word_freq_over	word_freq_remove	word_freq_internet	\
2940	0.15	0.1	0.00	0.00	
1303	0.43	0.6	0.43	0.26	
3468	0.00	0.0	0.00	0.00	
3181	0.00	0.0	0.00	0.00	
794	0.56	0.0	0.00	0.00	
...	
1861	0.00	0.0	0.00	0.00	
2366	4.16	0.0	0.00	0.00	
330	0.00	0.0	0.00	0.00	
536	0.00	0.0	0.00	0.00	
3114	0.00	0.0	0.00	0.00	
	word_freq_order	word_freq_mail	...	word_freq_conference	\
2940	0.55	0.00	...	0.0	
1303	0.69	0.52	...	0.0	
3468	0.00	0.00	...	0.0	
3181	0.00	0.00	...	0.0	
794	1.01	0.56	...	0.0	
...	
1861	0.00	0.00	...	0.0	
2366	0.00	0.00	...	0.0	
330	0.00	0.00	...	0.0	
536	0.00	0.00	...	0.0	
3114	0.00	0.00	...	0.0	
	char_freq_%3B	char_freq_%28	char_freq_%5B	char_freq_%21	\
2940	0.203	0.195	0.05	0.000	

1303	0.000	0.108	0.00	0.271
3468	0.000	0.000	0.00	0.153
3181	0.000	0.000	0.00	0.000
794	0.000	0.186	0.00	0.056
...
1861	0.000	0.000	0.00	0.613
2366	0.000	0.689	0.00	0.689
330	0.000	0.000	0.00	1.434
536	0.000	0.407	0.00	0.203
3114	0.000	0.484	0.00	0.484
char_freq_%24 char_freq_%23 capital_run_length_average \				
2940	0.014	0.000	2.880	
1303	0.243	0.013	6.395	
3468	0.000	0.000	1.933	
3181	0.000	0.000	4.333	
794	0.056	0.000	2.153	
...
1861	0.000	0.000	1.000	
2366	0.000	0.000	1.300	
330	0.000	0.000	7.055	
536	0.610	0.000	4.133	
3114	0.000	0.000	2.500	
capital_run_length_longest capital_run_length_total				
2940		45	1080	
1303		583	1375	
3468		7	58	
3181		20	26	
794		53	532	
...	
1861		1	14	
2366		4	13	
330		75	127	
536		17	62	
3114		15	65	
[3680 rows x 57 columns] word_freq_make word_freq_address word_freq_all				
word_freq_3d \				
1472	0.00	0.00	0.00	0.0
258	0.00	0.00	0.33	0.0
3564	0.00	0.00	0.00	0.0
65	0.66	0.00	0.66	0.0
4303	0.00	0.00	0.00	0.0
...
1405	0.00	0.00	0.35	0.0
2312	0.00	2.59	1.29	0.0
2804	0.00	0.00	0.00	0.0

2047	0.00	0.00	0.00	0.0
2597	0.00	0.00	0.00	0.0
word_freq_our word_freq_over word_freq_remove word_freq_internet \				
1472	0.00	0.00	0.00	0.00
258	0.99	0.99	0.33	0.33
3564	0.00	0.00	0.00	0.00
65	0.00	0.00	0.00	0.00
4303	0.00	0.00	0.00	0.00
...
1405	0.00	0.70	0.35	0.35
2312	1.29	0.00	0.00	0.00
2804	0.00	0.00	0.00	0.00
2047	0.34	0.00	0.00	0.00
2597	0.00	0.00	0.00	0.00
word_freq_order word_freq_mail ... word_freq_conference \				
1472	0.0	0.00	...	0.00
258	0.0	0.00	...	0.00
3564	0.0	0.00	...	0.00
65	0.0	0.66	...	0.00
4303	0.0	2.77	...	0.00
...
1405	0.0	0.00	...	0.00
2312	0.0	0.00	...	0.00
2804	0.0	0.00	...	0.00
2047	0.0	0.00	...	0.34
2597	0.0	0.00	...	0.00
char_freq_%3B char_freq_%28 char_freq_%5B char_freq_%21 \				
1472	0.144	0.000	0.000	3.907
258	0.000	0.108	0.000	0.000
3564	0.000	0.000	0.000	0.000
65	0.000	0.000	0.000	2.205
4303	0.000	0.000	0.000	0.438
...
1405	0.000	0.061	0.000	0.061
2312	0.000	0.000	0.000	0.000
2804	0.000	0.000	0.000	0.000
2047	0.088	0.132	0.000	0.000
2597	0.142	0.000	0.142	0.000
char_freq_%24 char_freq_%23 capital_run_length_average \				
1472	0.000	0.000	13.928	
258	0.162	0.054	2.195	
3564	0.000	0.000	1.214	
65	0.000	0.000	3.184	
4303	0.000	0.000	1.214	

```
...      ...      ...      ...
1405      0.000    0.122    2.302
2312      0.000    0.000    1.000
2804      0.000    0.000    2.000
2047      0.000    0.000    1.250
2597      0.000    0.000    1.717
```

```
capital_run_length_longest  capital_run_length_total
1472                  70          195
258                   50          202
3564                  4           17
65                    34          121
4303                  3           17
...
1405                  21          99
2312                  1           13
2804                  4           6
2047                  7           85
2597                  12          67
```

```
[921 rows x 57 columns] 2940      0
1303      1
3468      0
3181      0
794       1
..
1861      0
2366      0
330       1
536       1
3114      0
Name: class, Length: 3680, dtype: int64 1472      1
258       1
3564      0
65        1
4303      0
..
1405      1
2312      0
2804      0
2047      0
2597      0
Name: class, Length: 921, dtype: int64
```

Naive Bayes GaussianNB

```
[27]: start=time.time()
gnb=GaussianNB()
```

```
gnb.fit(X_train_scaled,y_train)
ttgnb=time.time()-start
y_pred=gnb.predict(X_test_scaled)
print("The training time for Gaussian Naive Bayes is: ",ttgnb)
```

The training time for Gaussian Naive Bayes is: 0.00466465950012207

```
[39]: start=time.time()
y_pred=gnb.predict(X_test_scaled)
ptgnb=time.time()-start
print("The Prediction time for Gaussian Naive Bayes is: ",ptgnb)
```

The Prediction time for Gaussian Naive Bayes is: 0.0009808540344238281

Accuracy

```
[30]: accuracy_score(y_test,y_pred)
```

```
[30]: 0.8327904451682954
```

Precision

```
[31]: precision_score(y_test,y_pred)
```

```
[31]: 0.7145790554414785
```

Recall

```
[32]: recall_score(y_test,y_pred)
```

```
[32]: 0.9586776859504132
```

F1 Score

```
[33]: f1_score(y_test,y_pred)
```

```
[33]: 0.8188235294117647
```

```
[34]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[419 139]
 [ 15 348]]
```

Specificity and False Positive Rate

```
[36]: TN, FP, FN, TP = cm.ravel()
specificity = TN / (TN + FP)
false_positive_rate = FP / (FP + TN)

print("Specificity:", specificity)
print("False Positive Rate:", false_positive_rate)
```

```
Specificity: 0.7508960573476703  
False Positive Rate: 0.24910394265232974
```

```
[37]: print("Overall Report of model \n\n",classification_report(y_test,y_pred))
```

Overall Report of model

	precision	recall	f1-score	support
0	0.97	0.75	0.84	558
1	0.71	0.96	0.82	363
accuracy			0.83	921
macro avg	0.84	0.85	0.83	921
weighted avg	0.87	0.83	0.83	921

MultinomialNB

```
[40]: start=time.time()  
mnb=MultinomialNB()  
mnb.fit(X_train,y_train)  
ttmnb=time.time()-start  
  
print("The training time for Multinomial Naive Bayes is: ",ttmnb)
```

The training time for Multinomial Naive Bayes is: 0.009273767471313477

```
[41]: start=time.time()  
y_pred=mnb.predict(X_test)  
ptmnb=time.time()-start  
print("The Prediction time for Multinomial Naive Bayes is: ",ptmnb)
```

The Prediction time for Multinomial Naive Bayes is: 0.009157180786132812

Accuracy

```
[42]: accuracy_score(y_test,y_pred)
```

```
[42]: 0.7763300760043431
```

Precision Score

```
[43]: precision_score(y_test,y_pred)
```

```
[43]: 0.7198879551820728
```

Recall-score

```
[44]: recall_score(y_test,y_pred)
```

```
[44]: 0.7079889807162535
```

F1-score

```
[45]: f1_score(y_test,y_pred)
```

```
[45]: 0.7138888888888889
```

```
[46]: print("Confusion matrix is \n",confusion_matrix(y_test,y_pred))
```

```
Confusion matrix is
[[458 100]
 [106 257]]
```

```
[48]: TN, FP, FN, TP = cm.ravel()
specificity = TN / (TN + FP)
false_positive_rate = FP / (FP + TN)

print("Specificity:", specificity)
print("False Positive Rate:", false_positive_rate)
```

Specificity: 0.7508960573476703

False Positive Rate: 0.24910394265232974

```
[47]: print("Overall Report of model \n\n",classification_report(y_test,y_pred))
```

Overall Report of model

	precision	recall	f1-score	support
0	0.81	0.82	0.82	558
1	0.72	0.71	0.71	363
accuracy			0.78	921
macro avg	0.77	0.76	0.77	921
weighted avg	0.78	0.78	0.78	921

BernoulliNB

```
[49]: start=time.time()
bnb=BernoulliNB()
bnb.fit(X_train,y_train)
ttbnb=time.time()-start
print("The training time for Bernoulli Naive Bayes is: ",ttbnb)
```

The training time for Bernoulli Naive Bayes is: 0.02569127082824707

```
[50]: start=time.time()
y_pred=bnb.predict(X_test)
```

```
ptbnb=time.time()-start  
  
print("The prediction time for Bernoulli Naive Bayes is: ",ptbnb)
```

The prediction time for Bernoulli Naive Bayes is: 0.003880739212036133

Accuracy

```
[51]: accuracy_score(y_test,y_pred)
```

```
[51]: 0.8762214983713354
```

Precision

```
[52]: precision_score(y_test,y_pred)
```

```
[52]: 0.8716417910447761
```

Recall Score

```
[53]: recall_score(y_test,y_pred)
```

```
[53]: 0.8044077134986226
```

F1-score

```
[54]: f1_score(y_test,y_pred)
```

```
[54]: 0.836676217765043
```

```
[56]: print("Confusion matrix is \n",confusion_matrix(y_test,y_pred))
```

Confusion matrix is

```
[[515  43]  
 [ 71 292]]
```

```
[57]: TN, FP, FN, TP = cm.ravel()  
 specificity = TN / (TN + FP)  
 false_positive_rate = FP / (FP + TN)  
  
print("Specificity:", specificity)  
print("False Positive Rate:", false_positive_rate)
```

Specificity: 0.7508960573476703

False Positive Rate: 0.24910394265232974

```
[58]: print("Overall Report of model \n\n",classification_report(y_test,y_pred))
```

Overall Report of model

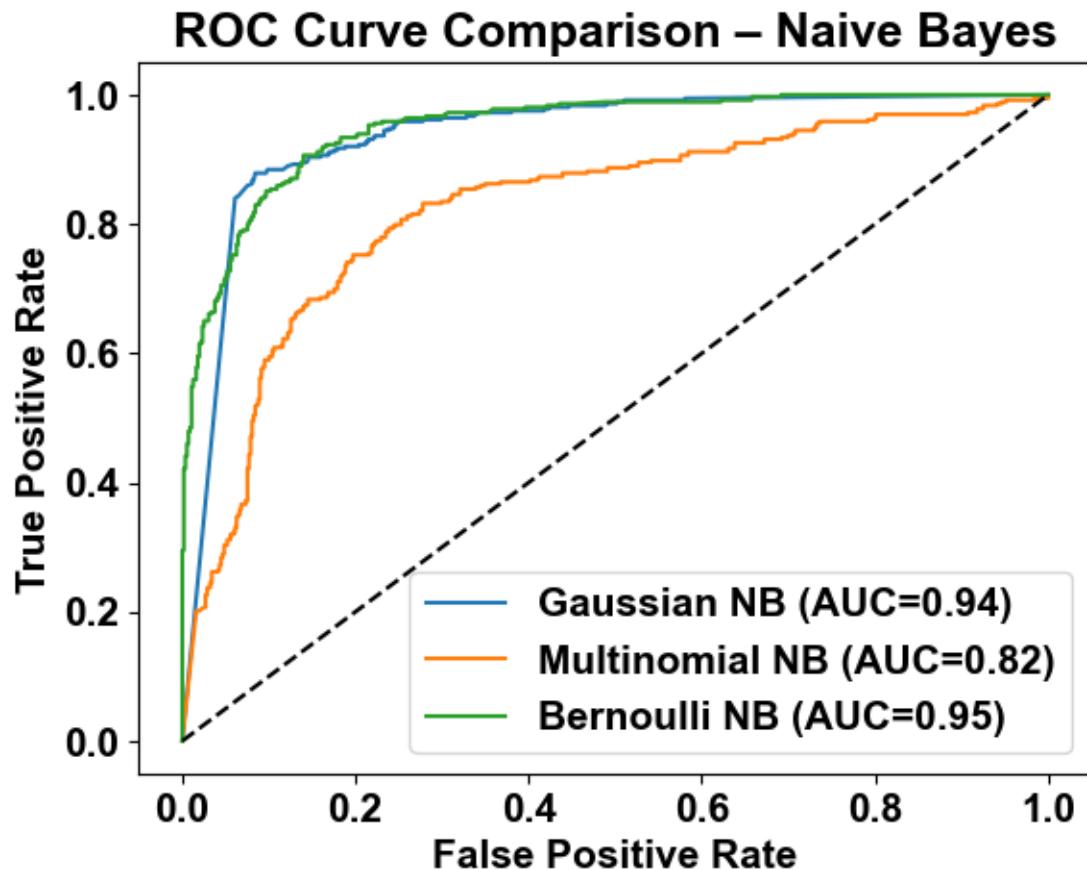
precision	recall	f1-score	support
-----------	--------	----------	---------

0	0.88	0.92	0.90	558
1	0.87	0.80	0.84	363
accuracy			0.88	921
macro avg	0.88	0.86	0.87	921
weighted avg	0.88	0.88	0.88	921

ROC curve

```
[59]: y_prob_gnb=gnb.predict_proba(X_test_scaled)[:,1]
y_prob_mnb=mnb.predict_proba(X_test)[:,1]
y_prob_bnb=bnb.predict_proba(X_test)[:,1]
```

```
[60]: fpr_gnb,tpr_gnb,_=roc_curve(y_test,y_prob_gnb)
fpr_mnb,tpr_mnb,_=roc_curve(y_test,y_prob_mnb)
fpr_bnb,tpr_bnb,_=roc_curve(y_test,y_prob_bnb)
auc_gnb=auc(fpr_gnb,tpr_gnb)
auc_mnb=auc(fpr_mnb,tpr_mnb)
auc_bnb=auc(fpr_bnb,tpr_bnb)
plt.plot(fpr_gnb,tpr_gnb,label=f'Gaussian NB (AUC={auc_gnb:.2f})')
plt.plot(fpr_mnb,tpr_mnb,label=f'Multinomial NB (AUC={auc_mnb:.2f})')
plt.plot(fpr_bnb,tpr_bnb,label=f'Bernoulli NB (AUC={auc_bnb:.2f})')
plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve Comparison - Naive Bayes")
plt.legend()
plt.savefig("roc_curve.png", dpi=300, bbox_inches="tight")
plt.show()
```



K-Nearest Neighbour

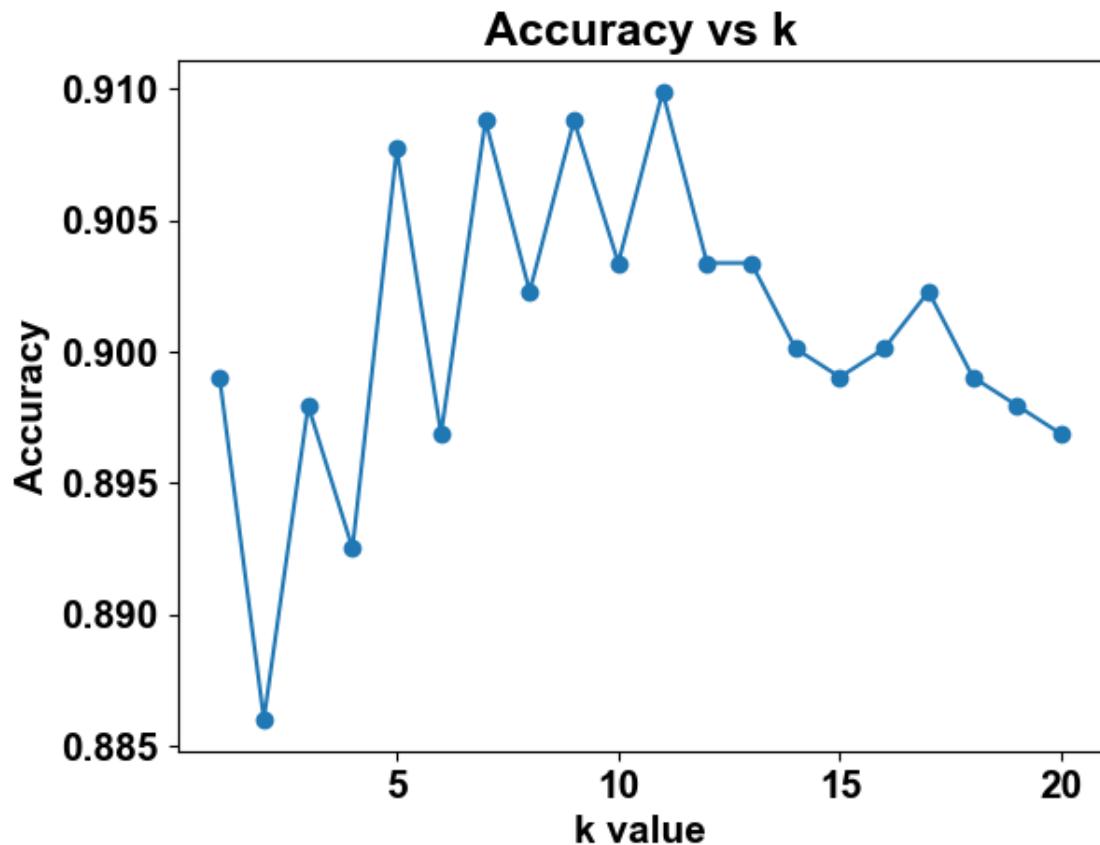
Basic Model

```
[61]: knn = KNeighborsClassifier()
knn.fit(X_train_scaled, y_train)
y_pred_knn = knn.predict(X_test_scaled)
```

Statistical significance test Anova-N

```
[76]: k_values = range(1, 21)
accuracies = []
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled, y_train)
    accuracies.append(knn.score(X_test_scaled, y_test))
plt.plot(k_values, accuracies, marker='o')
plt.xlabel('k value')
plt.ylabel('Accuracy')
plt.title('Accuracy vs k')
```

```
plt.savefig("accuracy_vs_k.png", dpi=300, bbox_inches="tight")
plt.show()
```



Stratified K-Fold

```
[77]: skf=StratifiedKFold(n_splits=5,shuffle=True,random_state=42)
cv_scores_base=cross_val_score(
    knn,
    X_train_scaled,
    y_train,
    cv=skf,
    scoring='accuracy'
)
print("Base KNN CV Accuracy:",cv_scores_base.mean())
```

Base KNN CV Accuracy: 0.8932065217391303

Grid Search

```
[78]: param_grid={
    'n_neighbors':list(range(1,31,2)),
```

```

    'weights':['uniform','distance'],
    'metric':['euclidean','manhattan']
}
grid=GridSearchCV(
    knn,
    param_grid,
    cv=skf,
    scoring='accuracy',
    n_jobs=-1
)
grid.fit(X_train_scaled,y_train)
print("Grid Best Params:",grid.best_params_)
print("Grid Best CV Accuracy:",grid.best_score_)

```

Grid Best Params: {'metric': 'manhattan', 'n_neighbors': 9, 'weights': 'distance'}

Grid Best CV Accuracy: 0.9252717391304348

Randomized Search

```
[79]: from scipy.stats import randint
param_dist={
    'n_neighbors':randint(1,30),
    'weights':['uniform','distance'],
    'metric':['euclidean','manhattan']
}
rand=RandomizedSearchCV(
    knn,
    param_distributions=param_dist,
    n_iter=15,
    cv=skf,
    scoring='accuracy',
    random_state=42,
    n_jobs=-1
)
rand.fit(X_train_scaled,y_train)
print("Random Best Params:",rand.best_params_)
print("Random Best CV Accuracy:",rand.best_score_)
```

Random Best Params: {'metric': 'manhattan', 'n_neighbors': 6, 'weights': 'distance'}

Random Best CV Accuracy: 0.9241847826086957

Final KNN Model

```
[80]: best_params=grid.best_params_
knn_final=KNeighborsClassifier(
    n_neighbors=best_params['n_neighbors'],
    weights=best_params['weights'],
```

```

        metric=best_params['metric']
    )
knn_final.fit(X_train_scaled,y_train)
y_pred_final=knn_final.predict(X_test_scaled)

```

Metrics

```
[81]: from sklearn.metrics import confusion_matrix,roc_curve,auc
def compute_metrics(y_true,y_pred):
    cm=confusion_matrix(y_true,y_pred)
    tn,fp,fn,tp=cm.ravel()
    accuracy=(tp+tn)/(tp+tn+fp+fn)
    precision=tp/(tp+fp)
    recall=tp/(tp+fn)
    f1=2*precision*recall/(precision+recall)
    specificity=tn/(tn+fp)
    fpr=fp/(fp+tn)
    return accuracy,precision,recall,f1,specificity,fpr,cm
```

```
[82]: start=time.time()
knn_final.fit(X_train_scaled,y_train)
train_time=time.time()-start

start=time.time()
y_pred_knn=knn_final.predict(X_test_scaled)
pred_time=time.time()-start
acc,prec,rec,f1,spec,fpr,cm=compute_metrics(y_test,y_pred_knn)
print("Final KNN Metrics")
print("Accuracy:",acc)
print("Precision:",prec)
print("Recall:",rec)
print("F1 Score:",f1)
print("Specificity:",spec)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)
```

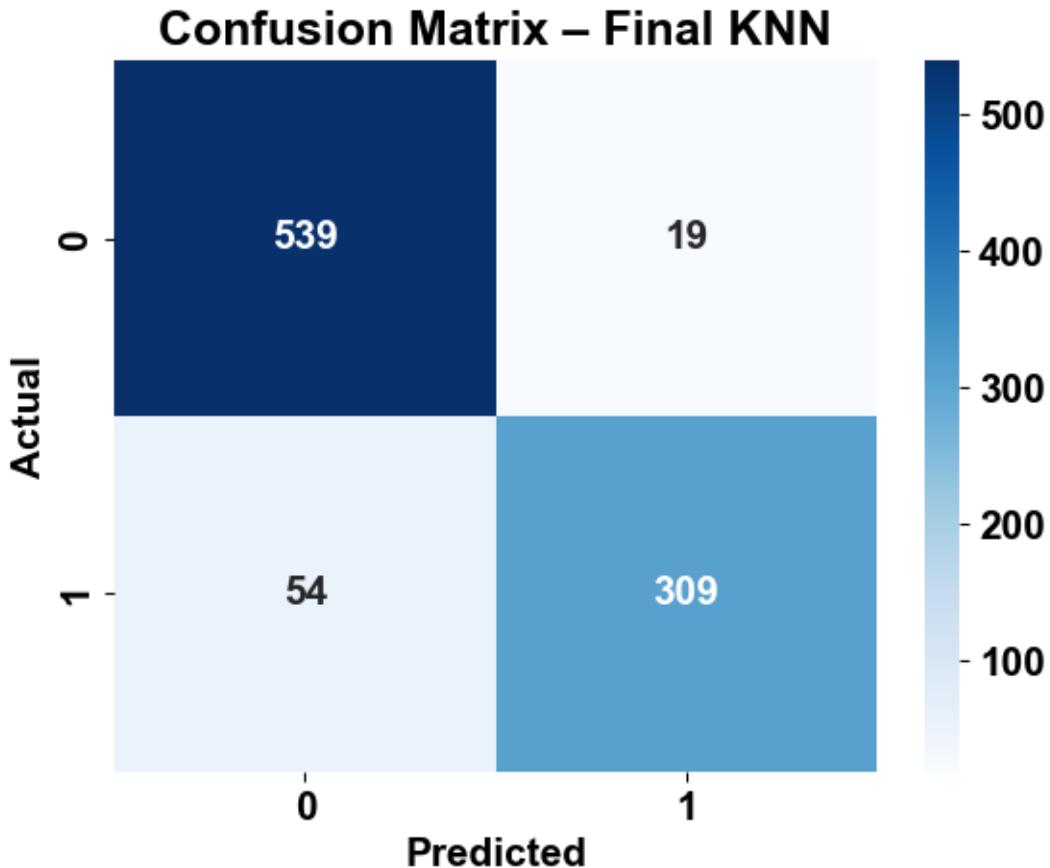
Final KNN Metrics

```

Accuracy: 0.9207383279044516
Precision: 0.9420731707317073
Recall: 0.8512396694214877
F1 Score: 0.894356005788712
Specificity: 0.9659498207885304
False Positive Rate: 0.034050179211469536
Training Time: 0.004589557647705078
Prediction Time: 0.09273862838745117
```

Confustion Matrix

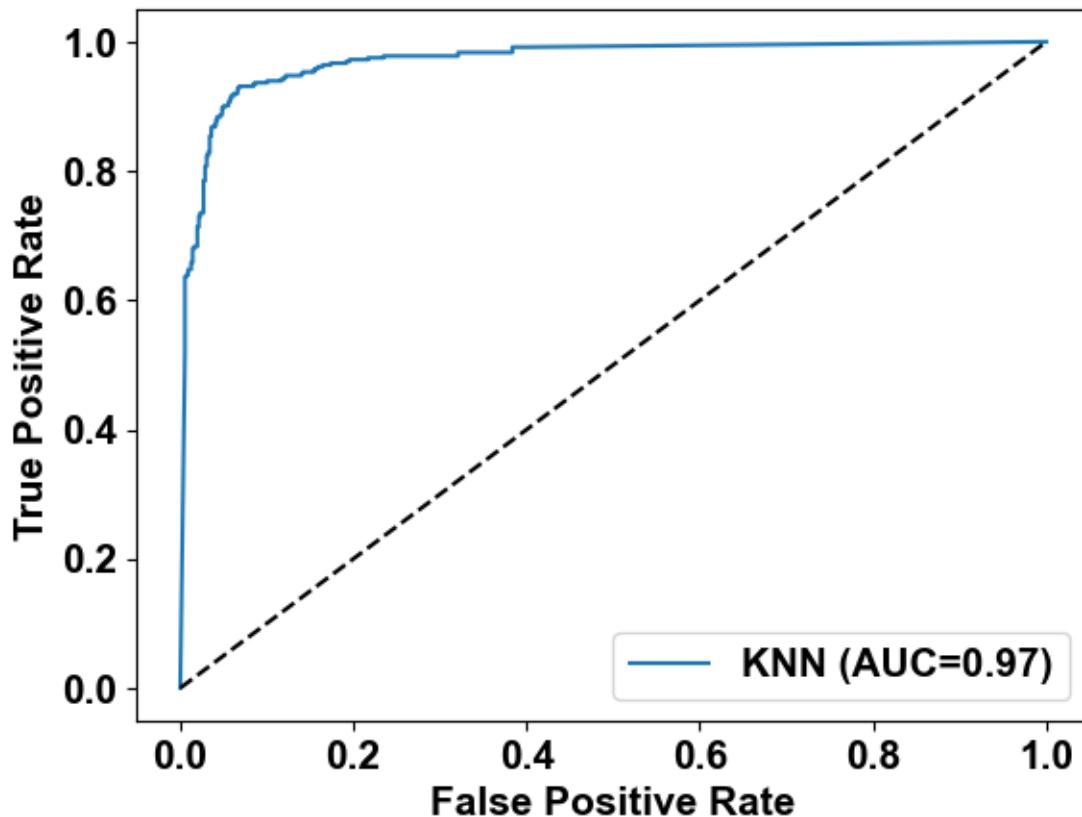
```
[83]: sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Final KNN")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



ROC Curve for KNN

```
[85]: y_prob_knn=knn_final.predict_proba(X_test_scaled)[:,1]
fpr_knn,tpr_knn,_=roc_curve(y_test,y_prob_knn)
auc_knn=auc(fpr_knn,tpr_knn)
plt.plot(fpr_knn,tpr_knn,label=f'KNN (AUC={auc_knn:.2f})')
plt.plot([0,1],[0,1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Final KNN")
plt.legend()
plt.savefig("roc_curve_finalKNN.png", dpi=300, bbox_inches="tight")
plt.show()
```

ROC Curve – Final KNN



```
[86]: best_params=grid.best_params_
optimal_k=best_params['n_neighbors']
```

Evaluating all models using multiple metrics

KDTree

```
[87]: start=time.time()
knn_kd=KNeighborsClassifier(
    n_neighbors=optimal_k,
    weights=best_params['weights'],
    metric=best_params['metric'],
    algorithm='kd_tree'
)
knn_kd.fit(X_train_scaled,y_train)
train_time_kd=time.time()-start
start=time.time()
y_pred_kd=knn_kd.predict(X_test_scaled)
pred_time_kd=time.time()-start
```

```
[88]: acc_kd,prec_kd,rec_kd,f1_kd,sp,fpr,cm=compute_metrics(y_test,y_pred_kd)
print("Final KDtree Metrics")
print("Accuracy:",acc_kd)
print("Precision:",prec_kd)
print("Recall:",rec_kd)
print("F1 Score:",f1_kd)
print("Specificity:",sp)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)
```

Final KDtree Metrics
 Accuracy: 0.9207383279044516
 Precision: 0.9420731707317073
 Recall: 0.8512396694214877
 F1 Score: 0.894356005788712
 Specificity: 0.9659498207885304
 False Positive Rate: 0.034050179211469536
 Training Time: 0.004589557647705078
 Prediction Time: 0.09273862838745117

BallTree

```
[89]: start=time.time()
knn_bt=KNeighborsClassifier(
    n_neighbors=optimal_k,
    weights=best_params['weights'],
    metric=best_params['metric'],
    algorithm='ball_tree'
)
knn_bt.fit(X_train_scaled,y_train)
train_time_bt=time.time()-start
start=time.time()
y_pred_bt=knn_bt.predict(X_test_scaled)
pred_time_bt=time.time()-start
```

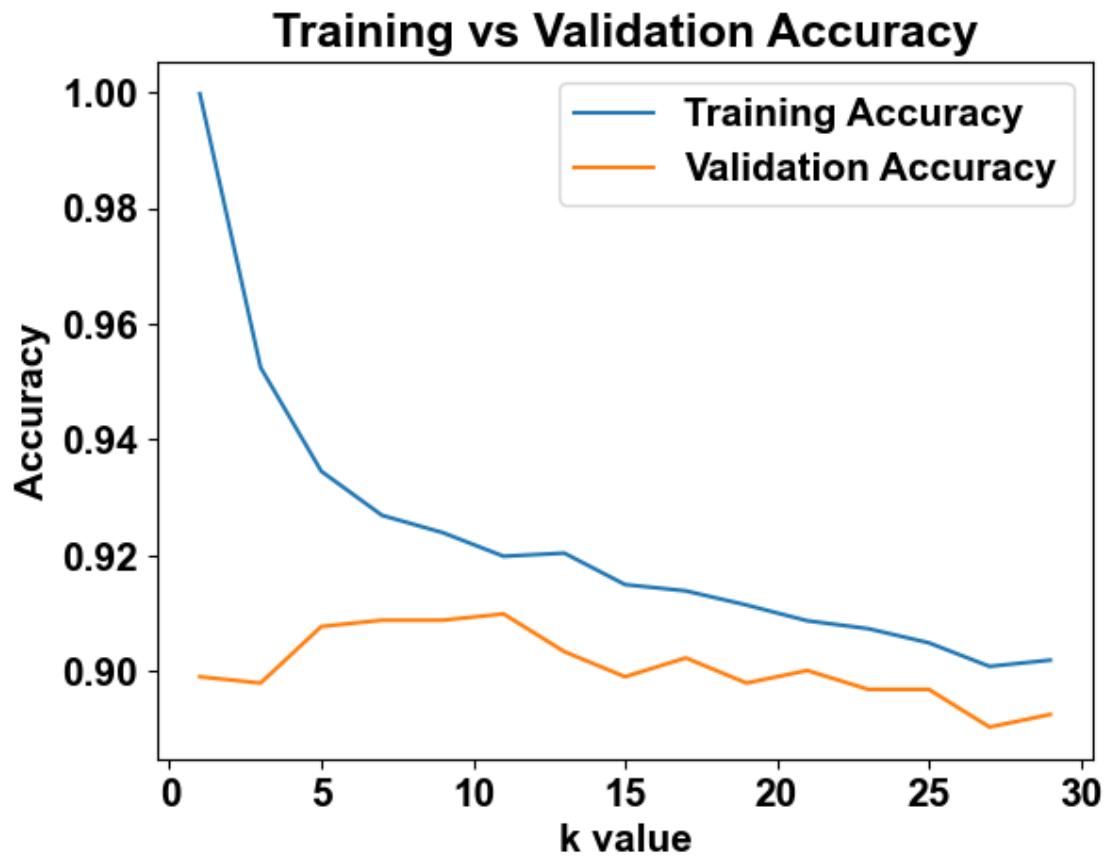
```
[90]: acc_bt,prec_bt,rec_bt,f1_bt,sp,fpr,cm=compute_metrics(y_test,y_pred_bt)
print("Final Balltree Metrics")
print("Accuracy:",acc_bt)
print("Precision:",prec_bt)
print("Recall:",rec_bt)
print("F1 Score:",f1_bt)
print("Specificity:",sp)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)
```

Final Balltree Metrics
 Accuracy: 0.9207383279044516

```
Precision: 0.9420731707317073
Recall: 0.8512396694214877
F1 Score: 0.894356005788712
Specificity: 0.9659498207885304
False Positive Rate: 0.034050179211469536
Training Time: 0.004589557647705078
Prediction Time: 0.09273862838745117
```

Training vs Validation

```
[92]: train_acc=[]
val_acc=[]
k_values=range(1,31,2)
for k in k_values:
    knn=KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled,y_train)
    train_acc.append(knn.score(X_train_scaled,y_train))
    val_acc.append(knn.score(X_test_scaled,y_test))
plt.plot(k_values,train_acc,label='Training Accuracy')
plt.plot(k_values,val_acc,label='Validation Accuracy')
plt.xlabel('k value')
plt.ylabel('Accuracy')
plt.title('Training vs Validation Accuracy')
plt.legend()
plt.savefig("training_vs_validation_accuracy.png", dpi=300, bbox_inches="tight")
plt.show()
```



Cross Validation Scores

```
[102]: gnb_scores=cross_val_score(  
        GaussianNB(),  
        X_train,  
        y_train,  
        cv=skf,  
        scoring='precision'  
)
```

```
[103]: mnb_scores=cross_val_score(  
        MultinomialNB(),  
        X_train,  
        y_train,  
        cv=skf,  
        scoring='precision'  
)
```

```
[104]: bnb_scores=cross_val_score(  
        BernoulliNB(),
```

```
X_train,  
y_train,  
cv=skf,  
scoring='precision'  
)
```

```
[105]: knn_kd_scores=cross_val_score(  
        knn_kd,  
        X_train_scaled,  
        y_train,  
        cv=skf,  
        scoring='precision'  
)
```

```
[106]: knn_bt_scores=cross_val_score(  
        knn_bt,  
        X_train_scaled,  
        y_train,  
        cv=skf,  
        scoring='precision'  
)
```

One way ANOVA Test

```
[107]: from scipy.stats import f_oneway  
F_stat,p_value=f_oneway(  
    gnb_scores,  
    mnb_scores,  
    bnb_scores,  
    knn_kd_scores,  
    knn_bt_scores  
)  
  
print("F-statistic:",F_stat)  
print("p-value:",p_value)
```

```
F-statistic: 147.4451411146714  
p-value: 1.537483731179684e-14
```

Mean Accuracy to find the best model

```
[108]: print("Gaussian NB Mean precision:",gnb_scores.mean())  
print("Multinomial NB Mean precision:",mnb_scores.mean())  
print("Bernoulli NB Mean precision:",bnb_scores.mean())  
print("KNN KDTree Mean precision:",knn_kd_scores.mean())  
print("KNN BallTree Mean precision:",knn_bt_scores.mean())
```

```
Gaussian NB Mean precision: 0.7011869673414239  
Multinomial NB Mean precision: 0.737146025168714
```

```
Bernoulli NB Mean precision: 0.8894670120013128
KNN KDTree Mean precision: 0.948881300360718
KNN BallTree Mean precision: 0.948881300360718
```

```
[109]: best_model=max(
    [
        ("Gaussian NB",gnb_scores.mean()),
        ("Multinomial NB",mnb_scores.mean()),
        ("Bernoulli NB",bnb_scores.mean()),
        ("KNN KDTree",knn_kd_scores.mean()),
        ("KNN BallTree",knn_bt_scores.mean())
    ],
    key=lambda x:x[1]
)

print("Best Model:",best_model)
```

```
Best Model: ('KNN KDTree', np.float64(0.948881300360718))
```

```
[ ]:
```

Naïve Bayes Performance Comparison

Metric	Gaussian NB	Multinomial NB	Bernoulli NB
Accuracy	0.8327	0.7763	0.8762
Precision	0.7145	0.7198	0.8716
Recall	0.9586	0.7079	0.8044
F1 Score	0.8188	0.7138	0.8366
Specificity	0.7508	0.7508	0.7508
Training Time (s)	0.0046	0.0092	0.0256

KNN Hyperparameter Tuning Results

Search Method	Best k	Best CV Accuracy	Best Parameters
Grid Search	9	0.9252	'metric': 'manhattan', 'weights': 'distance'
Randomized Search	6	0.9241	'metric': 'manhattan', 'weights': 'distance'

KNN Performance using Different Search Methods

Metric (KDTree)	Value
Optimal k	9
Accuracy	0.9207
Precision	0.9420
Recall	0.8512
F1 Score	0.8943
Training Time (s)	0.0056
Prediction Time (s)	0.0406

Metric (BallTree)	Value
Optimal k	9
Accuracy	0.9207
Precision	0.9420
Recall	0.8512
F1 Score	0.8943
Training Time (s)	0.0056
Prediction Time (s)	0.0406

KDTree vs BallTree Comparison

Criterion	KDTree	BallTree
Accuracy	0.9207	0.9207
Training Time (s)	0.0056	0.0056
Prediction Time (s)	0.0406	0.0406
Memory Usage	Low / Medium	Medium / High

Conclusion

Naïve Bayes provides fast and stable performance with high bias, whereas optimized KNN achieves better accuracy and generalization through careful hyperparameter tuning, with KDTree and BallTree improving computational efficiency.

References

- [Scikit-learn: Naïve Bayes](#)
- [Scikit-learn: KNN](#)
- [Scikit-learn: Hyperparameter Optimization](#)
- [Spambase Dataset](#)