

Degree & Branch	B.E. Computer Science & Engineering	Semester VI
Subject Code & Name	UCS2612 – Machine Learning Algorithms Laboratory	
Academic Year	2025–2026 (Even)	Batch 2023–2027
Due Date	27.01.2026	

Experiment 2: Binary Classification using Naïve Bayes and K-Nearest Neighbors

Objective

To implement Naïve Bayes and K-Nearest Neighbors (KNN) classifiers for a binary classification problem, evaluate them using multiple performance metrics, visualize model behavior, and analyze overfitting, underfitting, and bias-variance characteristics.

Dataset


A benchmark binary classification dataset containing numerical features and two class labels is used.

Dataset reference:

- Kaggle: urlcolorSpambase Dataset


```
[23]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
import math
from matplotlib import rcParams
from sklearn.model_selection import train_test_split,
    →StratifiedKFold, GridSearchCV,
    →RandomizedSearchCV,cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB, MultinomialNB,
    →BernoulliNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score,
    →recall_score, f1_score, confusion_matrix, roc_curve, auc
rcParams['font.family']='Arial'
rcParams['font.weight']='bold'
rcParams['font.size']=15
rcParams['axes.labelweight']='bold'
rcParams['axes.titleweight']='bold'
rcParams['xtick.labelsize']=15
```

```
plt.savefig("class_distribution.png", dpi=300,  
           bbox_inches="tight")
```



Histogram

```
[13]: bxwidth=1  
rows=math.ceil(len(df.columns)/3)  
fig,axes=plt.subplots(rows,3,figsize=(15,4*rows))  
axes=axes.flatten()  
columns=df.columns  
for i,(ax,col) in enumerate(zip(axes,columns)):  
    ax.hist(df[col],bins=20)  
    ax.set_title(f"{col.upper()} Distribution",pad=20)  
    ax.set_xlabel(col,labelpad=0)  
    ax.set_ylabel("Frequency",labelpad=10)  
for spine in ax.spines.values():  
    spine.set_linewidth(bxwidth)  
for j in range(i+1,len(axes)):  
    fig.delaxes(axes[j])  
plt.subplots_adjust(hspace=1.5,wspace=1.3)  
plt.tight_layout()  
plt.savefig("histogram.png", dpi=300, bbox_inches="tight")  
plt.show()
```



Train-Test Split

```
[19]: X=df.iloc[:, :-1]
y=df.iloc[:, -1]
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
→2,stratify=y,random_state=42)
print(X_train,X_test,y_train,y_test)
```

Z-score or Standard Scaling

```
[21]: scaler=StandardScaler()
X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.transform(X_test)
print(X_train,X_test,y_train,y_test)
```

Naive Bayes GaussianNB

```
[27]: start=time.time()
gnb=GaussianNB()
gnb.fit(X_train_scaled,y_train)
ttgnb=time.time()-start
y_pred=gnb.predict(X_test_scaled)
print("The training time for Gaussian Naive Bayes is: ",ttgnb)
```

```
[39]: start=time.time()
y_pred=gnb.predict(X_test_scaled)
ptgnb=time.time()-start
print("The Prediction time for Gaussian Naive Bayes is: "
→",ptgnb)
```

Accuracy

```
[30]: accuracy_score(y_test,y_pred)
```

Precision

```
[31]: precision_score(y_test,y_pred)
```

Recall

```
[32]: recall_score(y_test,y_pred)
```

F1 Score

```
[33]: f1_score(y_test,y_pred)
```

```
[34]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

Specificity and False Positive Rate

```
[36]: TN, FP, FN, TP = cm.ravel()
specificity = TN / (TN + FP)
false_positive_rate = FP / (FP + TN)

print("Specificity:", specificity)
print("False Positive Rate:", false_positive_rate)
```

```
[37]: print("Overall Report of model")
      ↪\n\n",classification_report(y_test,y_pred))
```

MultinomialNB

```
[40]: start=time.time()
mnb=MultinomialNB()
mnb.fit(X_train,y_train)
ttmnb=time.time()-start

print("The training time for Multinomial Naive Bayes is:")
      ↪",ttmnb)
```

The training time for Multinomial Naive Bayes is: 0.
↪009273767471313477

```
[41]: start=time.time()
y_pred=mnb.predict(X_test)
ptmnb=time.time()-start
print("The Prediction time for Multinomial Naive Bayes is:")
      ↪",ptmnb)
```

The Prediction time for Multinomial Naive Bayes is: 0.
↪009157180786132812

Accuracy

```
[42]: accuracy_score(y_test,y_pred)
```

```
[42]: 0.7763300760043431
```

Precision Score

```
[43]: precision_score(y_test,y_pred)
```

Recall-score

```
[44]: recall_score(y_test,y_pred)
```

F1-score

```
[45]: f1_score(y_test,y_pred)
```

```
[46]: print("Confusion matrix is:\n",confusion_matrix(y_test,y_pred))
```

```
[48]: TN, FP, FN, TP = cm.ravel()
specificity = TN / (TN + FP)
false_positive_rate = FP / (FP + TN)

print("Specificity:", specificity)
print("False Positive Rate:", false_positive_rate)
```

```
[47]: print("Overall Report of model:\n",classification_report(y_test,y_pred))
```

BernoulliNB

```
[49]: start=time.time()
bnb=BernoulliNB()
bnb.fit(X_train,y_train)
ttbnb=time.time()-start
print("The training time for Bernoulli Naive Bayes is:\n",ttbnb)
```

```
[50]: start=time.time()
y_pred=bnb.predict(X_test)
ptbnb=time.time()-start

print("The prediction time for Bernoulli Naive Bayes is:\n",ptbnb)
```

The prediction time for Bernoulli Naive Bayes is: 0.
→003880739212036133

Accuracy

```
[51]: accuracy_score(y_test,y_pred)
```

Precision

```
[52]: precision_score(y_test,y_pred)
```

Recall Score

```
[53]: recall_score(y_test,y_pred)
```

F1-score

```
[54]: f1_score(y_test,y_pred)
```

```
[56]: print("Confusion matrix is\n→\n",confusion_matrix(y_test,y_pred))
```

```
Confusion matrix is  
[[515 43]  
 [ 71 292]]
```

```
[57]: TN, FP, FN, TP = cm.ravel()  
specificity = TN / (TN + FP)  
false_positive_rate = FP / (FP + TN)  
  
print("Specificity:", specificity)  
print("False Positive Rate:", false_positive_rate)
```


```
Specificity: 0.7508960573476703  
False Positive Rate: 0.24910394265232974
```

```
[58]: print("Overall Report of model\n→\n",classification_report(y_test,y_pred))
```

```
ROC curve
```

```
[59]: y_prob_gnb=gnb.predict_proba(X_test_scaled)[:,1]  
y_prob_mnb=mnb.predict_proba(X_test)[:,1]  
y_prob_bnb=bnb.predict_proba(X_test)[:,1]
```

```
[60]: fpr_gnb,tpr_gnb,_=roc_curve(y_test,y_prob_gnb)  
fpr_mnb,tpr_mnb,_=roc_curve(y_test,y_prob_mnb)  
fpr_bnb,tpr_bnb,_=roc_curve(y_test,y_prob_bnb)  
auc_gnb=auc(fpr_gnb,tpr_gnb)  
auc_mnb=auc(fpr_mnb,tpr_mnb)  
auc_bnb=auc(fpr_bnb,tpr_bnb)  
plt.plot(fpr_gnb,tpr_gnb,label=f'Gaussian NB (AUC={auc_gnb:.  
→2f})')  
plt.plot(fpr_mnb,tpr_mnb,label=f'Multinomial NB (AUC={auc_mnb:  
→.2f})')  
plt.plot(fpr_bnb,tpr_bnb,label=f'Bernoulli NB (AUC={auc_bnb:.  
→2f})')  
plt.plot([0,1],[0,1],'k--')  
plt.xlabel("False Positive Rate")  
plt.ylabel("True Positive Rate")  
plt.title("ROC Curve Comparison - Naive Bayes")  
plt.legend()  
plt.savefig("roc_curve.png", dpi=300, bbox_inches="tight")  
plt.show()
```




K-Nearest Neighbour

Basic Model

```
[61]: knn = KNeighborsClassifier()
knn.fit(X_train_scaled, y_train)
y_pred_knn = knn.predict(X_test_scaled)
```

Statistical significance test Anova-N

```
[76]: k_values = range(1, 21)
accuracies = []
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled, y_train)
    accuracies.append(knn.score(X_test_scaled, y_test))
plt.plot(k_values, accuracies, marker='o')
plt.xlabel('k value')
plt.ylabel('Accuracy')
plt.title('Accuracy vs k')
plt.savefig("accuracy_vs_k.png", dpi=300, bbox_inches="tight")
plt.show()
```



Stratified K-Fold

```
[77]: skf=StratifiedKFold(n_splits=5,shuffle=True,random_state=42)
cv_scores_base=cross_val_score(
    knn,
    X_train_scaled,
    y_train,
    cv=skf,
    scoring='accuracy'
)
print("Base KNN CV Accuracy:",cv_scores_base.mean())
```

Grid Search

```
[78]: param_grid={
    'n_neighbors':list(range(1,31,2)),
    'weights':[‘uniform’,’distance’],
    ‘metric’:[‘euclidean’,’manhattan’]
}
grid=GridSearchCV(
    knn,
    param_grid,
    cv=skf,
    scoring='accuracy',
    n_jobs=-1
)
```

```
grid.fit(X_train_scaled,y_train)
print("Grid Best Params:",grid.best_params_)
print("Grid Best CV Accuracy:",grid.best_score_)
```

```
Grid Best Params: {'metric': 'manhattan', 'n_neighbors': 9, 'weights': 'distance'}
Grid Best CV Accuracy: 0.9252717391304348
```

Randomized Search

```
[79]: from scipy.stats import randint
param_dist={
    'n_neighbors':randint(1,30),
    'weights':['uniform','distance'],
    'metric':['euclidean','manhattan']
}
rand=RandomizedSearchCV(
    knn,
    param_distributions=param_dist,
    n_iter=15,
    cv=skf,
    scoring='accuracy',
    random_state=42,
    n_jobs=-1
)
rand.fit(X_train_scaled,y_train)
print("Random Best Params:",rand.best_params_)
print("Random Best CV Accuracy:",rand.best_score_)
```

```
Random Best Params: {'metric': 'manhattan', 'n_neighbors': 6, 'weights': 'distance'}
Random Best CV Accuracy: 0.9241847826086957
```

Final KNN Model

```
[80]: best_params=grid.best_params_
knn_final=KNeighborsClassifier(
    n_neighbors=best_params['n_neighbors'],
    weights=best_params['weights'],
    metric=best_params['metric']
)
knn_final.fit(X_train_scaled,y_train)
y_pred_final=knn_final.predict(X_test_scaled)
```

Metrics


```
[81]: from sklearn.metrics import confusion_matrix,roc_curve,auc
def compute_metrics(y_true,y_pred):
    cm=confusion_matrix(y_true,y_pred)
    tn,fp,fn,tp=cm.ravel()
    accuracy=(tp+tn)/(tp+tn+fp+fn)
    precision=tp/(tp+fp)
    recall=tp/(tp+fn)
    f1=2*precision*recall/(precision+recall)
    specificity=tn/(tn+fp)
    fpr=fp/(fp+tn)
    return accuracy,precision,recall,f1,specificity,fpr,cm
```

```
[82]: start=time.time()
knn_final.fit(X_train_scaled,y_train)
train_time=time.time()-start

start=time.time()
y_pred_knn=knn_final.predict(X_test_scaled)
pred_time=time.time()-start
acc,prec,rec,f1,spec,fpr,cm=compute_metrics(y_test,y_pred_knn)
print("Final KNN Metrics")
print("Accuracy:",acc)
print("Precision:",prec)
print("Recall:",rec)
print("F1 Score:",f1)
print("Specificity:",spec)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)
```


Confustion Matrix

```
[83]: sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.title("Confusion Matrix - Final KNN")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



ROC Curve for KNN

```
[85]: y_prob_knn=knn_final.predict_proba(X_test_scaled)[:,1]
fpr_knn,tpr_knn,_=roc_curve(y_test,y_prob_knn)
auc_knn=auc(fpr_knn,tpr_knn)
plt.plot(fpr_knn,tpr_knn,label=f'KNN (AUC={auc_knn:.2f})')
plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Final KNN")
plt.legend()
plt.savefig("roc_curve_finalKNN.png", dpi=300,
            bbox_inches="tight")
plt.show()
```



```
[86]: best_params=grid.best_params_
optimal_k=best_params['n_neighbors']
```

Evaluating all models using multiple metrics

KDTree

```
[87]: start=time.time()
knn_kd=KNeighborsClassifier(
    n_neighbors=optimal_k,
    weights=best_params['weights'],
    metric=best_params['metric'],
    algorithm='kd_tree',
)
knn_kd.fit(X_train_scaled,y_train)
train_time_kd=time.time()-start
start=time.time()
y_pred_kd=knn_kd.predict(X_test_scaled)
pred_time_kd=time.time()-start
```

```
[88]: acc_kd,prec_kd,rec_kd,f1_kd,sp,fpr,cm=compute_metrics(y_test,y_pred_kd)
print("Final KDtree Metrics")
print("Accuracy:",acc_kd)
print("Precision:",prec_kd)
print("Recall:",rec_kd)
```

```

print("F1 Score:",f1_kd)
print("Specificity:",sp)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)

```

BallTree

```
[89]: start=time.time()
knn_bt=KNeighborsClassifier(
    n_neighbors=optimal_k,
    weights=best_params['weights'],
    metric=best_params['metric'],
    algorithm='ball_tree'
)
knn_bt.fit(X_train_scaled,y_train)
train_time_bt=time.time()-start
start=time.time()
y_pred_bt=knn_bt.predict(X_test_scaled)
pred_time_bt=time.time()-start


```

```
[90]: acc_bt,prec_bt,rec_bt,f1_bt,sp,fpr,cm=compute_metrics(y_test,y_pred_bt)
print("Final Balltree Metrics")
print("Accuracy:",acc_bt)
print("Precision:",prec_bt)
print("Recall:",rec_bt)
print("F1 Score:",f1_bt)
print("Specificity:",sp)
print("False Positive Rate:",fpr)
print("Training Time:",train_time)
print("Prediction Time:",pred_time)
```

Training vs Validation

```
[92]: train_acc=[]
val_acc=[]
k_values=range(1,31,2)
for k in k_values:
    knn=KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled,y_train)
    train_acc.append(knn.score(X_train_scaled,y_train))
    val_acc.append(knn.score(X_test_scaled,y_test))
plt.plot(k_values,train_acc,label='Training Accuracy')
plt.plot(k_values,val_acc,label='Validation Accuracy')
plt.xlabel('k value')
plt.ylabel('Accuracy')
plt.title('Training vs Validation Accuracy')
plt.legend()
```

```
plt.savefig("training_vs_validation_accuracy.png", dpi=300, bbox_inches="tight")
plt.show()
```



Cross Validation Scores

```
[102]: gnb_scores=cross_val_score(
    GaussianNB(),
    X_train,
    y_train,
    cv=skf,
    scoring='precision'
)
```

```
[103]: mnb_scores=cross_val_score(
    MultinomialNB(),
    X_train,
    y_train,
    cv=skf,
    scoring='precision'
)
```

```
[104]: bnb_scores=cross_val_score(
    BernoulliNB(),
    X_train,
```

```
        y_train,  
        cv=skf,  
        scoring='precision'  
)
```

```
[105]: knn_kd_scores=cross_val_score(  
        knn_kd,  
        X_train_scaled,  
        y_train,  
        cv=skf,  
        scoring='precision'  
)
```

```
[106]: knn_bt_scores=cross_val_score(  
        knn_bt,  
        X_train_scaled,  
        y_train,  
        cv=skf,  
        scoring='precision'  
)
```

One way ANOVA Test

```
[107]: from scipy.stats import f_oneway  
F_stat,p_value=f_oneway(  
    gnb_scores,  
    mnb_scores,  
    bnb_scores,  
    knn_kd_scores,  
    knn_bt_scores  
)  
  
print("F-statistic:",F_stat)  
print("p-value:",p_value)
```

F-statistic: 147.4451411146714
p-value: 1.537483731179684e-14

Mean Accuracy to find the best model

```
[108]: print("Gaussian NB Mean precision:",gnb_scores.mean())  
print("Multinomial NB Mean precision:",mnb_scores.mean())  
print("Bernoulli NB Mean precision:",bnb_scores.mean())  
print("KNN KDTree Mean precision:",knn_kd_scores.mean())  
print("KNN BallTree Mean precision:",knn_bt_scores.mean())
```

```
[109]: best_model=max(  
    [  
        ("Gaussian NB",gnb_scores.mean()),
```

```

        ("Multinomial NB", mnb_scores.mean()),
        ("Bernoulli NB", bnb_scores.mean()),
        ("KNN KDTree", knn_kd_scores.mean()),
        ("KNN BallTree", knn_bt_scores.mean())
    ],
    key=lambda x:x[1]
)

print("Best Model:", best_model)

```

Best Model: ('KNN KDTree', np.float64(0.948881300360718))

[]:

Naïve Bayes Performance Comparison

Metric	Gaussian NB	Multinomial NB	Bernoulli NB
Accuracy	0.8327	0.7763	0.8762
Precision	0.7145	0.7198	0.8716
Recall	0.9586	0.7079	0.8044
F1 Score	0.8188	0.7138	0.8366
Specificity	0.7508	0.7508	0.7508
Training Time (s)	0.0046	0.0092	0.0256

KNN Hyperparameter Tuning Results

Search Method	Best k	Best CV Accuracy	Best Parameters
Grid Search	9	0.9252	'metric': 'manhattan', 'weights': 'distance'
Randomized Search	6	0.9241	'metric': 'manhattan', 'weights': 'distance'

KNN Performance using Different Search Methods

KDTree vs BallTree Comparison

Conclusion

Naïve Bayes provides fast and stable performance with high bias, whereas optimized KNN achieves better accuracy and generalization through careful

Metric (KDTree)	Value
Optimal k	9
Accuracy	0.9207
Precision	0.9420
Recall	0.8512
F1 Score	0.8943
Training Time (s)	0.0056
Prediction Time (s)	0.0406

Metric (BallTree)	Value
Optimal k	9
Accuracy	0.9207
Precision	0.9420
Recall	0.8512
F1 Score	0.8943
Training Time (s)	0.0056
Prediction Time (s)	0.0406

Criterion	KDTree	BallTree
Accuracy	0.9207	0.9207
Training Time (s)	0.0056	0.0056
Prediction Time (s)	0.0406	0.0406
Memory Usage	Low / Medium	Medium / High

hyperparameter tuning, with KDTree and BallTree improving computational efficiency.

References

- urlcolorScikit-learn: Naïve Bayes
- urlcolorScikit-learn: KNN
- urlcolorScikit-learn: Hyperparameter Optimization
- urlcolorSpambase Dataset