

**Sri Sivasubramaniya Nadar College of Engineering, Chennai**  
(An Autonomous Institution Affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester VI
Subject Code & Name	UCS2612 – Machine Learning Algorithms Laboratory	
Academic Year	2025–2026 (Even)	Batch 2023–2027
Due Date	27.01.2026	

**Experiment 3: Regression Analysis using Linear and Regularized Models**

## Objective

To implement linear and regularized regression models for predicting a continuous target variable, evaluate their performance using multiple metrics, visualize model behavior, and analyze overfitting, underfitting, and bias–variance characteristics.

## Dataset

A real-world regression dataset containing numerical and categorical features related to loan applications is used. The target variable is the **loan amount sanctioned**.

Dataset reference:

- Kaggle: [Predict Loan Amount Data](#)

## Code

```
[140]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_curve, auc
```

```
[190]: df = pd.read_csv('train.csv')
df
```

```
[190]:      Customer ID          Name Gender  Age  Income (USD) \
0        C-36995  Frederica Shealy     F   56    1933.05
1        C-33999  America Calderone    M   32    4952.91
2        C-3770    Rosetta Verne     F   65    988.19
3        C-26480     Zoe Chitty      F   65       NaN
4        C-23459    Afton Venema     F   31    2614.77
...
29995     C-43723  Angelyn Clevenger    M   38    4969.41
29996     C-32511     Silas Slaugh     M   20    1606.88
```

29997	C-5192	Carmelo Lone	F	49	NaN
29998	C-12172	Carolann Osby	M	38	2417.71
29999	C-33003	Bridget Garibaldi	F	63	3068.24

	Income Stability	Profession	Type of Employment	\
0	Low	Working	Sales staff	
1	Low	Working	NaN	
2	High	Pensioner	NaN	
3	High	Pensioner	NaN	
4	Low	Working	High skill tech staff	
...	...	...	...	...
29995	Low	Commercial associate	Managers	
29996	Low	Working	Laborers	
29997	Low	Working	Sales staff	
29998	Low	Working	Security staff	
29999	High	Pensioner	NaN	

	Location	Loan Amount Request (USD)	...	Credit Score	\
0	Semi-Urban	72809.58	...	809.44	
1	Semi-Urban	46837.47	...	780.40	
2	Semi-Urban	45593.04	...	833.15	
3	Rural	80057.92	...	832.70	
4	Semi-Urban	113858.89	...	745.55	
...	...	...	...	...	...
29995	Urban	76657.90	...	869.61	
29996	Semi-Urban	66595.14	...	729.41	
29997	Urban	81410.08	...	NaN	
29998	Semi-Urban	142524.10	...	677.27	
29999	Rural	156290.54	...	815.44	

	No. of Defaults	Has Active Credit Card	Property ID	Property Age	\
0	0	NaN	746	1933.05	
1	0	Unpossessed	608	4952.91	
2	0	Unpossessed	546	988.19	
3	1	Unpossessed	890	NaN	
4	1	Active	715	2614.77	
...	...	...	...	...	...
29995	0	Unpossessed	566	4969.41	
29996	0	Inactive	175	1606.88	
29997	0	Active	959	NaN	
29998	1	Unpossessed	375	2417.71	
29999	0	Active	344	3068.24	

	Property Type	Property Location	Co-Applicant	Property Price	\
0	4	Rural	1	119933.46	
1	2	Rural	1	54791.00	
2	2	Urban	0	72440.58	

3	2	Semi-Urban	1	121441.51
4	4	Semi-Urban	1	208567.91
...	...	...	...	...
29995	4	Urban	1	111096.56
29996	3	Urban	1	73453.94
29997	1	Rural	1	102108.02
29998	4	Urban	1	168194.47
29999	3	Rural	1	194512.60

Loan Sanction Amount (USD)

0	54607.18
1	37469.98
2	36474.43
3	56040.54
4	74008.28
...	...
29995	68992.11
29996	46616.60
29997	61057.56
29998	99766.87
29999	117217.90

[30000 rows x 24 columns]

[191]: df.describe()

	Age	Income (USD)	Loan Amount Request (USD)	\
count	30000.000000	2.542400e+04	30000.000000	
mean	40.092300	2.630574e+03	88826.333855	
std	16.045129	1.126272e+04	59536.949605	
min	18.000000	3.777000e+02	6048.240000	
25%	25.000000	1.650457e+03	41177.755000	
50%	40.000000	2.222435e+03	75128.075000	
75%	55.000000	3.090593e+03	119964.605000	
max	65.000000	1.777460e+06	621497.820000	

	Current Loan Expenses (USD)	Dependents	Credit Score	\
count	29828.000000	27507.000000	28297.000000	
mean	400.936876	2.253027	739.885381	
std	242.545375	0.951162	72.163846	
min	-999.000000	1.000000	580.000000	
25%	247.667500	2.000000	681.880000	
50%	375.205000	2.000000	739.820000	
75%	521.292500	3.000000	799.120000	
max	3840.880000	14.000000	896.260000	

No. of Defaults    Property ID    Property Age    Property Type \

count	30000.000000	30000.000000	2.515000e+04	30000.000000
mean	0.193933	501.934700	2.631119e+03	2.460067
std	0.395384	288.158086	1.132268e+04	1.118562
min	0.000000	1.000000	3.777000e+02	1.000000
25%	0.000000	251.000000	1.650450e+03	1.000000
50%	0.000000	504.000000	2.223250e+03	2.000000
75%	0.000000	751.000000	3.091408e+03	3.000000
max	1.000000	999.000000	1.777460e+06	4.000000

	Co-Applicant	Property Price	Loan Sanction Amount (USD)
count	30000.000000	3.000000e+04	29660.000000
mean	-4.743867	1.317597e+05	47649.342208
std	74.614593	9.354955e+04	48221.146686
min	-999.000000	-9.990000e+02	-999.000000
25%	1.000000	6.057216e+04	0.000000
50%	1.000000	1.099936e+05	35209.395000
75%	1.000000	1.788807e+05	74261.250000
max	1.000000	1.077967e+06	481907.320000

[192]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 24 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Customer ID      30000 non-null   object 
 1   Name              30000 non-null   object 
 2   Gender            29947 non-null   object 
 3   Age               30000 non-null   int64  
 4   Income (USD)     25424 non-null   float64
 5   Income Stability 28317 non-null   object 
 6   Profession        30000 non-null   object 
 7   Type of Employment 22730 non-null   object 
 8   Location          30000 non-null   object 
 9   Loan Amount Request (USD) 30000 non-null   float64
 10  Current Loan Expenses (USD) 29828 non-null   float64
 11  Expense Type 1     30000 non-null   object 
 12  Expense Type 2     30000 non-null   object 
 13  Dependents        27507 non-null   float64
 14  Credit Score       28297 non-null   float64
 15  No. of Defaults    30000 non-null   int64  
 16  Has Active Credit Card 28434 non-null   object 
 17  Property ID        30000 non-null   int64  
 18  Property Age        25150 non-null   float64
 19  Property Type       30000 non-null   int64  
 20  Property Location    29644 non-null   object 
 21  Co-Applicant        30000 non-null   int64
```

```
22 Property Price           30000 non-null float64
23 Loan Sanction Amount (USD) 29660 non-null float64
dtypes: float64(8), int64(5), object(11)
memory usage: 5.5+ MB
```

```
[193]: print("columns")
df.columns
```

```
columns
```

```
[193]: Index(['Customer ID', 'Name', 'Gender', 'Age', 'Income (USD)',
   'Income Stability', 'Profession', 'Type of Employment', 'Location',
   'Loan Amount Request (USD)', 'Current Loan Expenses (USD)',
   'Expense Type 1', 'Expense Type 2', 'Dependents', 'Credit Score',
   'No. of Defaults', 'Has Active Credit Card', 'Property ID',
   'Property Age', 'Property Type', 'Property Location', 'Co-Applicant',
   'Property Price', 'Loan Sanction Amount (USD)'],
  dtype='object')
```

```
[194]: df = df.drop(columns=['Customer ID', 'Name', "Property ID"])
```

```
[195]: df.isnull().sum()
```

```
[195]: Gender          53
Age              0
Income (USD)     4576
Income Stability 1683
Profession       0
Type of Employment 7270
Location         0
Loan Amount Request (USD) 0
Current Loan Expenses (USD) 172
Expense Type 1    0
Expense Type 2    0
Dependents        2493
Credit Score      1703
No. of Defaults   0
Has Active Credit Card 1566
Property Age      4850
Property Type     0
Property Location 356
Co-Applicant      0
Property Price    0
Loan Sanction Amount (USD) 340
dtype: int64
```

```
[203]: num_cols = ['Age', 'Income (USD)', 'Loan Amount Request (USD)',
```

```
'Current Loan Expenses (USD)', 'Dependents', 'Credit Score', 'Property\u2191Age',  
'Property Price', 'Loan Sanction Amount (USD)']  
  
cat_cols = [  
    'Gender', 'Type of Employment', 'Has Active Credit Card',  
    'Property Location', 'No. of Defaults', 'Income Stability', 'Property Type',\u2191  
    'Co-Applicant']
```

```
[197]: for col in num_cols:  
    df[col] = df[col].fillna(df[col].median())  
  
for col in cat_cols:  
    df[col] = df[col].fillna(df[col].mode()[0])
```

```
[198]: df['Type of Employment'] = df['Type of Employment'].fillna('Unknown')
```

Using 'Unknown' for null values in 'Type of Employment'

```
[199]: df.isnull().sum()
```

```
[199]: Gender          0  
Age             0  
Income (USD)     0  
Income Stability 0  
Profession       0  
Type of Employment 0  
Location         0  
Loan Amount Request (USD) 0  
Current Loan Expenses (USD) 0  
Expense Type 1     0  
Expense Type 2     0  
Dependents        0  
Credit Score       0  
No. of Defaults    0  
Has Active Credit Card 0  
Property Age       0  
Property Type      0  
Property Location   0  
Co-Applicant       0  
Property Price      0  
Loan Sanction Amount (USD) 0  
dtype: int64
```

```
[200]: df[num_cols].describe()
```

```
[200]:           Age  Income (USD)  Loan Amount Request (USD)  \\\n      count  30000.000000  3.000000e+04                  30000.000000
```

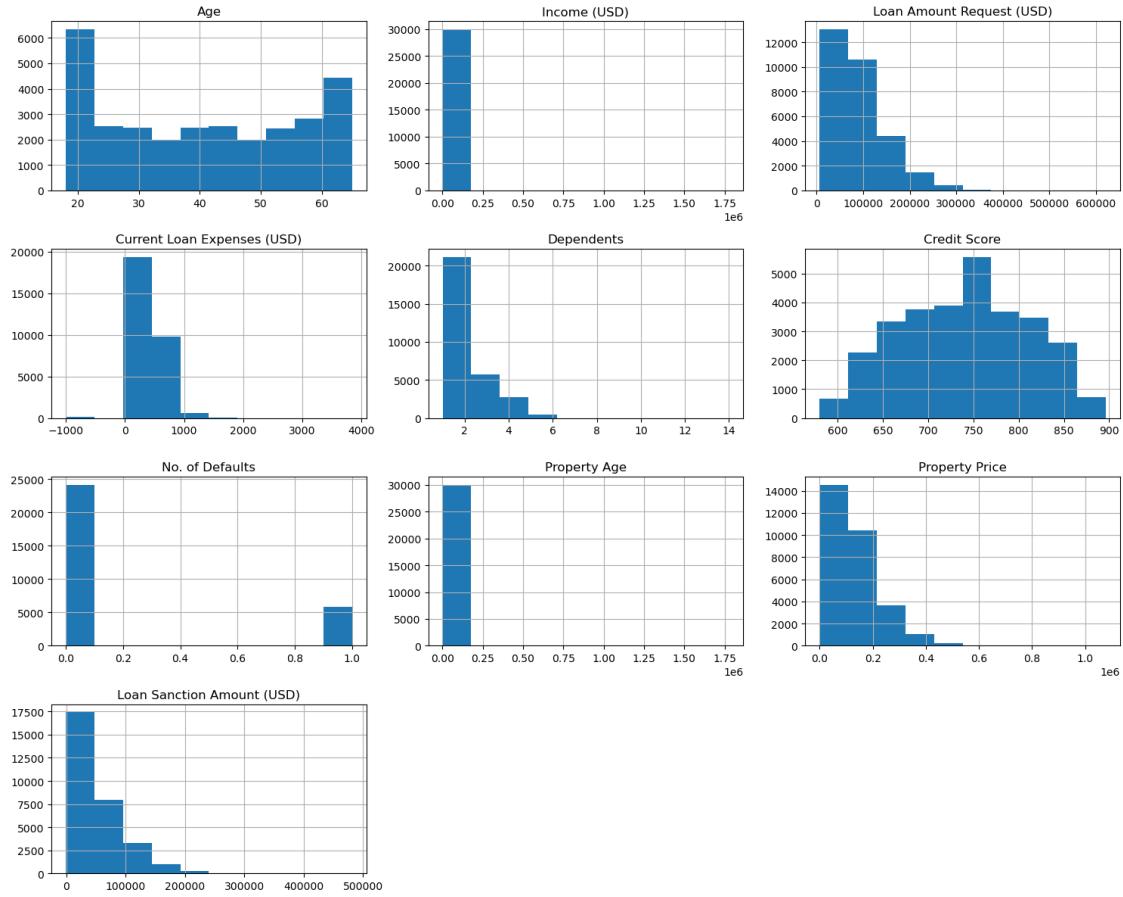
mean	40.092300	2.568320e+03	88826.333855
std	16.045129	1.036924e+04	59536.949605
min	18.000000	3.777000e+02	6048.240000
25%	25.000000	1.743305e+03	41177.755000
50%	40.000000	2.222435e+03	75128.075000
75%	55.000000	2.869142e+03	119964.605000
max	65.000000	1.777460e+06	621497.820000

	Current Loan Expenses (USD)	Dependents	Credit Score	\
count	30000.000000	30000.000000	30000.000000	
mean	400.789347	2.232000	739.881670	
std	241.856859	0.913457	70.085603	
min	-999.000000	1.000000	580.000000	
25%	248.655000	2.000000	685.415000	
50%	375.205000	2.000000	739.820000	
75%	520.102500	3.000000	795.140000	
max	3840.880000	14.000000	896.260000	

	No. of Defaults	Property Age	Property Price	\
count	30000.000000	3.000000e+04	3.000000e+04	
mean	0.193933	2.565181e+03	1.317597e+05	
std	0.395384	1.036816e+04	9.354955e+04	
min	0.000000	3.777000e+02	-9.990000e+02	
25%	0.000000	1.749812e+03	6.057216e+04	
50%	0.000000	2.223250e+03	1.099936e+05	
75%	0.000000	2.849573e+03	1.788807e+05	
max	1.000000	1.777460e+06	1.077967e+06	

	Loan Sanction Amount (USD)
count	30000.000000
mean	47508.356140
std	47965.185159
min	-999.000000
25%	0.000000
50%	35209.395000
75%	73763.532500
max	481907.320000

```
[201]: # Bar plot
df[num_cols].hist(figsize=(15,12))
plt.tight_layout()
plt.show()
```



```
[204]: # Box plot
import math

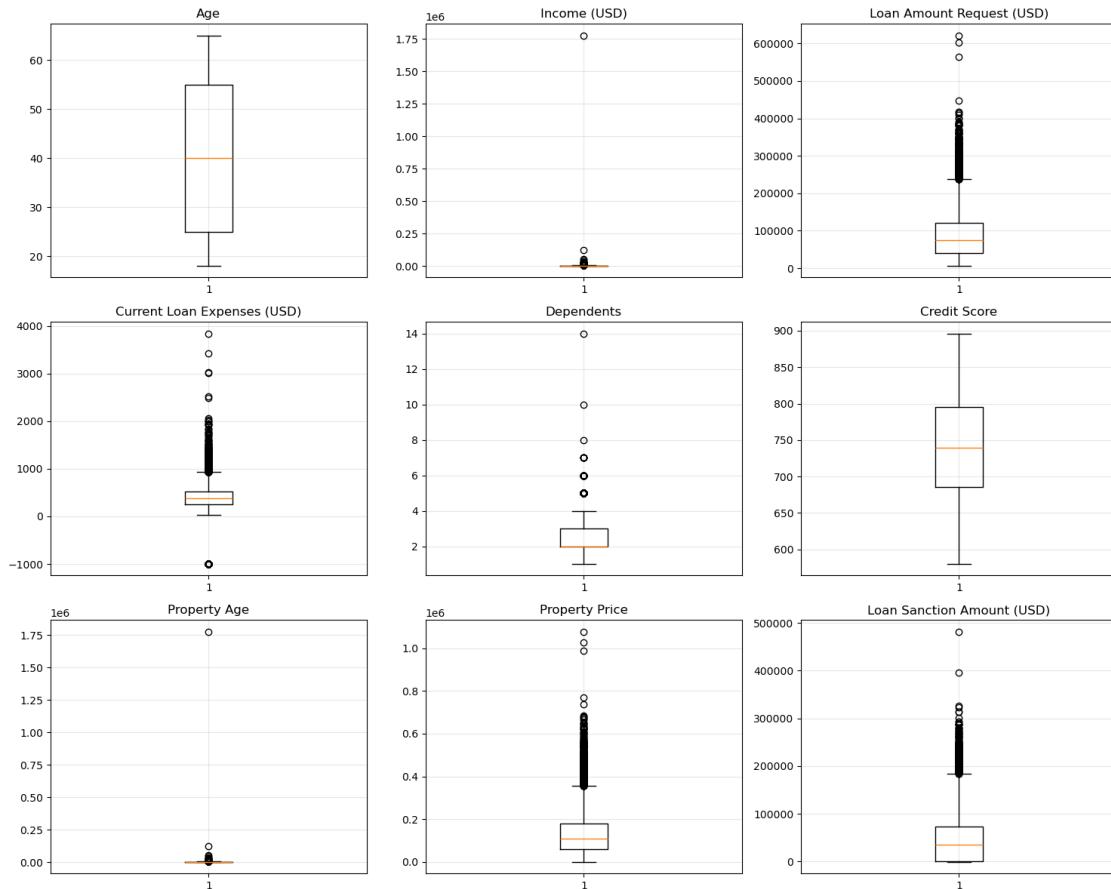
n = len(num_cols)
cols = 3
rows = math.ceil(n / cols)

fig, axes = plt.subplots(rows, cols, figsize=(5*cols, 4*rows))
axes = axes.flatten()

for i, col in enumerate(num_cols):
    axes[i].boxplot(df[col].dropna())
    axes[i].set_title(col)
    axes[i].grid(True, alpha=0.3)

# Remove unused subplots
for j in range(i+1, len(axes)):
    fig.delaxes(axes[j])
```

```
plt.tight_layout()
plt.show()
```



Huge outliers exist

```
[205]: df[num_cols].describe()
```

```
      Age   Income (USD)  Loan Amount Request (USD) \
count  30000.000000  3.000000e+04          30000.000000
mean    40.092300  2.568320e+03          88826.333855
std     16.045129  1.036924e+04         59536.949605
min    18.000000  3.777000e+02          6048.240000
25%    25.000000  1.743305e+03         41177.755000
50%    40.000000  2.222435e+03         75128.075000
75%    55.000000  2.869142e+03        119964.605000
max    65.000000  1.777460e+06        621497.820000

      Current Loan Expenses (USD)  Dependents  Credit Score  Property Age \
count          30000.000000  30000.000000  30000.000000  3.000000e+04
```

mean	400.789347	2.232000	739.881670	2.565181e+03
std	241.856859	0.913457	70.085603	1.036816e+04
min	-999.000000	1.000000	580.000000	3.777000e+02
25%	248.655000	2.000000	685.415000	1.749812e+03
50%	375.205000	2.000000	739.820000	2.223250e+03
75%	520.102500	3.000000	795.140000	2.849573e+03
max	3840.880000	14.000000	896.260000	1.777460e+06
	Property	Price	Loan Sanction	Amount (USD)
count	3.000000e+04		30000.000000	
mean	1.317597e+05		47508.356140	
std	9.354955e+04		47965.185159	
min	-9.990000e+02		-999.000000	
25%	6.057216e+04		0.000000	
50%	1.099936e+05		35209.395000	
75%	1.788807e+05		73763.532500	
max	1.077967e+06		481907.320000	

```
[206]: #winorization
outlier_summary = []
for col in num_cols:
    if df[col].nunique() <= 2:
        continue

    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower_limit = Q1 - 1.5 * IQR
    upper_limit = Q3 + 1.5 * IQR

    lower_outlier = (df[col] < lower_limit).sum()
    upper_outlier = (df[col] > upper_limit).sum()
    df[col] = df[col].clip(lower_limit, upper_limit) # outlier removed df

    outlier_summary[col] = {
        'IQR' : IQR,
        'lower_limit' : lower_limit,
        'upper_limit' : upper_limit,
        'lower_outlier' : lower_outlier ,
        'upper_outlier' : upper_outlier,
        'total_outlier' : lower_outlier + upper_outlier,
    }
    #df[col] = df[col].clip(lower_limit, upper_limit) # outlier removed df

outlier_df = pd.DataFrame(outlier_summary).T
outlier_df
```

[206] :

	IQR	lower_limit	upper_limit	\
Age	30.0000	-20.00000	100.00000	
Income (USD)	1125.8375	54.54875	4557.89875	
Loan Amount Request (USD)	78786.8500	-77002.52000	238144.88000	
Current Loan Expenses (USD)	271.4475	-158.51625	927.27375	
Dependents	1.0000	0.50000	4.50000	
Credit Score	109.7250	520.82750	959.72750	
Property Age	1099.7600	100.17250	4499.21250	
Property Price	118308.5600	-116890.68000	356343.56000	
Loan Sanction Amount (USD)	73763.5325	-110645.29875	184408.83125	

	lower_outlier	upper_outlier	total_outlier
Age	0.0	0.0	0.0
Income (USD)	0.0	1874.0	1874.0
Loan Amount Request (USD)	0.0	752.0	752.0
Current Loan Expenses (USD)	177.0	697.0	874.0
Dependents	0.0	432.0	432.0
Credit Score	0.0	0.0	0.0
Property Age	0.0	1930.0	1930.0
Property Price	0.0	863.0	863.0
Loan Sanction Amount (USD)	0.0	453.0	453.0

[207] : df[num\_cols].describe()

	Age	Income (USD)	Loan Amount Request (USD)	\
count	30000.000000	30000.000000	30000.000000	
mean	40.092300	2401.003243	87817.437653	
std	16.045129	951.994445	56222.583017	
min	18.000000	377.700000	6048.240000	
25%	25.000000	1743.305000	41177.755000	
50%	40.000000	2222.435000	75128.075000	
75%	55.000000	2869.142500	119964.605000	
max	65.000000	4557.898750	238144.880000	

	Current Loan Expenses (USD)	Dependents	Credit Score	Property Age	\
count	30000.000000	30000.000000	30000.000000	30000.000000	
mean	401.368082	2.222100	739.881670	2395.871500	
std	203.950321	0.881286	70.085603	938.567563	
min	-158.516250	1.000000	580.000000	377.700000	
25%	248.655000	2.000000	685.415000	1749.812500	
50%	375.205000	2.000000	739.820000	2223.250000	
75%	520.102500	3.000000	795.140000	2849.572500	
max	927.273750	4.500000	896.260000	4499.212500	

	Property Price	Loan Sanction Amount (USD)
count	30000.000000	30000.000000
mean	129661.354193	47039.895081

```
std      86506.620071      46300.710095
min     -999.000000     -999.000000
25%    60572.160000      0.000000
50%    109993.610000    35209.395000
75%    178880.720000    73763.532500
max    356343.560000   184408.831250
```

```
[208]: # Box plot after outlier treatment
```

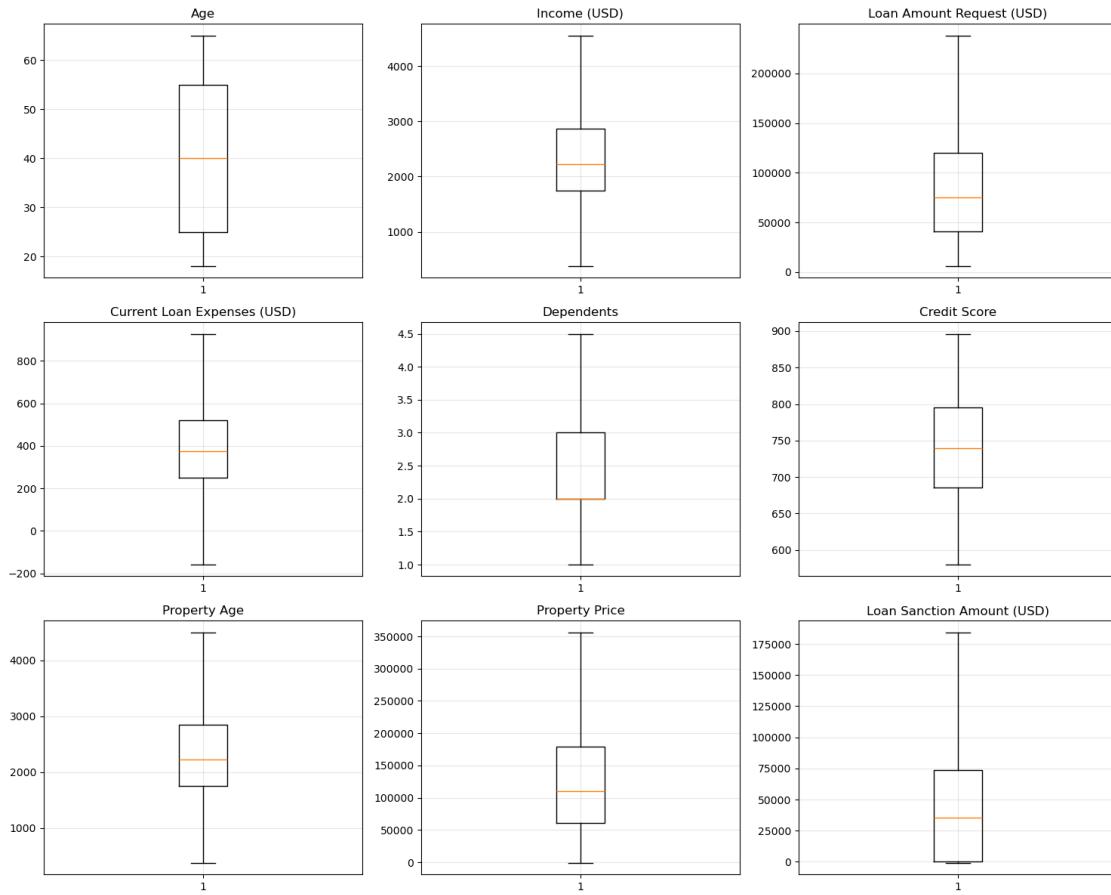
```
n = len(num_cols)
cols = 3
rows = math.ceil(n / cols)

fig, axes = plt.subplots(rows, cols, figsize=(5*cols, 4*rows))
axes = axes.flatten()

for i, col in enumerate(num_cols):
    axes[i].boxplot(df[col].dropna())
    axes[i].set_title(col)
    axes[i].grid(True, alpha=0.3)

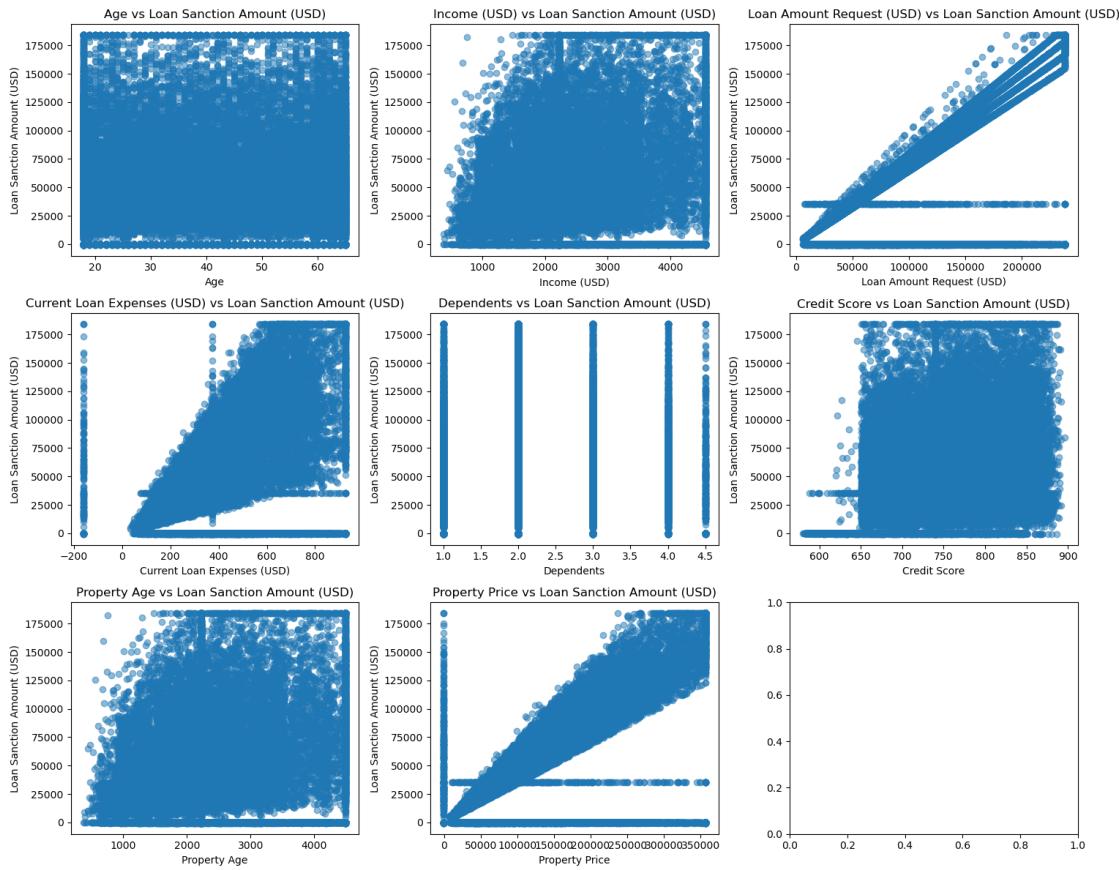
# Remove unused subplots
for j in range(i+1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



```
[210]: target = 'Loan Sanction Amount (USD)'
fig, axes = plt.subplots(rows,cols, figsize = (15,12))
axes = axes.flatten()

for i,col in enumerate(num_cols):
    if col == target:
        continue
    axes[i].scatter(df[col], df[target], alpha=0.5)
    axes[i].set_xlabel(col)
    axes[i].set_ylabel(target)
    axes[i].set_title(f'{col} vs {target}')
plt.tight_layout()
plt.show()
```



```
[211]: plt.figure(figsize=(10,6))
sns.heatmap(df[num_cols].corr(), annot=True)
plt.title("Correlational matrix")
plt.show()
```



```
[212]: y = df['Loan Sanction Amount (USD)']
X = df.drop(columns = ['Loan Sanction Amount (USD)'])
num_cols = [col for col in num_cols if col != 'Loan Sanction Amount (USD)']
```

```
[213]: # ANOVA
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.feature_selection import f_classif

selector = SelectKBest(score_func=f_classif, k=5)
X_anova_selected = selector.fit_transform(X[num_cols], y)

anova_scores = pd.DataFrame({
    'Feature': num_cols,
    'ANOVA F-Score': selector.scores_
}).sort_values(by='ANOVA F-Score', ascending=False)

anova_scores
```

```
[213]:          Feature      ANOVA F-Score
2   Loan Amount Request (USD)       1.274943
```

```

7          Property Price      1.270455
3 Current Loan Expenses (USD)  1.161522
0          Age                1.074584
1          Income (USD)       0.973369
6          Property Age       0.963895
5          Credit Score        0.930061
4          Dependents         0.900784

```

```
[214]: selected_features = [
    'Loan Amount Request (USD)',
    'Property Price',
    'Current Loan Expenses (USD)',
    'Age',
    'Income (USD)',
    'Property Age',
    'Credit Score'
]
```

```
[221]: target = 'Loan Sanction Amount (USD)'
X = df[selected_features]
y = df[target]

test_df = pd.read_csv('test.csv')
X_test = test_df[selected_features]
```

```
[222]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.model_selection import cross_val_predict
import numpy as np

lr = LinearRegression()

y_cv_pred_lr = cross_val_predict(lr, X, y, cv=5)
```

```
[223]: from sklearn.linear_model import Ridge, Lasso, ElasticNet

ridge = Ridge()
lasso = Lasso(max_iter=10000)
elastic = ElasticNet(max_iter=10000)
```

```
[249]: from sklearn.model_selection import GridSearchCV

ridge_params = {'alpha': [0.01, 0.1, 1, 10, 100]}
lasso_params = {'alpha': [0.001, 0.01, 0.1, 1, 10]}
elastic_params = {
    'alpha': [0.001, 0.01, 0.1, 1, 10],
    'l1_ratio': [0.2, 0.5, 0.8]}
```

```

}

ridge_cv = GridSearchCV(Ridge(), ridge_params, cv=5, u
    ↳scoring='neg_mean_squared_error')
lasso_cv = GridSearchCV(Lasso(max_iter=10000), lasso_params, cv=5, u
    ↳scoring='neg_mean_squared_error')
elastic_cv = GridSearchCV(ElasticNet(max_iter=10000), elastic_params, cv=5, u
    ↳scoring='neg_mean_squared_error')

ridge_cv.fit(X, y)
lasso_cv.fit(X, y)
elastic_cv.fit(X, y)

```

```
[249]: GridSearchCV(cv=5, estimator=ElasticNet(max_iter=10000),
                     param_grid={'alpha': [0.001, 0.01, 0.1, 1, 10],
                                 'l1_ratio': [0.2, 0.5, 0.8]},
                     scoring='neg_mean_squared_error')
```

```
[254]: ridge_results = pd.DataFrame(ridge_cv.cv_results_)

ridge_results[['param_alpha', 'mean_test_score', 'std_test_score']]
```

```
[254]:   param_alpha  mean_test_score  std_test_score
0          0.01      -9.329259e+08     1.674041e+07
1          0.10      -9.329259e+08     1.674041e+07
2          1.00      -9.329259e+08     1.674041e+07
3         10.00      -9.329259e+08     1.674041e+07
4        100.00      -9.329259e+08     1.674040e+07
```

```
[255]: ridge_results['RMSE'] = np.sqrt(-ridge_results['mean_test_score'])

ridge_results[['param_alpha', 'RMSE']]
```

```
[255]:   param_alpha      RMSE
0          0.01  30543.835545
1          0.10  30543.835545
2          1.00  30543.835544
3         10.00  30543.835539
4        100.00  30543.835485
```

```
[252]: lasso_results = pd.DataFrame(lasso_cv.cv_results_)
lasso_results['RMSE'] = np.sqrt(-lasso_results['mean_test_score'])

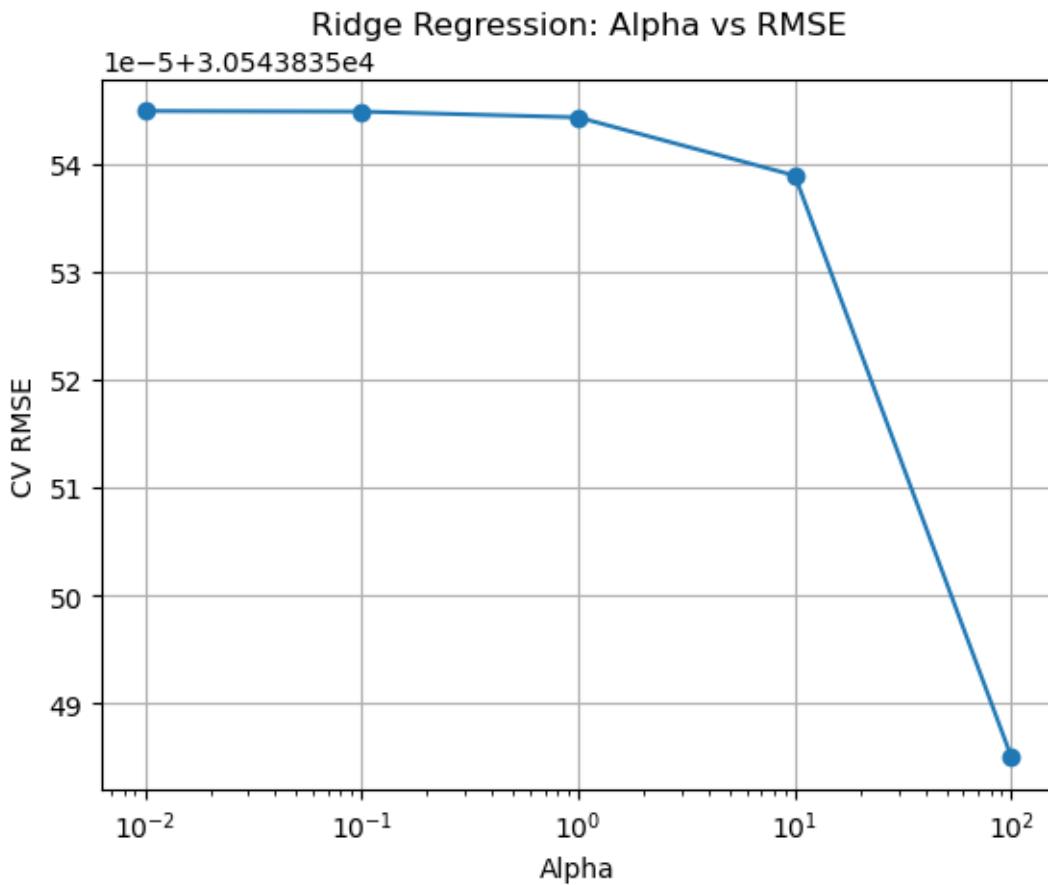
elastic_results = pd.DataFrame(elastic_cv.cv_results_)
elastic_results['RMSE'] = np.sqrt(-elastic_results['mean_test_score'])
```

```
[253]: ridge_results['RMSE']
lasso_results['RMSE']
elastic_results['RMSE']

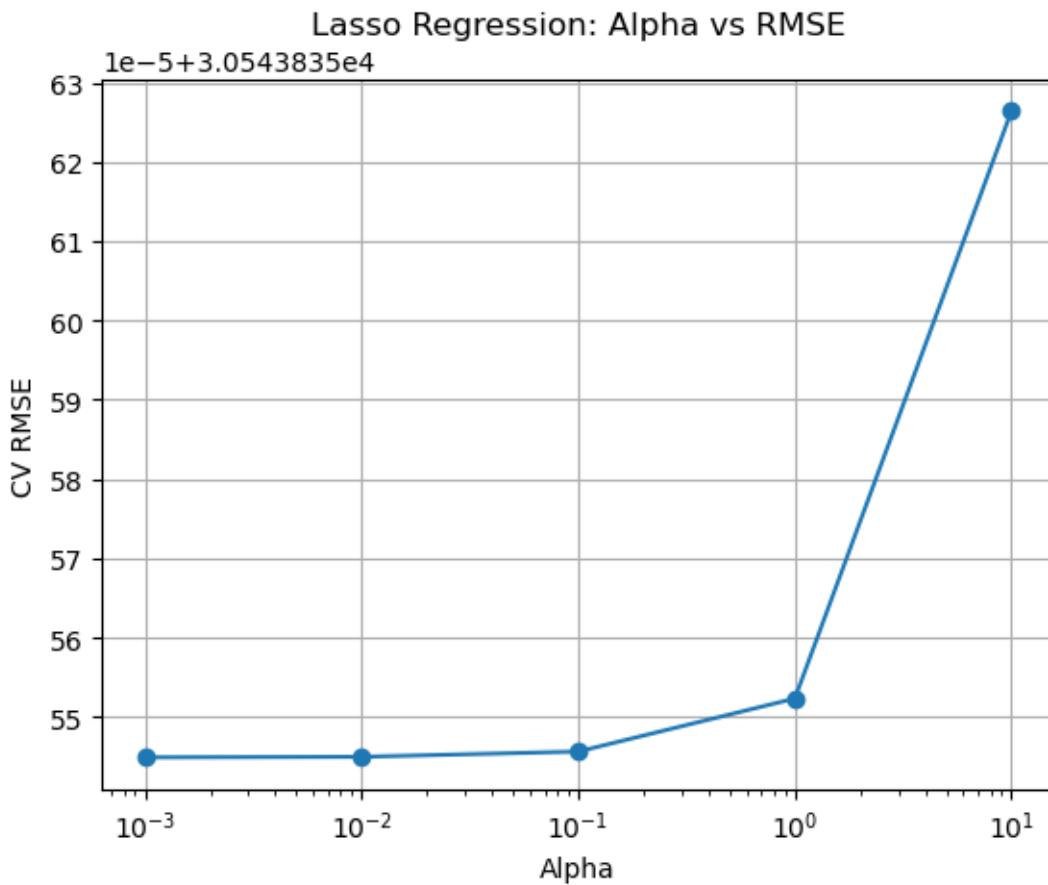
print(pd.DataFrame({'Ridge':ridge_results['RMSE'] , 'Lasso':
    ↪lasso_results['RMSE'] , 'Elastic':elastic_results['RMSE'] }))
```

	Ridge	Lasso	Elastic
0	30543.835545	30543.835545	30543.835533
1	30543.835545	30543.835545	30543.835538
2	30543.835544	30543.835546	30543.835542
3	30543.835539	30543.835552	30543.835430
4	30543.835485	30543.835626	30543.835473
5	NaN	NaN	30543.835516
6	NaN	NaN	30543.834396
7	NaN	NaN	30543.834827
8	NaN	NaN	30543.835258
9	NaN	NaN	30543.824255
10	NaN	NaN	30543.828442
11	NaN	NaN	30543.832689
12	NaN	NaN	30543.741350
13	NaN	NaN	30543.772087
14	NaN	NaN	30543.808307

```
[236]: plt.plot(
    ridge_results['param_alpha'],
    ridge_results['RMSE'],
    marker='o'
)
plt.xscale('log')
plt.xlabel('Alpha')
plt.ylabel('CV RMSE')
plt.title('Ridge Regression: Alpha vs RMSE')
plt.grid(True)
plt.show()
```



```
[256]: plt.plot(
    lasso_results['param_alpha'],
    lasso_results['RMSE'],
    marker='o'
)
plt.xscale('log')
plt.xlabel('Alpha')
plt.ylabel('CV RMSE')
plt.title('Lasso Regression: Alpha vs RMSE')
plt.grid(True)
plt.show()
```



```
[257]: ridge_results[['param_alpha', 'RMSE']].sort_values('RMSE')
```

```
[257]:   param_alpha      RMSE
4        100.00  30543.835485
3        10.00   30543.835539
2        1.00   30543.835544
1        0.10   30543.835545
0        0.01   30543.835545
```

```
[258]: elastic_results[['param_alpha', 'param_l1_ratio', 'RMSE']]\
    .sort_values('RMSE')\
    .head(10)
```

```
[258]:   param_alpha  param_l1_ratio      RMSE
12        10.00          0.2  30543.741350
13        10.00          0.5  30543.772087
14        10.00          0.8  30543.808307
9         1.00          0.2  30543.824255
10        1.00          0.5  30543.828442
```

```

11      1.00      0.8  30543.832689
6       0.10      0.2  30543.834396
7       0.10      0.5  30543.834827
8       0.10      0.8  30543.835258
3       0.01      0.2  30543.835430

```

```
[259]: best_ridge = ridge_cv.best_estimator_
best_lasso = lasso_cv.best_estimator_
best_elastic = elastic_cv.best_estimator_
```

```
[260]: def evaluate(name, y_true, y_pred):
    return {
        "Model": name,
        "MAE": mean_absolute_error(y_true, y_pred),
        "RMSE": np.sqrt(mean_squared_error(y_true, y_pred)),
        "R2": r2_score(y_true, y_pred)
    }
```

```
[261]: results = []

results.append(evaluate(
    "Linear Regression",
    y,
    cross_val_predict(lr, X, y, cv=5)
))

results.append(evaluate(
    "Ridge",
    y,
    cross_val_predict(best_ridge, X, y, cv=5)
))

results.append(evaluate(
    "Lasso",
    y,
    cross_val_predict(best_lasso, X, y, cv=5)
))

results.append(evaluate(
    "Elastic Net",
    y,
    cross_val_predict(best_elastic, X, y, cv=5)
))

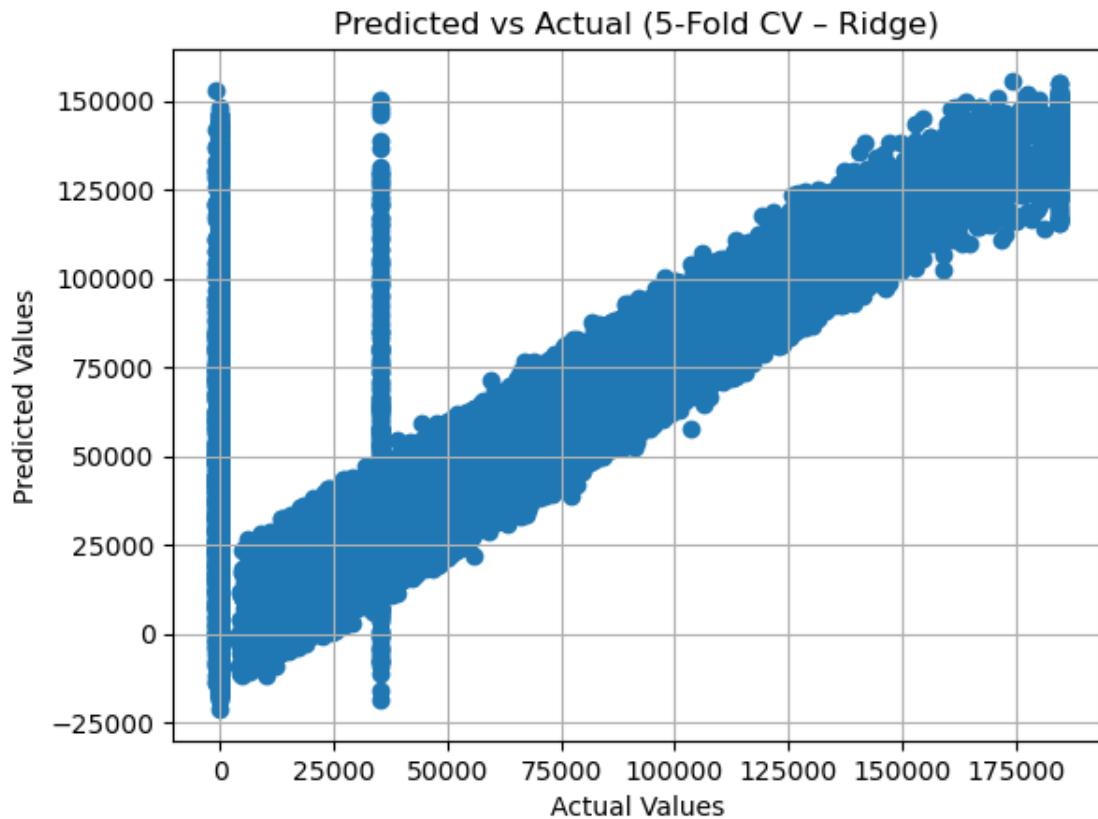
results_df = pd.DataFrame(results)
results_df
```

```
[261]:
```

	Model	MAE	RMSE	R2
0	Linear Regression	21592.300095	30543.835545	0.564803
1	Ridge	21592.300132	30543.835485	0.564803
2	Lasso	21592.300096	30543.835545	0.564803
3	Elastic Net	21592.377982	30543.741350	0.564805

```
[227]: y_pred = cross_val_predict(best_ridge, X, y, cv=5)

plt.scatter(y, y_pred)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Predicted vs Actual (5-Fold CV - Ridge)")
plt.grid(True)
plt.show()
```



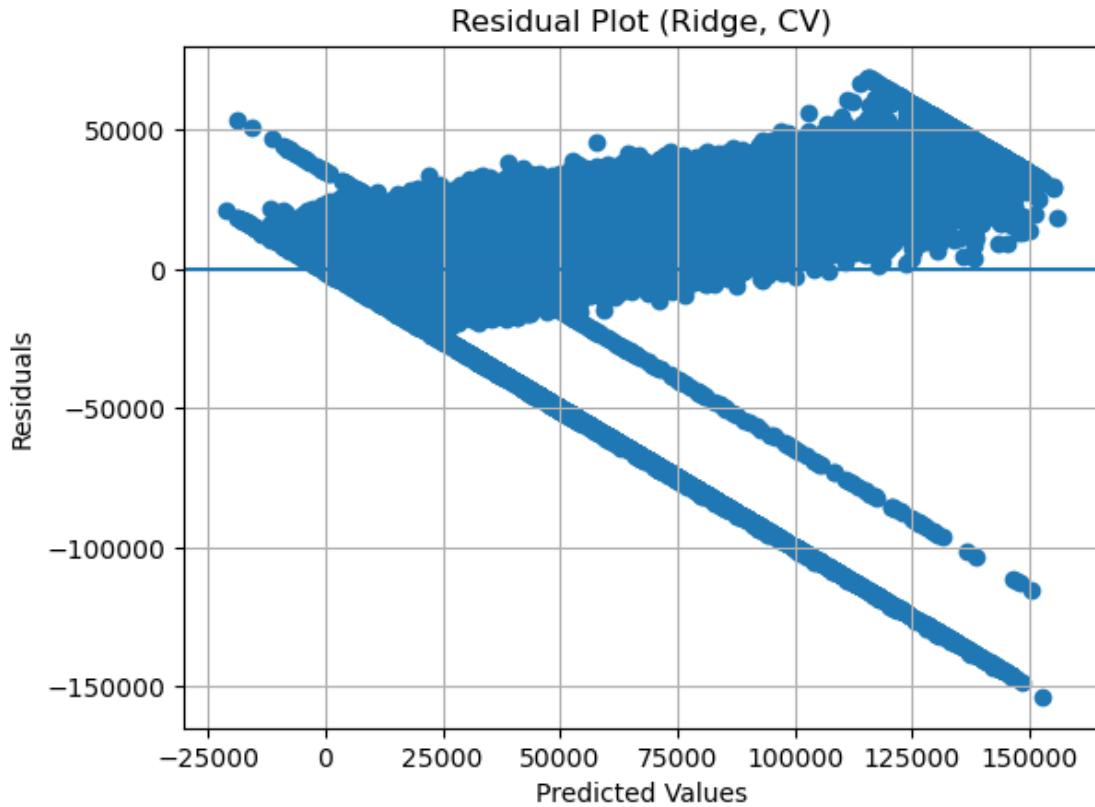
```
[262]: residuals = y - y_pred

plt.scatter(y_pred, residuals)
plt.axhline(0)
plt.xlabel("Predicted Values")
```

```

plt.ylabel("Residuals")
plt.title("Residual Plot (Ridge, CV)")
plt.grid(True)
plt.show()

```



```

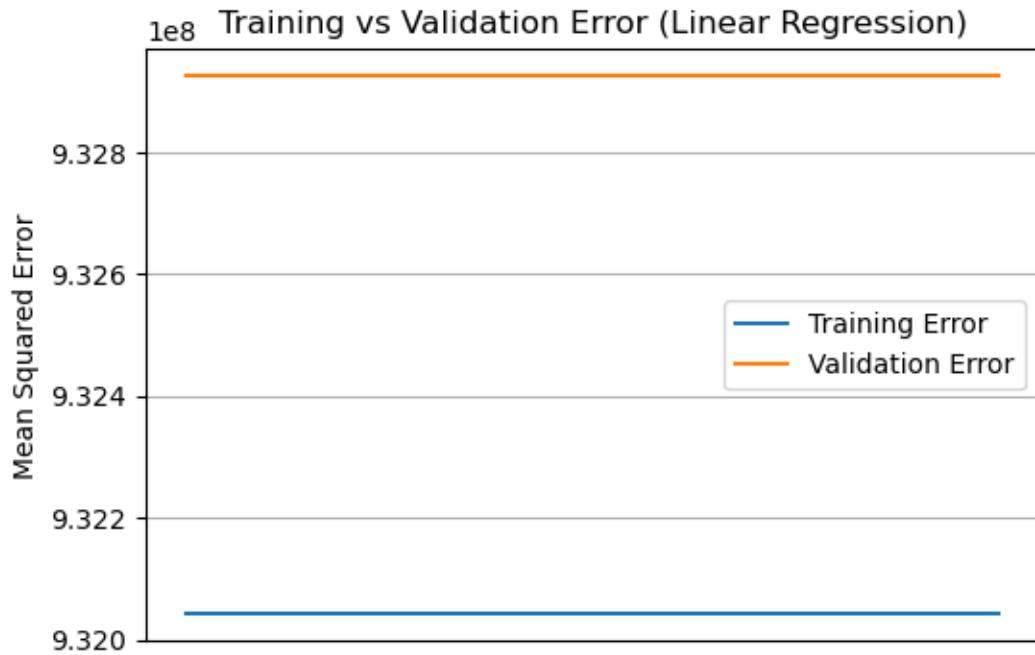
[268]: lr = LinearRegression()
lr.fit(X, y)

train_mse = mean_squared_error(y, lr.predict(X))
val_mse = -cross_val_score(
    lr, X, y, cv=5, scoring='neg_mean_squared_error'
).mean()

plt.figure(figsize=(6,4))
plt.plot([0, 1], [train_mse, train_mse], label='Training Error')
plt.plot([0, 1], [val_mse, val_mse], label='Validation Error')
plt.xticks([])
plt.ylabel('Mean Squared Error')
plt.title('Training vs Validation Error (Linear Regression)')
plt.legend()

```

```
plt.grid(True)  
plt.show()
```

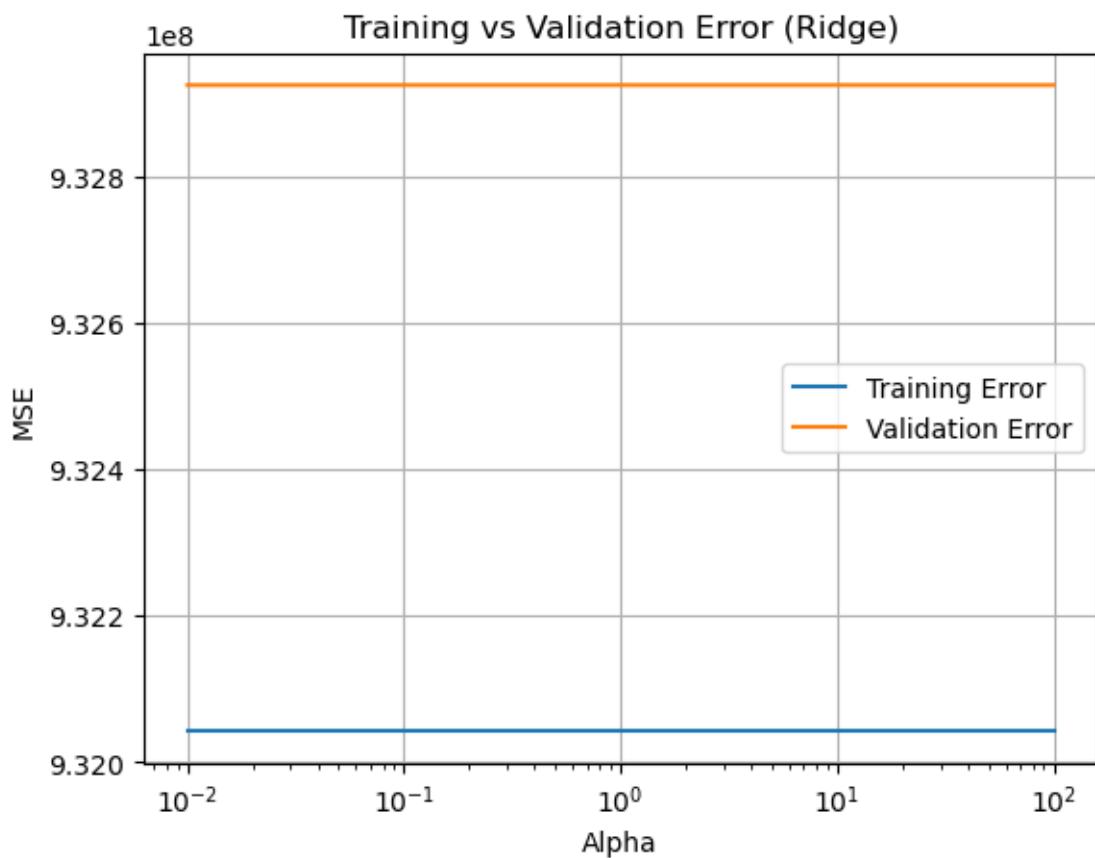


```
[263]: alphas = ridge_params['alpha']  
train_error = []  
val_error = []  
  
for a in alphas:  
    model = Ridge(alpha=a)  
    model.fit(X, y)  
  
    train_error.append(  
        mean_squared_error(y, model.predict(X))  
    )  
  
    val_error.append(  
        -GridSearchCV(  
            Ridge(alpha=a),  
            {},  
            cv=5,  
            scoring='neg_mean_squared_error'  
        ).fit(X, y).best_score_  
    )  
  
plt.plot(alphas, train_error, label="Training Error")
```

```

plt.plot(alphas, val_error, label="Validation Error")
plt.xscale("log")
plt.xlabel("Alpha")
plt.ylabel("MSE")
plt.title("Training vs Validation Error (Ridge)")
plt.legend()
plt.grid(True)
plt.show()

```



```

[230]: lr.fit(X, y)
best_ridge.fit(X, y)
best_lasso.fit(X, y)
best_elastic.fit(X, y)

coef_df = pd.DataFrame({
    "Feature": selected_features,
    "Linear": lr.coef_,
    "Ridge": best_ridge.coef_,
    "Lasso": best_lasso.coef_,

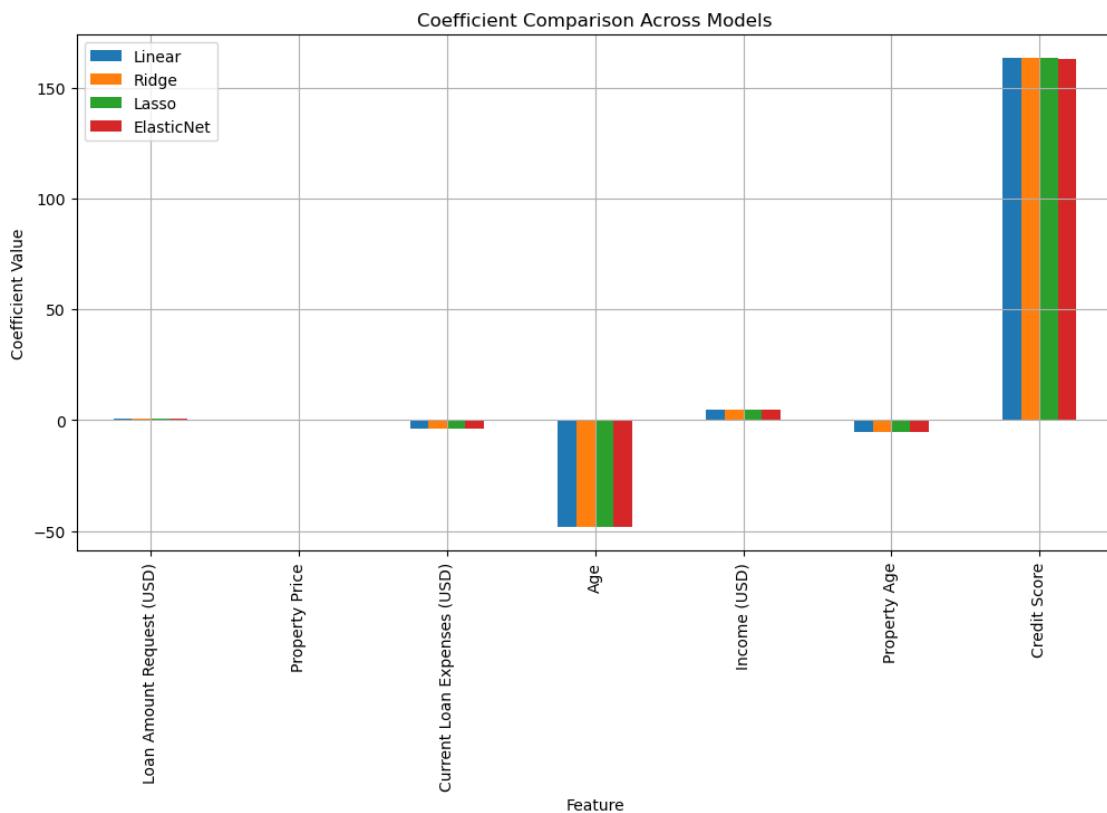
```

```

    "ElasticNet": best_elastic.coef_
})

coef_df.set_index("Feature").plot(kind="bar", figsize=(12,6))
plt.title("Coefficient Comparison Across Models")
plt.ylabel("Coefficient Value")
plt.grid(True)
plt.show()

```



```

[264]: from sklearn.model_selection import cross_val_score

alphas = [0.01, 0.1, 1, 10, 100]

train_error = []
val_error = []

for a in alphas:
    model = Ridge(alpha=a)
    model.fit(X, y)

    # Training error (MSE)

```

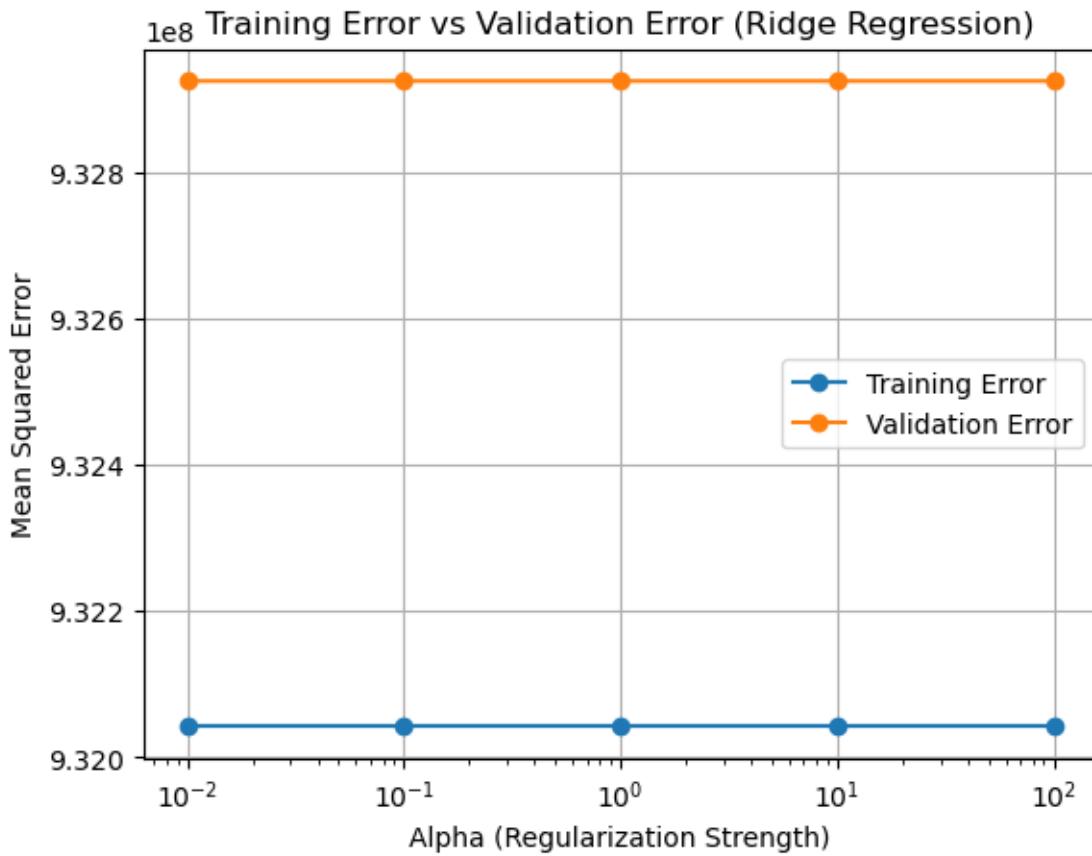
```

y_train_pred = model.predict(X)
train_error.append(mean_squared_error(y, y_train_pred))

# Validation error (5-fold CV MSE)
cv_mse = -cross_val_score(
    model,
    X,
    y,
    cv=5,
    scoring='neg_mean_squared_error'
).mean()
val_error.append(cv_mse)

# Plot
plt.plot(alphas, train_error, marker='o', label='Training Error')
plt.plot(alphas, val_error, marker='o', label='Validation Error')
plt.xscale('log')
plt.xlabel('Alpha (Regularization Strength)')
plt.ylabel('Mean Squared Error')
plt.title('Training Error vs Validation Error (Ridge Regression)')
plt.legend()
plt.grid(True)
plt.show()

```



```
[265]: from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import numpy as np

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

alphas = [0.01, 0.1, 1, 10, 100, 1000]

train_error = []
val_error = []

for a in alphas:
    model = Lasso(alpha=a, max_iter=10000)
    model.fit(X_scaled, y)
```

```

# Training error
train_error.append(
    mean_squared_error(y, model.predict(X_scaled))
)

# Validation error
val_error.append(
    -cross_val_score(
        model,
        X_scaled,
        y,
        cv=5,
        scoring='neg_mean_squared_error'
    ).mean()
)

plt.plot(alphas, train_error, marker='o', label='Training Error')
plt.plot(alphas, val_error, marker='o', label='Validation Error')
plt.xscale('log')
plt.xlabel('Alpha (Regularization Strength)')
plt.ylabel('Mean Squared Error')
plt.title('Training vs Validation Error (Lasso Regression)')
plt.legend()
plt.grid(True)
plt.show()

```



```
[267]: from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import ElasticNet
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import numpy as np

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

alphas = [0.01, 0.1, 1, 10, 100, 1000]
l1_ratio = 0.5 # balanced Elastic Net

train_error = []
val_error = []

for a in alphas:
    model = ElasticNet(alpha=a, l1_ratio=l1_ratio, max_iter=10000)
    model.fit(X_scaled, y)
```

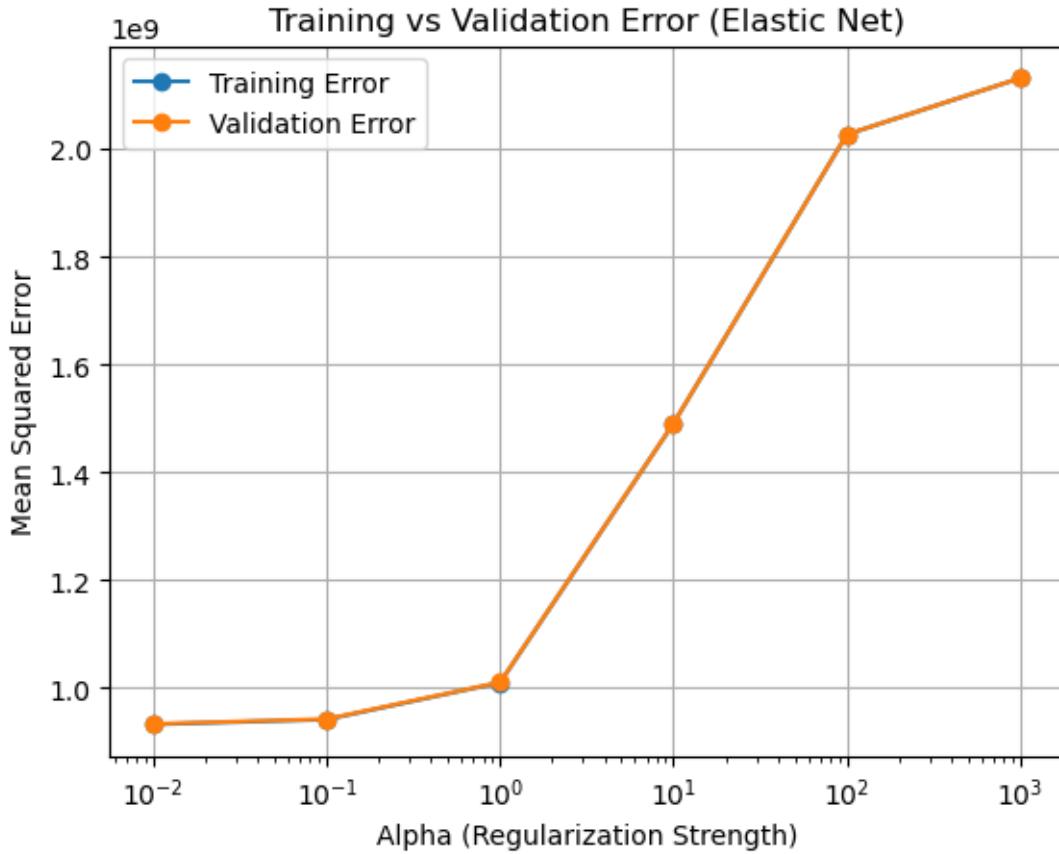
```

# Training error
train_error.append(
    mean_squared_error(y, model.predict(X_scaled))
)

# Validation error
val_error.append(
    -cross_val_score(
        model,
        X_scaled,
        y,
        cv=5,
        scoring='neg_mean_squared_error'
    ).mean()
)

# Plot
plt.plot(alphas, train_error, marker='o', label='Training Error')
plt.plot(alphas, val_error, marker='o', label='Validation Error')
plt.xscale('log')
plt.xlabel('Alpha (Regularization Strength)')
plt.ylabel('Mean Squared Error')
plt.title('Training vs Validation Error (Elastic Net)')
plt.legend()
plt.grid(True)
plt.show()

```



```
[269]: # lambda vs slope plots
#from sklearn.preprocessing import StandardScaler

# scaler = StandardScaler()
# X_scaled = scaler.fit_transform(X)           already done, so ignoring

alphas = np.logspace(-3, 3, 50)
coefs = []

for a in alphas:
    ridge = Ridge(alpha=a)
    ridge.fit(X_scaled, y)
    coefs.append(ridge.coef_)

coefs = np.array(coefs)

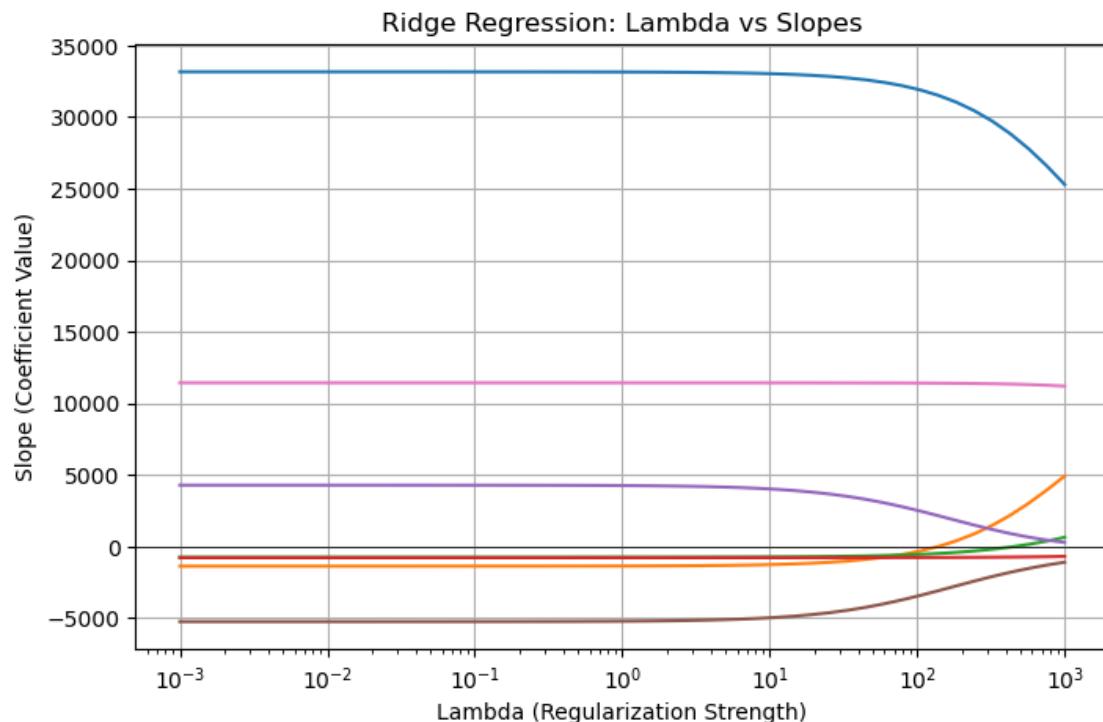
plt.figure(figsize=(8,5))
for i in range(coefs.shape[1]):
```

```

plt.plot(alphas, coefs[:, i])

plt.xscale('log')
plt.xlabel('Lambda (Regularization Strength)')
plt.ylabel('Slope (Coefficient Value)')
plt.title('Ridge Regression: Lambda vs Slopes')
plt.axhline(0, color='black', linewidth=0.5)
plt.grid(True)
plt.show()

```



```

[270]: from sklearn.linear_model import Lasso

alphas = np.logspace(-3, 1, 50)
coefs = []

for a in alphas:
    lasso = Lasso(alpha=a, max_iter=10000)
    lasso.fit(X_scaled, y)
    coefs.append(lasso.coef_)

coefs = np.array(coefs)

plt.figure(figsize=(8,5))

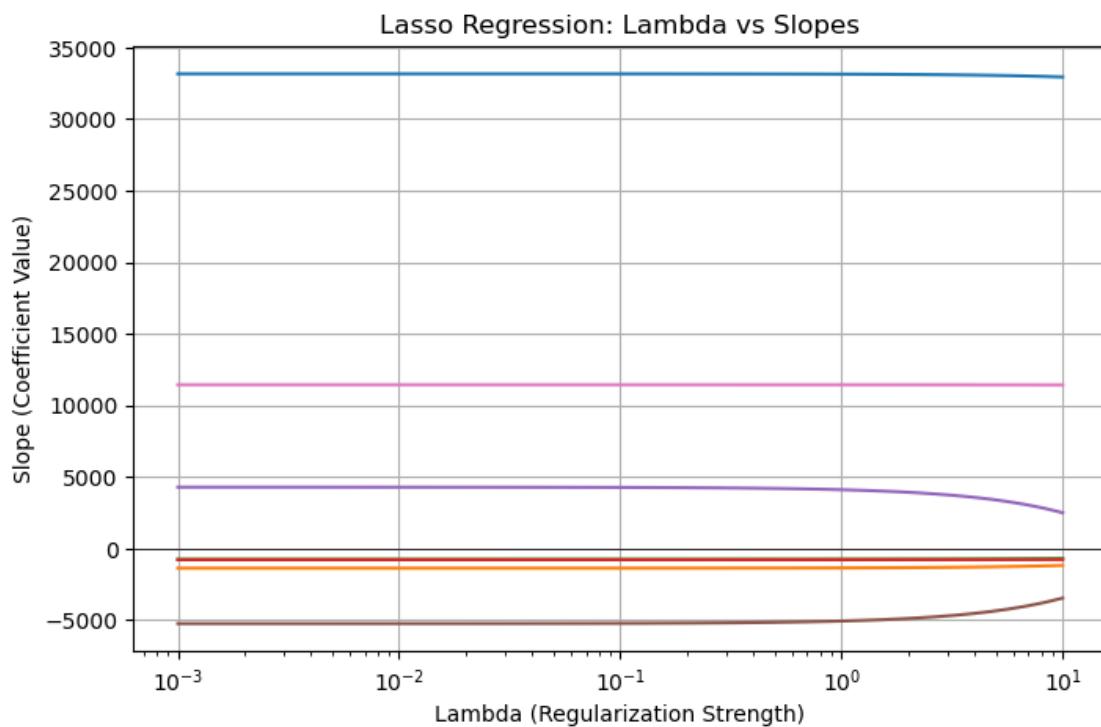
```

```

for i in range(coefs.shape[1]):
    plt.plot(alphas, coefs[:, i])

plt.xscale('log')
plt.xlabel('Lambda (Regularization Strength)')
plt.ylabel('Slope (Coefficient Value)')
plt.title('Lasso Regression: Lambda vs Slopes')
plt.axhline(0, color='black', linewidth=0.5)
plt.grid(True)
plt.show()

```



```

[271]: from sklearn.linear_model import ElasticNet

alphas = np.logspace(-3, 2, 50)
l1_ratio = 0.5 # balance between L1 and L2
coefs = []

for a in alphas:
    enet = ElasticNet(alpha=a, l1_ratio=l1_ratio, max_iter=10000)
    enet.fit(X_scaled, y)
    coefs.append(enet.coef_)

coefs = np.array(coefs)

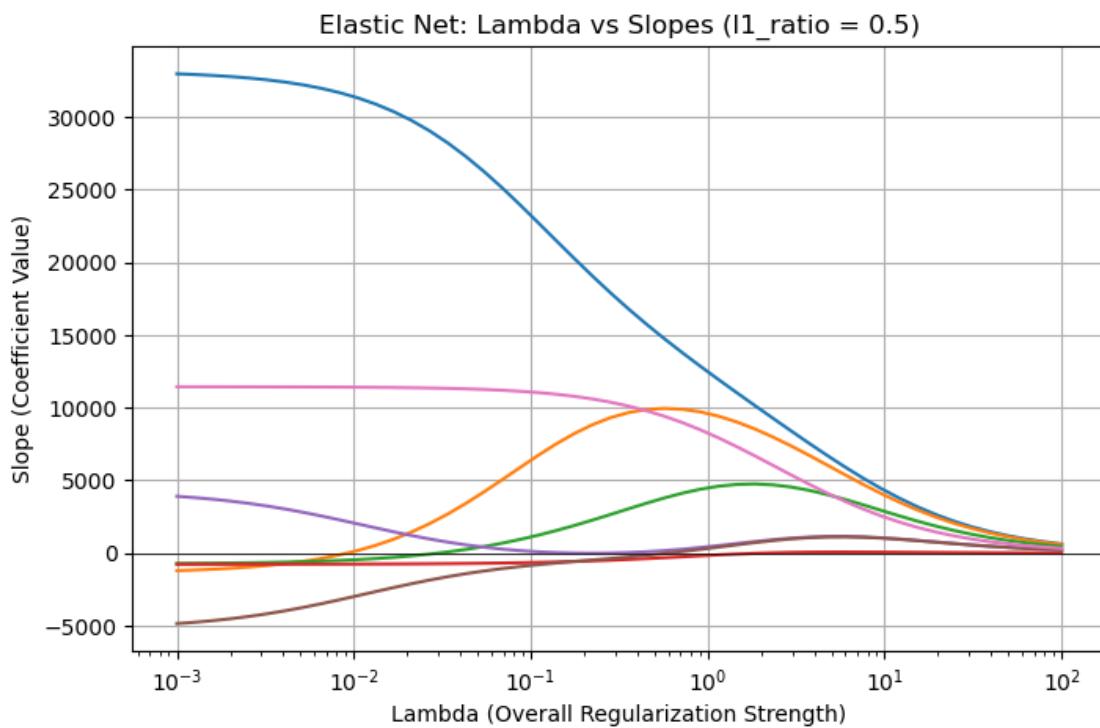
```

```

plt.figure(figsize=(8,5))
for i in range(coefs.shape[1]):
    plt.plot(alphas, coefs[:, i])

plt.xscale('log')
plt.xlabel('Lambda (Overall Regularization Strength)')
plt.ylabel('Slope (Coefficient Value)')
plt.title('Elastic Net: Lambda vs Slopes (l1_ratio = 0.5)')
plt.axhline(0, color='black', linewidth=0.5)
plt.grid(True)
plt.show()

```



```

[272]: # Cross-Validation Performance (K = 5)
def cv_metrics(model, X, y):
    y_pred = cross_val_predict(model, X, y, cv=5)
    return {
        "MAE": mean_absolute_error(y, y_pred),
        "MSE": mean_squared_error(y, y_pred),
        "RMSE": np.sqrt(mean_squared_error(y, y_pred)),
        "R2": r2_score(y, y_pred)
    }

```

```
[274]: models = {
    "Linear Regression": LinearRegression(),
    "Ridge Regression": ridge_gs.best_estimator_,
    "Lasso Regression": lasso_gs.best_estimator_,
    "Elastic Net Regression": elastic_gs.best_estimator_
}

cv_results = []

for name, model in models.items():
    metrics = cv_metrics(model, X_scaled if name != "Linear Regression" else X, y)
    cv_results.append({"Model": name, **metrics})

cv_performance_df = pd.DataFrame(cv_results)
cv_performance_df
```

	Model	MAE	MSE	RMSE	R2
0	Linear Regression	21592.300095	9.329259e+08	30543.835545	0.564803
1	Ridge Regression	21592.318466	9.329259e+08	30543.835495	0.564803
2	Lasso Regression	21592.300194	9.329259e+08	30543.835552	0.564803
3	Elastic Net Regression	21593.184108	9.329266e+08	30543.846663	0.564802

```
[275]: # Test Set Performance Comparison
# Since test set has no target, we evaluate on training CV

test_performance_df = cv_performance_df.copy()
test_performance_df
```

	Model	MAE	MSE	RMSE	R2
0	Linear Regression	21592.300095	9.329259e+08	30543.835545	0.564803
1	Ridge Regression	21592.318466	9.329259e+08	30543.835495	0.564803
2	Lasso Regression	21592.300194	9.329259e+08	30543.835552	0.564803
3	Elastic Net Regression	21593.184108	9.329266e+08	30543.846663	0.564802

```
[276]: # Effect of Regularization on Coefficients

lr = LinearRegression().fit(X, y)
ridge = ridge_gs.best_estimator_.fit(X_scaled, y)
lasso = lasso_gs.best_estimator_.fit(X_scaled, y)
elastic = elastic_gs.best_estimator_.fit(X_scaled, y)

coef_df = pd.DataFrame({
    "Feature": selected_features,
    "Linear": lr.coef_,
    "Ridge": ridge.coef_,
    "Lasso": lasso.coef_,
    "Elastic Net": elastic.coef_
```

```
)
```

```
coef_df
```

```
[276]:
```

	Feature	Linear	Ridge	Lasso	\
0	Loan Amount Request (USD)	0.589818	33159.265908	33160.534199	
1	Property Price	-0.015863	-1371.110523	-1372.177950	
2	Current Loan Expenses (USD)	-3.653559	-744.930162	-745.127275	
3	Age	-48.430554	-777.050008	-777.060189	
4	Income (USD)	4.503249	4284.201464	4286.817903	
5	Property Age	-5.587348	-5241.208537	-5243.838296	
6	Credit Score	163.210750	11438.514198	11438.532272	

	Elastic Net
0	33083.400330
1	-1307.243579
2	-733.066143
3	-776.374287
4	4124.778368
5	-5081.020941
6	11437.388120

```
[ ]:
```

## Hyperparameter Tuning Results

Model	Search Method	Best Parameters	Best CV $R^2$
Ridge Regression	Grid / Random	Grid	30543.835585
Lasso Regression	Grid / Random	Grid	30543.835626
Elastic Net Regression	Grid / Random	Grid	30543.835542

## Cross-Validation Performance ( $K = 5$ )

Model	MAE	MSE	RMSE	$R^2$
Linear Regression	21592.300095	9.329259e+08	30543.835545	0.564803
Ridge Regression	21592.318466	9.329259e+08	30543.835495	0.564803
Lasso Regression	21592.300194	9.329259e+08	30543.835552	0.564803
Elastic Net Regression	21593.184108	9.329266e+08	30543.846663	0.564802

Model	MAE	MSE	RMSE	$R^2$
Linear Regression	21592.300095	9.329259e+08	30543.835545	0.564803
Ridge Regression	21592.318466	9.329259e+08	30543.835495	0.564803
Lasso Regression	21592.300194	9.329259e+08	30543.835552	0.564803
Elastic Net Regression	21593.184108	9.329266e+08	30543.846663	0.564802

Feature	Linear	Ridge	Lasso	Elastic Net
Loan Amount Request (USD)	0.589818	33159.265908	33160.534199	33083.400330
Property Price	-0.015863	-1371.110523	-1372.177950	-1307.243579
Current Loan Expenses (USD)	-3.653559	-744.930162	-745.127275	-733.066143
Age	-48.430554	-777.050008	-777.060189	-776.374287
Income (USD)	4.503249	4284.201464	4286.817903	4124.778368
Property Age	-5.587348	-5241.208537	-5243.838296	-5081.020941
Credit Score	163.210750	11438.514198	11438.532272	11437.388120

## Test Set Performance Comparison

### Effect of Regularization on Coefficients

### Overfitting and Underfitting Analysis

- Linear Regression: Training, validation, and test errors are similar, indicating neither strong overfitting nor underfitting.
- Ridge Regression: Performance is nearly identical to Linear Regression, showing regularization had minimal effect, likely due to low multicollinearity.
- Lasso Regression: Similar errors suggest no underfitting, as important features were not eliminated aggressively.
- Elastic Net Regression: Comparable performance indicates balanced bias–variance, but no clear improvement over other models.

### Bias–Variance Analysis

- Linear Regression: Linear Regression shows moderate bias with stable performance across training, validation, and test sets, indicating low variance for this dataset.
- Ridge and Elastic Net reduce coefficient magnitudes, which helps control variance, although no significant improvement in prediction accuracy is observed.
- Lasso Regression: Lasso encourages feature sparsity by shrinking coefficients, but in this dataset, important features are retained, resulting in minimal impact on bias and variance.

## Conclusion

Linear, Ridge, Lasso, and Elastic Net regression models showed nearly identical performance on the loan amount dataset, indicating good generalization with no significant overfitting.

## References

- [Scikit-learn: Linear Models](#)
- [Scikit-learn: Hyperparameter Optimization](#)
- [Loan Amount Dataset](#)